

Jason Brownlee

Clever Algorithms

Modern Artificial Intelligence Recipes

Clever Algorithms: Modern Artificial Intelligence Recipes

© Copyright 2010 Jason Brownlee. Some Rights Reserved.

License Summary

This work is licensed under a Creative Commons Attribution-Noncommercial-Share Alike 2.5 Australia License.

You are free:

- **to Share** - to copy, distribute and transmit the work
- **to Remix** - to adapt the work

Under the following conditions:

- **Attribution** - You must attribute the work in the manner specified by the author or licensor (but not in any way that suggests that they endorse you or your use of the work).
- **Noncommercial** - You may not use this work for commercial purposes.
- **Share Alike** - If you alter, transform, or build upon this work, you may distribute the resulting work only under the same or similar license to this one.

With the understanding that:

- **Waiver** - Any of the above conditions can be waived if you get permission from the copyright holder.
- **Other Rights** - In no way are any of the following rights affected by the license:
 - Your fair dealing or fair use rights;
 - The author's moral rights;
 - Rights other persons may have either in the work itself or in how the work is used, such as publicity or privacy rights.
- **Notice** - For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to this web page:
<http://creativecommons.org/licenses/by-nc-sa/2.5/au>

The full terms of the license are located on this web page:

<http://creativecommons.org/licenses/by-nc-sa/2.5/au/legalcode>

Preface

About the book

The need for this project was born of frustration while working towards my Ph.D. I was investigating optimization algorithms and was implementing a large number of them for a software platform called the Optimization Algorithm Toolkit (OAT)¹. Each algorithm required considerable effort to locate the relevant source material (from books, papers, articles, and existing implementations), decipher and interpret the technique, then to finally attempt to piece together a functional implementation.

Taking a broader perspective, I realized that the communication of algorithmic techniques in the field of Artificial Intelligence was clearly a difficult and outstanding open problem. Generally, algorithm descriptions are:

- *Incomplete*: many techniques are ambiguously described, partially described, or not described at all.
- *Inconsistent*: a given technique may be described using a variety of formal and semi-formal methods that vary across different techniques, limiting the transferability of skills an audience requires (such as mathematics, pseudo code, program code, and narratives). An inconsistent representation for techniques mean that the skills used to understand and internalize one technique may not be transferable to realizing different techniques or even extensions of the same technique.
- *Distributed*: the description of data structures, operations, and parameterization of a given technique may span a collection of papers, articles, books, and source code published over a number of years, the access to which may be restricted and difficult to obtain.

For the practitioner, a badly described algorithm may be simply frustrating, where the gaps in available information are filled with intuition and ‘best guess’. At the other end of the spectrum, a badly described algorithm may an example of bad science and the failure of the scientific method, where the inability to understand and implement a technique may prevent the replication of results, the application, or the investigation and extension of a technique.

¹OAT located at <http://optalgtoolkit.sourceforge.net>

The software I produced provided a first step solution to this problem: a set of working algorithms implemented in a (somewhat) consistent way and downloaded from a single location (features likely provided of any library of artificial intelligence techniques). The next logical step needed to address this problem is to develop a methodology that anybody can follow. The strategy to address the open problem of poor technique communication is to present complete algorithm descriptions (rather than implementations) in a consistent manner, and in a centralized location. This book is the outcome of developing such a strategy that not only provides a methodology for standardized algorithm descriptions, but provides a large corpus of complete and consistent algorithm descriptions in a single centralized location.

The algorithms described in this work are practical, interesting, and fun, and the goal of this project was to promote these features by making algorithms from the field more accessible, usable, and understandable. This project was developed over a number years though a lot of writing, discussion, and revision. The content was developed and released publicly under a permissive license on the website <http://www.CleverAlgorithms.com>, where forerunning technical reports and the content of this book are freely available. I hope that this project has succeeded in some small way and that you too can enjoy applying, learning, and playing with Clever Algorithms.

About the author

Jason Brownlee has a Bachelors degree in Applied Science, a Masters in Information Technology and a Ph.D. in Computer Science from Swinburne University of Technology in Melbourne, Australia. The subject of Jason's Masters research was Niching Genetic Algorithms. Jason's Ph.D. work was in the area of Artificial Immune Systems and involved research into extending the state of Clonal Selection inspired machine learning algorithms and devising new techniques inspired by the structure and function of the acquired immune system. Jason has earned a living as a Consultant on numerous enterprise-level information technology projects in retail, energy, and information services sectors. Jason has also worked as a Software Engineer investigating the use of intelligent agent technology in geospatial and information services domains in the defense sector. Jason has a long standing passion for both practical software engineering and basic research into machine learning and has developed and released many reports, software plug-ins, and software tools. Jason also enjoys writing and maintains a blog located at <http://www.neverreadpassively.com> and can be followed on twitter at <http://twitter.com/jbrownlee>.

Acknowledgments

Jason Brownlee would like to sincerely thank Daniel Angus for early discussions that lead to the inception of this book project. Jason would like to thank Ying Liu for her unrelenting support and patience throughout the development of the project.

Contents

Preface	iii
I Background	1
1 Introduction	3
1.1 What is AI	3
1.2 Problems	8
1.3 Unconventional Optimization	12
1.4 Book Organization	15
1.5 How to Read this Book	17
1.6 Further Reading	18
II Algorithms	21
2 Stochastic Algorithms	23
2.1 Overview	23
2.2 Random Search	24
2.3 Adaptive Random Search	27
2.4 Stochastic Hill Climbing	32
2.5 Iterated Local Search	35
2.6 Guided Local Search	39
2.7 Variable Neighborhood Search	44
2.8 Greedy Randomized Adaptive Search	49
2.9 Scatter Search	54
2.10 Tabu Search	60
2.11 Reactive Tabu Search	65
3 Physical Algorithms	73
3.1 Overview	73
3.2 Simulated Annealing	74
3.3 Adaptive Simulated Annealing	75
3.4 Memetic Algorithm	76

3.5	Extremal Optimization	77
3.6	Cultural Algorithm	78
3.7	Summary	79
4	Evolutionary Algorithms	81
4.1	Overview	81
4.2	Genetic Algorithm	82
4.3	Genetic Programming	83
4.4	Evolutionary Programming	84
4.5	Evolution Strategies	85
4.6	Learning Classifier System	86
4.7	Differential Evolution	87
4.8	Grammatical Evolution	88
4.9	Non-dominated Sorting Genetic Algorithm	89
4.10	Strength Pareto Evolutionary Algorithm	90
4.11	Island Population Genetic Algorithm	91
4.12	Summary	92
5	Probabilistic Algorithms	93
5.1	Overview	93
5.2	Cross-Entropy Method	94
5.3	Population-Based Incremental Learning	95
5.4	Probabilistic Incremental Program Evolution	96
5.5	Compact Genetic Algorithm	97
5.6	Extended Compact Genetic Algorithm	98
5.7	Bayesian Optimization Algorithm	99
5.8	Hierarchical Bayesian Optimization Algorithm	100
5.9	Univariate Marginal Distribution Algorithm	101
5.10	Bivariate Marginal Distribution Algorithm	102
5.11	Gaussian Adaptation	103
5.12	Summary	104
6	Swarm Algorithms	105
6.1	Overview	105
6.2	Particle Swarm Optimization	106
6.3	AntNet	107
6.4	Ant System	108
6.5	MAX-MIN Ant System	109
6.6	Rank-Based Ant System	110
6.7	Ant Colony System	111
6.8	Multiple Ant Colony System	112
6.9	Population-based Ant Colony Optimization	113
6.10	Bees Algorithm	114
6.11	Bacterial Foraging Optimization Algorithm	115

6.12 Summary	116
7 Immune Algorithms	117
7.1 Overview	117
7.2 Clonal Selection Algorithm	118
7.3 Negative Selection Algorithm	119
7.4 Artificial Immune Recognition System	120
7.5 Immune Network Algorithm	121
7.6 Dendritic Cell Algorithm	122
7.7 Summary	123
 III Extensions	 125
8 Advanced Topics	127
8.1 Programming Paradigms	127
8.2 Devising New Algorithms	128
8.3 Testing Algorithms	128
8.4 Visualizing Algorithms	129
8.5 Saving Algorithm Results	129
8.6 Comparing Algorithms	130
8.7 Summary	130
 Index	 131
 Bibliography	 133

Part I

Background

Chapter 1

Introduction

Welcome to Clever Algorithms! This is a handbook of recipes for computational problem solving techniques from the fields of Computational Intelligence, Biologically Inspired Computation, and Metaheuristics. Clever Algorithms are interesting, practical, and fun to learn about and implement. Research scientists may be interested in browsing algorithm inspirations in search of an interesting system or process analogues to investigate. Developers and software engineers may compare various problem solving algorithms and technique-specific guidelines. Practitioners, students, and interested amateurs may implement state-of-the-art algorithms to address business or scientific needs, or simply play with the fascinating systems they represent.

This introduction chapter provides relevant background information on Artificial Intelligence and Algorithms. The core of the book provides a large corpus of algorithm described in a complete and consistent manner. The final chapter covers some advanced topics to consider once a number of algorithms have been mastered. This book has been designed as a reference text rather than being read cover-to-cover, where specific techniques are looked up, or where the algorithms across whole fields of study are browsed. This book is an algorithm handbook and a technique guidebook, and I hope you find something useful.

1.1 What is AI

1.1.1 Artificial Intelligence

The field of classical *Artificial Intelligence* (AI) coalesced in the 1950s drawing on an understanding of the brain from neuroscience, the new mathematics of information theory, control theory referred to as cybernetics, and the dawn of the digital computer. AI is a cross-disciplinary field of research generally concerned with developing and investigating systems that operate or act intelligently. It is generally considered a discipline in the field of computer science given the strong focus on computation.

Russell and Norvig provide a perspective that defines Artificial Intelligence in four categories: (1) systems that think like humans, (2) systems that act like humans, (3)

systems that think rationally, (4) systems that act rationally [116]. In their definition, acting like a human suggests that a system can do some specific things humans can do, this includes fields such as the Turing test, natural language processing, automated reasoning, knowledge representation, machine learning, computer vision, and robotics. Thinking like a human suggests systems that model the cognitive information processing properties of humans, for example a general problem solver and systems that build internal models of their world. Thinking rationally suggests laws of rationalism and structured thought, such as syllogisms and formal logic. Finally, acting rationally suggests systems that do rational things such as expected utility maximization and rational agents.

Luger and Stubblefield suggest that AI is a sub-field of computer science concerned with the automation of intelligence, and like other sub-fields of computer science has both theoretical concerns (*how and why do the systems work?*) and application concerns (*where and when can the systems be used?*) [90]. They suggest a strong empirical focus to research, because although there may be a strong desire for mathematical analysis, the systems themselves defy analysis given their complexity. The machines and software investigated in AI are not black boxes, rather analysis proceeds by observing the systems interactions with their environment, followed by an internal assessment of the system to relate its structure back to their behavior.

Artificial Intelligence is therefore concerned with investigating mechanisms that underlie intelligence and intelligence behavior. The traditional approach toward designing and investigating AI (the so-called ‘good old fashioned’ AI) has been to employ a symbolic basis for these mechanisms. A newer approach historically referred to as scruffy artificial intelligence or soft computing does not necessarily use a symbolic basis, instead patterning these mechanisms after biological or natural processes. This represents a modern paradigm shift in interest from symbolic knowledge representations, to inference strategies for adaptation and learning, and has been referred to as neat versus scruffy approaches to AI. The neat philosophy is concerned with formal symbolic models of intelligence that can explain *why* they work, whereas the scruffy philosophy is concerned with intelligent strategies that explain *how* they work [121].

Neat AI

The traditional stream of AI concerns a top down perspective of problem solving, generally involving symbolic representations and logic processes that most importantly can explain why they work. The successes of this prescriptive stream include a multitude of specialist approaches such as rule-based expert systems, automatic theorem provers, and operations research techniques that underly modern planning and scheduling software. Although traditional approaches have resulted in significant success they have their limits, most notably scalability. Increases in problem size result in an unmanageable increase in the complexity of such problems meaning that although traditional techniques can guarantee an optimal, precise, or true solution, the computational execution time or computing memory required can be intractable.

Scruffy AI

There have been a number of thrusts in the field of AI toward less crisp techniques that are able to locate approximate, imprecise, or partially-true solutions to problems with a reasonable cost of resources. Such approaches are typically *descriptive* rather than *prescriptive*, describing a process for achieving a solution (how), but not explaining why they work (like the neater approaches).

Scruffy AI approaches are defined as relatively simple procedures that result in complex emergent and self-organizing behavior that can defy traditional reductionist analyses, the effects of which can be exploited for quickly locating approximate solutions to intractable problems. A common characteristic of such techniques is the incorporation of randomness in their processes resulting in robust probabilistic and stochastic decision making contrasted to the sometimes more fragile determinism of the crisp approaches. Another important common attribute is the adoption of an inductive rather than deductive approach to problem solving, generalizing solutions or decisions from sets of specific observations made by the system.

1.1.2 Natural Computation

An important perspective on scruffy Artificial Intelligence is the motivation and inspiration for the core information processing strategy of a given technique. Computers can only do what they are instructed, therefore a consideration is to distill information processing from other fields of study, such as the physical world and biology. The study of biologically motivated computation is called Biologically Inspired Computing [33], and is one of three related fields of Natural Computing [47, 48, 107]. Natural Computing is an interdisciplinary field concerned with the relationship of computation and biology, which in addition to Biologically Inspired Computing is also comprised of Computationally Motivated Biology and Computing with Biology [108, 92].

Biologically Inspired Computation

Biologically Inspired Computation is computation inspired by biological metaphor, also referred to as *Biomimicry*, and *Biomemetics* in other engineering disciplines [29, 20]. The intent of this field is to devise mathematical and engineering tools to generate solutions to computation problems. The field involves using procedures for finding solutions abstracted from the natural world for addressing computationally phrased problems.

Computationally Motivated Biology

Computationally Motivated Biology involves investigating biology using computers. The intent of this area is to use information sciences and simulation to model biological systems in digital computers with the aim to replicate and better understand behaviors in biological systems. The field facilitates the ability to better understand life-as-it-is and investigate life-as-it-could-be. Typically, work in this sub-field is not concerned with the construction of mathematical and engineering tools, rather it is focused on simulating

natural phenomena. Common examples include Artificial Life, Fractal Geometry (L-systems, Iterative Function Systems, Particle Systems, Brownian motion), and Cellular Automata. A related field is that of Computational Biology generally concerned with modeling biological systems and the application of statistical methods such as in the sub-field of Bioinformatics.

Computation with Biology

Computation with Biology is the investigation of substrates other than silicon in which to implement computation [1]. Common examples include molecular or DNA Computing and Quantum Computing.

1.1.3 Computational Intelligence

Computational Intelligence is a modern name for the sub-field of AI concerned with sub-symbolic (also called messy, scruffy, and soft) techniques. Computational Intelligence describes techniques that focus on *strategy* and *outcome*. The field broadly covers sub-disciplines that focus on adaptive and intelligence systems, not limited to: Evolutionary Computation, Swarm Intelligence (Particle Swarm and Ant Colony Optimization), Fuzzy Systems, Artificial Immune Systems, and Artificial Neural Networks [37, 109]. This section provides a brief summary of each of the five primary areas of study.

Evolutionary Computation

A paradigm that is concerned with the investigation of systems inspired by the neo-Darwinian theory of evolution by means of natural selection. Popular evolutionary algorithms include the Genetic Algorithm, Evolution Strategy, Genetic and Evolutionary Programming, and Differential Evolution [5, 6]. The evolutionary process is considered an adaptive strategy and is typically applied to search and optimization domains [66, 74].

Swarm Intelligence

A paradigm that considers collective intelligence as a behavior that emerges through the interaction and cooperation of large numbers of lesser intelligent agents. The paradigm consists of two dominant sub-fields (1) Ant Colony Optimization that investigates probabilistic algorithms inspired by the stigmergy and foraging behavior of ants [24, 35], and (2) Particle Swarm Optimization that investigates probabilistic algorithms inspired by the flocking and foraging behavior of birds and fish [119]. Like evolutionary computation, swarm intelligence-based techniques are considered adaptive strategies and are typically applied to search and optimization domains.

Artificial Neural Networks

Neural Networks are a paradigm that is concerned with the investigation of architectures and learning strategies inspired by the modeling of neurons in the brain [22]. Learning

strategies are typically divided into supervised and unsupervised which manage environmental feedback in different ways. Neural network learning processes are considered adaptive learning and are typically applied to function approximation and pattern recognition domains.

Fuzzy Intelligence

Fuzzy Intelligence is a paradigm that is concerned with the investigation of fuzzy logic, which is a form of logic that is not constrained to true and false like propositional logic, but rather functions which define approximate truth or degrees of truth [146]. Fuzzy logic and fuzzy systems are a logic system used as a reasoning strategy and are typically applied to expert system and control system domains.

Artificial Immune Systems

A collection of approaches inspired by the structure and function of the acquired immune system of vertebrates. Popular approaches include clonal selection, negative selection, dendritic cell algorithm, and immune network algorithms. The immune-inspired adaptive processes vary in strategy and show similarities to the fields of Evolutionary Computation and Artificial Neural Networks, and are typically used for optimization and pattern recognition domains [34].

1.1.4 Metaheuristics

Another popular name for the strategy-outcome perspective of scruffy AI is *Metaheuristics*. A heuristic is an algorithm that locates ‘good enough’ solutions to a problem without concern for whether the solution can be proven to be correct or optimal [97]. Heuristic methods trade-off concerns such as precision, quality, and accuracy in favor of computational effort (space and time efficiency). The Greedy search procedure that only takes cost-improving steps is an example of heuristic method.

Like heuristics, Metaheuristic may be considered a general algorithmic framework that can be applied to different optimization problems with relative few modifications to make them adapted to a specific problem [63, 129]. The difference is that Metaheuristics are intended to extend the capabilities of heuristics by combining one or more heuristic methods (referred to as procedures) using a higher-level strategy (hence ‘meta’). A procedure in a metaheuristic is considered black-box in that little (if any) prior knowledge is known about it by the meta-heuristic, and as such it may be replaced with a different procedure. Procedures may be as simple as the manipulation of a representation, or as complex as another complete metaheuristic. Some examples of metaheuristics include iterated local search, tabu search, the genetic algorithm, ant colony optimization, and simulated annealing.

Blum and Roli outline nine properties of metaheuristics [23], as follows:

- Metaheuristics are strategies that “guide” the search process.

- The goal is to efficiently explore the search space in order to find (near-)optimal solutions.
- Techniques which constitute metaheuristic algorithms range from simple local search procedures to complex learning processes.
- Metaheuristic algorithms are approximate and usually non-deterministic.
- They may incorporate mechanisms to avoid getting trapped in confined areas of the search space.
- The basic concepts of metaheuristics permit an abstract level description.
- Metaheuristics are not problem-specific.
- Metaheuristics may make use of domain-specific knowledge in the form of heuristics that are controlled by the upper level strategy.
- Today's more advanced metaheuristics use search experience (embodied in some form of memory) to guide the search.

Hyperheuristics are yet another extension that focuses on heuristics that modify their parameters (online or offline) to improve the efficacy of solution, or the efficiency of the computation. Hyperheuristics provide high-level strategies that may employ machine learning and adapt their search behavior by modifying the application of the sub-procedures or even which procedures are used (operating on the space of heuristics which in turn operate within the problem domain) [27, 28].

1.1.5 Clever Algorithms

This book is concerned with Clever Algorithms which are algorithms drawn from many sub-fields of Artificial Intelligence not limited to the scruffy fields of Biologically Inspired Computation, Computational Intelligence and Metaheuristics. The term *Clever Algorithms* is intended to unify a collection of interesting and useful computational tools under a consistent and accessible banner. An alternative name (*Inspired Algorithms*) was considered, although ultimately rejected given that not all of the algorithms to be described in the project have an inspiration (specifically a biological or physical inspiration) for their computational strategy. The set of algorithms described in this book may generally be referred to as ‘unconventional optimization algorithms’ (for example, see [30]), as optimization is the main form of computation provided by the listed approaches. A technically more appropriate name for these approaches is Stochastic Global Optimization (for example, see [141] and [91]).

1.2 Problems

Algorithms from the fields of Computational Intelligence, Biologically Inspired Computing, and Metaheuristics are applied to difficult problems, to which more traditional

approaches may not be suited. Michalewicz and Fogel propose five reasons why problems may be difficult [97] (page 11):

- The number of possible solutions in the search space is so large as to forbid an exhaustive search for the best answer.
- The problem is so complicated that just to facilitate any answer at all, we have to use such simplified models of the problem that any result is essentially useless.
- The evaluation function that describes the quality of any proposed solution is noisy or varies with time, thereby requiring not just a single solution but an entire series of solutions.
- The possible solutions are so heavily constrained that constructing even one feasible answer is difficult, let alone searching for an optimal solution.
- The person solving the problem is inadequately prepared or imagines some psychological barrier that prevents them from discovering a solution.

This section introduces two problem formalisms that embody many of the most difficult problems faced by Artificial and Computational Intelligence. They are: Function Optimization and Function Approximation. Each class of problem is described in terms of its general properties, a formalism, and a set of specialized sub-problems. These problem classes provide a tangible framing of the algorithmic techniques described throughout the work.

1.2.1 Function Optimization

Real-world optimization problems and generalizations thereof can be drawn from most fields of science, engineering, and information technology (for a sample see [2, 132]). Importantly, optimization problems have had a long tradition in the fields of Artificial Intelligence in motivating basic research into new problem solving techniques, and for investigating and verifying systemic behavior against benchmark problem instances.

Problem Description

Mathematically, optimization is defined as the search for a combination of parameters commonly referred to as decision variables ($x = \{x_1, x_2, x_3, \dots, x_n\}$) which minimize or maximize some ordinal quantity (c) (typically a scalar called a score or cost) assigned by an objective function or cost function (f), under a set of constraints ($g = \{g_1, g_2, g_3, \dots, g_n\}$). For example, a general minimization case would be as follows: $f(x') \leq f(x), \forall x_i \in x$. Constraints may provide boundaries on decision variables (for example in a real-value hypercube \mathbb{R}^n), or may generally define regions of feasibility and in-feasibility in the decision variable space. In applied mathematics the field may be referred to as Mathematical Programming. More generally the field may be referred to as Global or Function Optimization given the focus on the objective function (for more general information on optimization refer to [75]).

Sub-Fields of Study

The study of optimization is comprised of many specialized sub-fields, based on an overlapping taxonomy that focuses on the principle concerns in the general formalism. For example, with regard to the decision variables, one may consider univariate and multivariate optimization problems. The type of decision variables promotes specialities for continuous, discrete, and permutations of variables. Dependencies between decision variables under a cost function define the fields of Linear Programming, Quadratic Programming, and Nonlinear Programming. A large class of optimization problems can be reduced to discrete sets and are considered in the field of Combinatorial Optimization, to which many theoretical properties are known, most importantly that many interesting and relevant problems cannot be solved by an approach with polynomial time complexity (so-called NP-complete, for example see [105]).

The topography of the response surface for the decision variables under the cost function may be convex, which is a class of functions to which many important theoretical findings have been made, not limited to the fact that location of the local optimal configuration also means the global optimal configuration of decisional variables has been located [25]. Many interesting and real-world optimization problems produce cost surfaces that are non-convex or so called multi-modal¹ (rather than uni-modal) suggesting that there are multiple peaks and valleys. Further, many real-world optimization problems with continuous decision variables cannot be differentiated given their complexity or limited information availability meaning that derivative-based gradient decent methods that are well understood are not applicable, requiring the use of so-called ‘direct search’ (sample or pattern-based) methods [88]. Real-world objective function evaluation may be noisy, discontinuous, dynamic, and the constraints of real-world problem solving may require an approximate solution in limited time or using resources, motivating the need for heuristic approaches.

1.2.2 Function Approximation

The phrasing of real-world problems in the Function Approximation formalism are among the most computationally difficult considered in the broader field of Artificial Intelligence for reasons including: incomplete information, high-dimensionality, noise in the sample observations, and non-linearities in the target function. This section considers the Function Approximation Formalism and related specialization’s as a general motivating problem to contrast and compare with Function Optimization.

Problem Description

Function Approximation is the problem of finding a function (f) that approximates a target function (g), where typically the approximated function is selected based on a

¹Taken from statistics referring to the centers of mass in distributions, although in optimization it refers to ‘regions of interest’ in the search space, in particular valleys in minimization, and peaks in maximization cost surfaces.

sample of observations (x , also referred to as the training set) taken from the unknown target function. In machine learning, the function approximation formalism is used to describe general problem types commonly referred to as pattern recognition, such as classification, clustering, and curve fitting (called a decision or discrimination function). Such general problem types are described in terms of approximating an unknown Probability Density Function (PDF), which underlies the relationships in the problem space, and is represented in the sample data. This ‘function approximation’ perspective of such problems is commonly referred to as statistical machine learning and/or density estimation [50, 22].

Sub-Fields of Study

The function approximation formalism can be used to phrase some of the hardest problems faced by Computer Science, and Artificial Intelligence in particular, such as natural language processing and computer vision. The general process focuses on (i) the collection and preparation of the observations from the target function, (ii) the selection and/or preparation of a model of the target function, and (ii) the application and ongoing refinement of the prepared model. Some important problem-based sub-fields include:

- *Feature Selection* where a feature is considered an aggregation of one-or-more attributes, where only those features that have meaning in the context of the target function are necessary to the modeling function [83, 67].
- *Classification* where observations are inherently organized into labelled groups (classes) and a supervised process models an underlying discrimination function to classify unobserved samples.
- *Clustering* where observations may be organized into groups based on underlying common features, although the groups are unlabeled requiring a process to model an underlying discrimination function without corrective feedback.
- *Curve or Surface Fitting* where a model is prepared that provides a ‘best-fit’ (called a regression) for a set of observations that may be used for interpolation over known observations and extrapolation for observations outside what has been modelled.

The field of Function Optimization is related to Function Approximation, as many-sub-problems of Function Approximation may be defined as optimization problems. Many of the technique paradigms used for function approximation are differentiated based on the representation and the optimization process used to minimize error or maximize effectiveness on a given approximation problem. The difficulty of Function Approximation problems centre around (i) the nature of the unknown relationships between attributes and features, (ii) the number (dimensionality) of attributes and features, and (iii) general concerns of noise in such relationships and the dynamic availability of samples from the target function. Additional difficulties include the incorporation of

prior knowledge (such as imbalance in samples, incomplete information and the variable reliability of data), and problems of invariant features (such as transformation, translation, rotation, scaling and skewing of features).

1.3 Unconventional Optimization

Not all algorithms described in this book are for optimization, although, those that are may be referred to as ‘unconventional’ to differentiate them from the more traditional approaches. Examples of traditional approaches include (but are not not limited) to mathematical optimization algorithms (such as Newton’s method and Gradient descent that uses derivatives to locate a local minimum) and direct search methods (such as the Simplex method and the Nelder-Mead method that use a search pattern to locate optima). Unconventional optimization algorithms are designed for the more difficult problem instances, the attributes of which were introduced in Section 1.2.1. This section introduces some common attributes of this class of algorithm.

1.3.1 Black Box Algorithms

Black Box optimization algorithms are those that exploit little, if any, information from a problem domain in order to devise a solution. They are generalized problem solving procedures that may be applied to a range of problems with very little modification [36]. Domain specific knowledge refers to making use of known relationships between solution representations and the objective cost function. Generally speaking, the less domain specific information incorporated into a technique, the more flexible the technique, although the less efficient it will be for a given problem. For example, ‘random search’ is the most general black box approach and is also the most flexible requiring only the generation of random solutions for a given problem. Random search also has a worst case behavior, that is worse than enumerating an entire search domain given the freedom it has to resample. In practice, the more prior knowledge available about a problem, the more information that should be exploited by a technique in order to efficiently locate a solution for the problem, heuristically or otherwise. Therefore, black box methods are those methods suitable for those problems where little information from the problem domain is available to be used by a problem solving approach.

1.3.2 No Free Lunch

The *No Free Lunch Theorem* of search and optimization by Wolpert and Macready proposes that all black box optimization algorithms are the same for searching for the extremum of a cost function when averaged over all possible functions [144, 143]. The theorem has caused a lot of pessimism and misunderstanding, particularly in relation to the evaluation and comparison of Metaheuristic and Computational Intelligence algorithms.

The implication of the theorem is that searching for the ‘best’ general-purpose black box optimization algorithm is irresponsible as no such procedure is theoretically possi-

ble. The theory applies to stochastic and deterministic optimization algorithms as well as to algorithms that learn and adjust their search strategy over time. It is invariant to the performance measure used and the representation selected. The theorem is an important contribution to computer science, although its implications are theoretical. The original paper was produced at a time when grandiose generalizations were being made as to algorithm, representation, or configuration superiority. The practical impact of the theory is to encourage practitioners to bound claims of applicability for search and optimization algorithms. Wolpert and Macready encouraged effort be put into devising practical problem classes and into the matching of suitable algorithms to problem classes. Further, they compelled practitioners to exploit domain knowledge in optimization algorithm application, which is now an axiom in the field.

1.3.3 Stochastic Optimization

Stochastic optimization algorithms those that use randomness to elicit non-deterministic behaviors, contrasted to purely deterministic procedures. Most algorithms from the fields of Computational Intelligence, Biologically Inspired Computation, and Metaheuristics may be considered to belong the field of Stochastic Optimization. Algorithms that exploit randomness, are not random in behavior, rather they sample a problem space in a biased manner, focusing on areas of interest and neglecting less interesting areas [123]. A class of techniques that focus on the stochastic sampling of a domain are called Markov Chain Monte Carlo (MCMC) algorithms that provide good average performance, quickly, and generally offer a low chance of the worst case performance. Such approaches are suited to problems with many coupled degrees of freedom, for example large, high-dimensional spaces. MCMC approaches involve stochastically sampling from a target distribution function similar to Monte Carlo simulation methods using a process that resembles a biased Markov chain.

- *Monte Carlo* methods are used for selecting a statistical sample to approximate a given target probability density function and are traditionally used in statistical physics. Samples are drawn sequentially and the process may include criteria for rejecting samples and biasing the sampling locations within high-dimensional spaces.
- *Markov Chain* processes provide a probabilistic model for state transitions or moves within a discrete domain called a walk or a chain of steps. A Markov system is only dependent on the current position in the domain in order to probabilistically determine the next step in the walk.

MCMC techniques combine these two approaches to solve integration and optimization problems in large dimensional spaces by generating samples while exploring the space using a Markov chain process, rather than sequentially or independently [3]. The step generation is configured to bias sampling in more important regions of the domain.

Three examples of MCMC techniques include the Metropolis-Hastings algorithm, Simulated annealing for global optimization, and the Gibbs sampler which are commonly employed in the fields of physics, chemistry, statistics, and economics.

1.3.4 Inductive Learning

Many unconventional optimization algorithms employ a process that includes the iterative improvement of candidate solutions against an objective cost function. This process of adaptation is generally a method by which the process obtains characteristics that improve the system's (candidate solution) relative performance in an environment (cost function). This adaptive behavior is commonly achieved through a 'selectionist process' of repetition of the steps: generation, test, and selection. The use of non-deterministic processes mean that the sampling of the domain (the generation step) is typically non-parametric, although guided by past experience.

The method of acquiring information is called inductive learning or learning from example, where the approach uses the implicit assumption that specific examples are representative of the broader information content of the environment, specifically with regard to anticipated need. Many unconventional optimization approaches maintain a single candidate solution, a population of samples, or a compression thereof that provides both an instantaneous representation of all of the information acquired by the process, and the basis for generating and making future decisions.

This method of simultaneously acquiring and improving information from the domain and the optimization of decision making (where to direct future effort) is called the k -armed bandit (two-armed and multi-armed bandit) problem from the field of statistical decision making known as game theory [115, 21]. This formalism considers the capability of a strategy to allocate available resources proportional to the future payoff the strategy is expected to receive. The classic example is the 2-armed bandit problem used by Goldberg to describe the behavior of the genetic algorithm [66]. The example involves an agent that learns which one of the two slot machines provides more return by pulling the handle of each (sampling the domain) and biasing future handle pulls proportional to the expected utility, based on the probabilistic experience with the past distribution of the payoff. The formalism may also be used to understand the properties of inductive learning demonstrated by the adaptive behavior of most unconventional optimization algorithms.

The stochastic iterative process of generate and test can be computationally wasteful, potentially re-searching areas of the problem space already searched, and requiring many trials or samples in order to achieve a 'good enough' solution. The limited use of prior knowledge from the domain (black box) coupled with the stochastic sampling process mean that the adapted solutions are created without top-down insight or instruction can sometimes be interesting, innovative, and even competitive with decades of human expertise [81].

1.4 Book Organization

The remainder of this book is organized into two parts: *Algorithms* that describes a large number of techniques in a complete and a consistent manner presented in a rough algorithm groups, and *Extensions* that reviews more advanced topics suitable for when a number of algorithms have been mastered.

1.4.1 Algorithms

Algorithms are presented in six groups or kingdoms distilled from the broader fields of study each in their own chapter, as follows:

- *Stochastic Algorithms* that focus on the introduction of randomness into heuristic methods (Chapter 2).
- *Physical Algorithms* that focus on methods inspired by physical and social systems (Chapter 3).
- *Evolutionary Algorithms* that focus on methods inspired by evolution by means of natural selection (Chapter 4).
- *Probabilistic Algorithms* that focus on methods that build models and estimate distributions in search domains (Chapter 5).
- *Swarm Algorithms* that focus on methods that exploit the properties of collective intelligence (Chapter 6).
- *Immune Algorithms* that focus on methods inspired by the adaptive immune system of mammals (Chapter 7).

A given algorithm is more than just a procedure or code listing. Each approach is an island of research and the meta-information that define the context of a technique are just as important to understanding and application as abstract recipes and concrete implementations. A standardized algorithm description was adopted to provide a consistent presentation of algorithms with a mixture of softer narrative descriptions, programmatic descriptions both abstract and concrete, and most importantly useful sources for finding out more information about the technique.

The standardized algorithm description template covers the following subjects:

- *Name*: The algorithm name defines the canonical name used to refer to the technique, in addition to common aliases, abbreviations, and acronyms. The name is used as the heading of an algorithm descriptions.
- *Taxonomy*: The algorithm taxonomy defines where a techniques fits into the field, both the specific subfields of Computational Intelligence and Biologically Inspired Computation as well as the broader field of Artificial Intelligence. The taxonomy also provides a context for determining the relationships between algorithms.

- *Inspiration*: (optional) The inspiration describes the specific system or process that provoked the inception of the algorithm. The inspiring system may non-exclusively be natural, biological, physical, or social. The description of the inspiring system may include relevant domain specific theory, observation, nomenclature, and most important must include those salient attributes of the system that are somehow abstractly or conceptually manifest in the technique.
- *Metaphor*: (optional) The metaphor is a description of the technique in the context of the inspiring system or a different suitable system. The features of the technique are made apparent through an analogous description of the features of the inspiring system. The explanation through analogy is not expected to be literal, rather the method is used as an allegorical communication tool. The inspiring system is not explicitly described, this is the role of the ‘inspiration’ topic, which represents a loose dependency for this topic.
- *Strategy*: The strategy is an abstract description of the computational model. The strategy describes the information processing actions a technique shall take in order to achieve an objective. The strategy provides a logical separation between a computational realization (procedure) and a analogous system (metaphor). A given problem solving strategy may be realized as one of a number specific algorithms or problem solving systems.
- *Procedure*: The algorithmic procedure summarizes the specifics of realizing a strategy as a systemized and parameterized computation. It outlines how the algorithm is organized in terms of the computation, data structures, and representations.
- *Heuristics*: The heuristics section describes the commonsense, best practice, and demonstrated rules for applying and configuring a parameterized algorithm. The heuristics relate to the technical details of the techniques procedure and data structures for general classes of application (neither specific implementations nor specific problem instances).
- *Code Listing*: The code listing description provides a minimal but functional version of the technique implemented with a programming language. The code description can be typed into an computer and provide a working execution of the technique. The technique implementation also includes a minimal problem instance to which it is applied, and both the problem and algorithm implementations are complete enough to demonstrate the techniques procedure. The description is presented as a programming source code listing with a terse introductory summary.
- *References*: The references section includes a listing of both primary sources of information about the technique as well as useful introductory sources for novices to gain a deeper understanding of the theory and application of the technique. The description consists of hand-selected reference material including books, peer reviewed conference papers, journal articles, and potentially websites.

Source code examples are included in the algorithm descriptions, and the Ruby Programming Language was selected for use throughout the book. Ruby was selected because it supports the procedural programming paradigm that was adopted to ensure that examples can be easily ported to object-oriented and other paradigms. Additionally, Ruby is interpreted meaning the code can be directly executed without an introduced compilation step, and it is free to download and use from the website². Finally, Ruby is concise, expressive, and supports meta-programming features that improve the readability of code examples. All of the source code for the algorithms presented in this book is available from the books website at <http://www.CleverAlgorithms.com>.

1.4.2 Extensions

There are some some advanced topics that cannot be meaningfully considered until one has a firm grasp of a number of algorithms, and these are discussed at the back of the book. The Advanced Topics chapter addresses topics such as: the use of alternative programming paradigms when implementing clever algorithms, methodologies used when devising entirely new approaches, strategies to consider when testing clever algorithms, visualizing the behavior and results of algorithms, and finally comparing algorithms based on the results they produce using statistical methods. Like the background information provided in this chapter, the extensions provide a gentle introduction and starting point into some advanced topics and plenty of references for seeking a deeper understanding.

1.5 How to Read this Book

This book is a reference text that provides a large compendium of algorithm descriptions. It is a trusted handbook of practical computational recipes to consulted when one is confronted with difficult function optimization and approximation problems. It is also an encompassing guidebook of modern heuristic methods that may be browsed for inspiration, exploration, and general interest.

The audience for this work may be interested with the fields of Computational Intelligence, Biologically Inspired Computation, and Metaheuristics and may count themselves as belonging to one of the following broader groups:

- *Scientists*: Research scientists concerned with theoretically or empirically investigating algorithms, addressing questions such as: *What is the motivating system and strategy for a given technique? What are some algorithms that may be used in a comparison within a given subfield or across subfields?*
- *Engineers*: Programmers and developers concerned with implementing, applying, or maintaining algorithms, addressing questions such as: *What is the algorithm procedure for a given technique? What are the best practice heuristics for employing a given technique?*

²Ruby can be downloaded for free from <http://www.ruby-lang.org>

- *Students*: Undergraduate and graduate students interested in learning about techniques, addressing questions such as: *What are some interesting algorithms to study? How to implement a given approach?*
- *Amateurs*: Practitioners interested in knowing more about algorithms, addressing questions such as: *What classes of techniques exist and what algorithms do they provide? How to conceptualize the computation of a technique?*

1.6 Further Reading

This book is not an introduction to Artificial Intelligence or related sub-fields, nor is it a field guide for a specific class of algorithms. This section provides some pointers to selected books and articles for those readers seeking a deeper understanding of the fields of study to which the Clever Algorithms described in this book belong.

1.6.1 Artificial Intelligence

Artificial Intelligence is large field of study and many excellent texts have been written to introduce the subject. Russell and Novig’s “*Artificial Intelligence: A Modern Approach*” is an excellent introductory text providing a broad and deep review of what the field has to offer and is useful for students and practitioners alike [116]. Luger and Stubblefield’s “*Artificial Intelligence: Structures and Strategies for Complex Problem Solving*” is also an excellent reference text, providing a more empirical approach to the field.

1.6.2 Computational Intelligence

Introductory books for the field of Computational Intelligence generally focus on a handful of specific sub-fields and their techniques. Engelbrecht’s “*Computational Intelligence: An Introduction*” provides a modern and detailed introduction to the field covering classic subjects such as Evolutionary Computation and Artificial Neural Networks, as well as more recent techniques such as Swarm Intelligence and Artificial Immune Systems [37]. Pedrycz’s slightly more dated “*Computational Intelligence: An Introduction*” also provides a solid coverage of the core of the field with some deeper insights into fuzzy logic and fuzzy systems [109].

1.6.3 Biologically Inspired Computation

Computational methods inspired by natural and biologically systems represent a large fraction of the algorithms described in this book. The collection of articles published in de Castro and Von Zuben’s “*Recent Developments in Biologically Inspired Computing*” provide a good overview of the state of the field, and the introductory chapter on need for such methods does an excellent job to motivate the field of study [29]. Forbes’s “*Imitation of Life: How Biology Is Inspiring Computing*” set’s the scene for Natural Computing and the interrelated disciplines, of which Biologically Inspired Computing

is but one useful example [47]. Finally, Benyus’s “*Biomimicry: Innovation Inspired by Nature*” provides a good introduction into the broader related field of a new frontier in science and technology that involves building systems inspired by an understanding of biological systems [20].

1.6.4 Metaheuristics

The field of Metaheuristics was initially constrained to heuristics for applying classical optimization procedures, although has expanded to encompass a broader and diverse set of techniques. Michalewicz and Fogel’s “*How to Solve It: Modern Heuristics*” provides a practical tour of heuristic methods with a consistent set of worked examples [97]. Glover and Kochenberger’s “*Handbook of Metaheuristics*” provides a solid introduction into a broad collection of techniques and their capabilities [63].

1.6.5 The Ruby Programming Language

The Ruby Programming Language is a multi-paradigm dynamic language that appeared in approximately 1995. It’s meta-programming capabilities coupled with concise and readable syntax have made it a popular language of choice for web development, scripting, and application development. The classic reference text for the language is Thomas, Fowler, and Hunt’s “*Programming Ruby: The Pragmatic Programmers’ Guide*” referred to as the ‘pickaxe book’ because of the picture of the pickaxe on the cover [130]. An updated edition is available that covers version 1.9 (compared to 1.8 in the cited version) that will work just as well for use as a reference for the examples in this book. Flanagan and Matsumoto’s “*The Ruby Programming Language*” also provides a seminal reference text with contributions from Yukihiro Matsumoto, the author of the language [46].

Part II

Algorithms

Chapter 2

Stochastic Algorithms

2.1 Overview

This chapter describes Stochastic Algorithms. The majority of the algorithms to be described in the Clever Algorithms project are comprised of probabilistic and stochastic processes. What differentiates the ‘stochastic algorithms’ in this chapter from the remaining algorithms is the specific lack of i) an inspiring system, and ii) a metaphorical explanation. Both ‘inspiration’ and ‘metaphor’ refer to the descriptive elements in the standardized algorithm description.

These described algorithms are predominately global optimization algorithms and metaheuristics that manage the application of an embedded neighborhood exploring (local) search procedure. As such, with the exception of ‘Stochastic Hill Climbing’ and ‘Random Search’ the algorithms may be considered extensions of the multi-start search (also known as multi-restart search). The set of algorithms provide various different strategies by which ‘better’ and varied starting points can be generated and issued to a neighborhood searching technique for refinement, a process that is repeated with potentially improving or unexplored areas to search.

2.2 Random Search

Random Search, RS, Blind Random Search, Blind Search, Pure Random Search, PRS

2.2.1 Taxonomy

Random search belongs to the fields of Stochastic Optimization and Global Optimization. Random search is a direct search method as it does not require derivatives to search a continuous domain. This base approach is related to techniques that provide small improvements such as Directed Random Search, and Adaptive Random Search (Section 2.3).

2.2.2 Strategy

The strategy of Random Search is to sample solutions from across the entire search space using a uniform probability distribution. Each future sample is independent of the samples that come before it.

2.2.3 Procedure

Algorithm 1 provides a pseudo-code listing of the Random Search Algorithm for minimizing a cost function.

Algorithm 1: Pseudo Code Listing for the Random Search Algorithm.

Input: NumIterations, ProblemSize, SearchSpace

Output: Best

```

1 Best  $\leftarrow$  0;
2 foreach  $iter_i \in$  NumIterations do
3    $candidate_i = \text{RandomSolution}(\text{ProblemSize}, \text{SearchSpace});$ 
4   if  $\text{Cost}(candidate_i) < \text{Cost}(\text{Best})$  then
5     Best  $\leftarrow candidate_i$ ;
6   end
7 end
8 return Best;
```

2.2.4 Heuristics

- Random search is minimal in that it only requires a candidate solution construction routine and a candidate solution evaluation routine, both of which may be calibrated using the approach.
- The worst case performance for Random Search for locating the optima is worse than an Enumeration of the search domain, given that Random Search has no memory and can blindly resample.

- Random Search can return a reasonable approximation of the optimal solution within a reasonable time under low problem dimensionality, although the approach does not scale well with problem size (such as the number of dimensions).
- Care must be taken with some problem domains to ensure that random candidate solution construction is unbiased
- The results of a Random Search can be used to seed another search technique, like a local search technique (such as the Hill Climbing algorithm) that can be used to locate the best solution in the neighborhood of the ‘good’ candidate solution.

2.2.5 Code Listing

Listing 2.1 provides an example of the Random Search Algorithm implemented in the Ruby Programming Language. In the example, the algorithm runs for a fixed number of iterations and returns the best candidate solution discovered. The example problem is an instance of a continuous function optimization that seeks $\min f(x)$ where $f = \sum_{i=1}^n x_i^2$, $-5.0 \leq x_i \leq 5.0$ and $n = 2$. The optimal solution for this basin function is $(v_0, \dots, v_{n-1}) = 0.0$.

```

1 def cost(candidate_vector)
2   return candidate_vector.inject(0) {|sum, x| sum + (x ** 2.0)}
3 end
4
5 def random_solution(problem_size, search_space)
6   return Array.new(problem_size) do |i|
7     search_space[i][0] + ((search_space[i][1] - search_space[i][0]) * rand())
8   end
9 end
10
11 def search(max_iterations, problem_size, search_space)
12   best = nil
13   max_iterations.times do |iter|
14     candidate = {}
15     candidate[:vector] = random_solution(problem_size, search_space)
16     candidate[:cost] = cost(candidate[:vector])
17     best = candidate if best.nil? or candidate[:cost] < best[:cost]
18     puts " > iteration #{(iter+1)}, best=#{best[:cost]}"
19   end
20   return best
21 end
22
23 max_iterations = 100
24 problem_size = 2
25 search_space = Array.new(problem_size) {|i| [-5, +5]}
26
27 best = search(max_iterations, problem_size, search_space)
28 puts "Done. Best Solution: c=#{best[:cost]}, v=#{best[:vector].inspect}"

```

Listing 2.1: Random Search Algorithm in the Ruby Programming Language

2.2.6 References

Primary Sources

There is no seminal specification of the Random Search algorithm, rather there are discussions of the general approach and related random search methods from the 1950's through to the 1970's. This was around the time that pattern and direct search methods were actively researched. Brooks is credited with the so-called 'pure random search' [26]. Two seminal reviews of 'random search methods' of the time include: Karnopp [79] and perhaps Kul'chitskii [84].

Learn More

For overviews of into Random Search Methods see Zhigljavsky [147], Solis and Wets [122], and also White [142] who provides an excellent review article. Spall provides a detailed overview of the field of Stochastic Optimization, including the Random Search method [123] (for example, see Chapter 2). For a shorter introduction by Spall, see [124] (specifically Section 6.2). Also see Zabinsky for another detailed review of the broader field [145].

2.3 Adaptive Random Search

Adaptive Random Search, ARS, Adaptive Step Size Random Search, ASSRS, Variable Step-Size Random Search.

2.3.1 Taxonomy

The Adaptive Random Search algorithm belongs to the general set of approaches known as Stochastic Optimization and Global Optimization. It is a direct search method in that it does not require derivatives to navigate the search space. Adaptive Random Search is an extension of the Random Search (Section 2.2) and Localized Random Search algorithms.

2.3.2 Strategy

The Adaptive Random Search algorithm was designed to address the limitations of the fixed step size in the Localized Random Search algorithm. The strategy for Adaptive Random Search is to continually approximate the optimal step size required to reach the global optimum in the search space. This is achieved by trialling and adopting smaller or larger step sizes only if they result in an improvement in the search performance.

The Strategy of the Adaptive Step Size Random Search algorithm (the specific technique reviewed) is to trial a larger step in each iteration and adopt the larger step if it results in an improved result. Very large step sizes are trialled in the same manner although with a much lower frequency. This strategy of preferring large moves is intended to allow the technique to escape local optimal. Smaller step sizes are adopted if no improvement is made for an extended period.

2.3.3 Procedure

Algorithm 2 provides a pseudo-code listing of the Adaptive Random Search Algorithm for minimizing a cost function based on the specification for ‘Adaptive Step-Size Random Search’ by Schummer and Steiglitz [118].

2.3.4 Heuristics

- Adaptive Random Search was designed for continuous function optimization problem domains.
- Candidates with equal cost should be considered improvements to allow the algorithm to make progress across plateaus in the response surface.
- Adaptive Random Search may adapt the search direction in addition to the step size.
- The step size may be adapted for all parameters, or for each parameter individually.

Algorithm 2: Pseudo Code Listing for the Adaptive Random Search Algorithm.

Input: $Iter_{max}$, ProblemSize, SearchSpace, $StepSize_{init}$, $StepSizeF_{small}$,
 $StepSizeF_{large}$, $StepSizeIter_{large}$, $NoChng_{max}$

Output: Current

```

1   $nochng_{count} \leftarrow 0$ ;
2   $step_{size} \leftarrow \text{InitializeStepSize}(\text{SearchSpace}, StepSize_{init})$ ;
3  Current  $\leftarrow \text{RandomSolution}(\text{ProblemSize}, \text{SearchSpace})$ ;
4  foreach  $iter_i \in Iter_{max}$  do
5       $candidate_1 \leftarrow \text{TakeStep}(\text{SearchSpace}, \text{Current}, step_{size})$ ;
6       $largestep_{size} \leftarrow 0$ ;
7      if  $iter_i \bmod StepSizeIter_{large}$  then
8           $largestep_{size} \leftarrow step_{size} \times StepSizeF_{large}$ ;
9      end
10     else
11          $largestep_{size} \leftarrow step_{size} \times StepSizeF_{small}$ ;
12     end
13      $candidate_2 \leftarrow \text{TakeStep}(\text{SearchSpace}, \text{Current}, largestep_{size})$ ;
14     if  $\text{Cost}(candidate_1) \leq \text{Cost}(\text{Current})$  or  $\text{Cost}(candidate_2) \leq \text{Cost}(\text{Current})$ 
15     then
16         if  $\text{Cost}(candidate_2) < \text{Cost}(candidate_1)$  then
17             Current  $\leftarrow candidate_2$ ;
18              $step_{size} \leftarrow largestep_{size}$ ;
19         end
20         else
21             Current  $\leftarrow candidate_1$ ;
22         end
23          $nochng_{count} \leftarrow 0$ ;
24     end
25     else
26          $nochng_{count} \leftarrow nochng_{count} + 1$ ;
27         if  $nochng_{count} > NoChng_{max}$  then
28              $nochng_{count} \leftarrow 0$ ;
29              $step_{size} \leftarrow \frac{step_{size}}{StepSizeF_{small}}$ ;
30         end
31     end
32 return Current;

```

2.3.5 Code Listing

Listing 2.2 provides an example of the Adaptive Random Search Algorithm implemented in the Ruby Programming Language, based on the specification for ‘Adaptive Step-Size Random Search’ by Schummer and Steiglitz [118]. In the example, the algorithm runs for a fixed number of iterations and returns the best candidate solution discovered. The example problem is an instance of a continuous function optimization that seeks $\min f(x)$ where $f = \sum_{i=1}^n x_i^2$, $-5.0 < x_i < 5.0$ and $n = 2$. The optimal solution for this basin function is $(v_0, \dots, v_{n-1}) = 0.0$.

```

1 def cost(candidate_vector)
2   return candidate_vector.inject(0) {|sum, x| sum + (x ** 2.0)}
3 end
4
5 def random_solution(problem_size, search_space)
6   return Array.new(problem_size) do |i|
7     search_space[i][0] + ((search_space[i][1] - search_space[i][0]) * rand())
8   end
9 end
10
11 def take_step(problem_size, search_space, current, step_size)
12   step = []
13   problem_size.times do |i|
14     max, min = current[i]+step_size, current[i]-step_size
15     max = search_space[i][1] if max > search_space[i][1]
16     min = search_space[i][0] if min < search_space[i][0]
17     step << min + ((max - min) * rand)
18   end
19   return step
20 end
21
22 def large_step_size(iteration, step_size, small_factor, large_factor, factor_multiple)
23   if iteration.modulo(factor_multiple)
24     return step_size * large_factor
25   end
26   return step_size * small_factor
27 end
28
29 def search(max_iterations, problem_size, search_space, init_factor, small_factor,
30           large_factor, factor_multiple, max_no_improvements)
31   step_size = (search_space[0][1]-search_space[0][0]) * init_factor
32   current, count = {}, 0
33   current[:vector] = random_solution(problem_size, search_space)
34   current[:cost] = cost(current[:vector])
35   max_iterations.times do |iter|
36     step, bigger_step = {}, {}
37     step[:vector] = take_step(problem_size, search_space, current[:vector], step_size)
38     step[:cost] = cost(step[:vector])
39     bigger_step_size = large_step_size(iter, step_size, small_factor, large_factor,
40                                       factor_multiple)
41     bigger_step[:vector] = take_step(problem_size, search_space, current[:vector],
42                                     bigger_step_size)
43     bigger_step[:cost] = cost(bigger_step[:vector])

```

```

41     if step[:cost] <= current[:cost] or bigger_step[:cost] <= current[:cost]
42         if bigger_step[:cost] < step[:cost]
43             step_size, current = bigger_step_size, bigger_step
44         else
45             current = step
46         end
47         count = 0
48     else
49         count += 1
50         count, stepSize = 0, (step_size/small_factor) if count >= max_no_improvements
51     end
52     puts " > iteration #{(iter+1)}, best=#{current[:cost]}"
53 end
54 return current
55 end
56
57 max_iterations = 1000
58 problem_size = 2
59 search_space = Array.new(problem_size) {|i| [-5, +5]}
60 init_factor = 0.05
61 small_factor = 1.3
62 large_factor = 3.0
63 factor_multiple = 10
64 max_no_improvements = 30
65
66 best = search(max_iterations, problem_size, search_space, init_factor, small_factor,
67               large_factor, factor_multiple, max_no_improvements)
68 puts "Done. Best Solution: cost=#{best[:cost]}, v=#{best[:vector].inspect}"

```

Listing 2.2: Adaptive Random Search Algorithm in the Ruby Programming Language

2.3.6 References

Primary Sources

Many works in the 1960s and 1970s experimented with variable step sizes for Random Search methods. Schummer and Steiglitz are commonly credited the adaptive step size procedure, which they called ‘Adaptive Step-Size Random Search’ [118]. Their approach only modifies the step size based on an approximation of the optimal step size required to reach the global optima. Kregting and White review adaptive random search methods and propose an approach called ‘Adaptive Directional Random Search’ that modifies both the algorithms step size and direction in response to the cost function [82].

Learn More

White reviews extensions to Rastrigin’s ‘Creeping Random Search’ [113] (fixed step size) that use probabilistic step sizes drawn stochastically from uniform and probabilistic distributions [142]. White also reviews works that propose dynamic control strategies for the step size, such as Karnopp [79] who proposes increases and decreases to the step size based on performance over very small numbers of trials. Schrack and Choit review

random search methods that modify their step size in order to approximate optimal moves while searching, including the property of reversal [117]. Masri, et al. describe an adaptive random search strategy that alternates between periods of fixed and variable step sizes [96].

2.4 Stochastic Hill Climbing

Stochastic Hill Climbing, SHC, Random Hill Climbing, RHC, Random Mutation Hill Climbing, RMHC.

2.4.1 Taxonomy

The Stochastic Hill Climbing algorithm is a Stochastic Optimization algorithm and is a Local Optimization algorithm (contrasted to Global Optimization). It is a direct search technique, as it does not require derivatives of the search space. Stochastic Hill Climbing is an extension of deterministic hill climbing algorithms such as Simple Hill Climbing (first-best neighbor), Steepest-Ascent Hill Climbing (best neighbor), and a parent of approaches such as Parallel Hill Climbing and Random-Restart Hill Climbing.

2.4.2 Strategy

The strategy of the Stochastic Hill Climbing algorithm is iterate the process of randomly selecting a neighbor for a candidate solution and only accept it if it results in an improvement. The strategy was proposed to address the limitations of deterministic hill climbing techniques that were likely to get stuck in local optima due to their greedy acceptance of neighboring moves.

2.4.3 Procedure

Algorithm 3 provides a pseudo-code listing of the Stochastic Hill Climbing algorithm for minimizing a cost function, specifically the Random Mutation Hill Climbing algorithm described by Forrest and Mitchell applied to a maximization optimization problem [49].

Algorithm 3: Pseudo Code Listing for the Stochastic Hill Climbing algorithm.

Input: $Iter_{max}$, ProblemSize
Output: Current

```

1 Current  $\leftarrow$  RandomSolution(ProblemSize);
2 foreach  $iter_i \in Iter_{max}$  do
3   Candidate  $\leftarrow$  RandomNeighbor(Current);
4   if Cost(Candidate)  $\geq$  Cost(Current) then
5     Current  $\leftarrow$  Candidate;
6   end
7 end
8 return Current;
```

2.4.4 Heuristics

- Stochastic Hill Climbing was designed to be used in discrete domains with explicit neighbors such as combinatorial optimization (compared to continuous function

optimization).

- The algorithm’s strategy may be applied to continuous domains by making use of a step-size to define candidate-solution neighbors (such as Localized Random Search and Fixed Step-Size Random Search).
- Stochastic Hill Climbing is a local search technique (compared to global search) and may be used to refine a result after the execution of a global search algorithm.
- Even though the technique uses a stochastic process, it can still get stuck in local optima.
- Neighbors with better or equal cost should be accepted, allowing the technique to navigate across plateaus in the response surface.
- The algorithm can be restarted and repeated a number of times after it converges to provide an improved result (called Multiple Restart Hill Climbing).
- The procedure can be applied to multiple candidate solutions concurrently, allowing multiple algorithm runs to be performed at the same time (called Parallel Hill Climbing).

2.4.5 Code Listing

Listing 2.3 provides an example of the Stochastic Hill Climbing algorithm implemented in the Ruby Programming Language, specifically the Random Mutation Hill Climbing algorithm described by Forrest and Mitchell [49]. The algorithm is executed for a fixed number of iterations and is applied to a binary string optimization problem called ‘One Max’. The objective of this maximization problem is to prepare a string of all ‘1’ bits, where the cost function only reports the number of bits in a given string.

```

1 def cost(bitstring)
2   return bitstring.inject(0) {|sum,x| sum = sum + ((x=='1') ? 1 : 0)}
3 end
4
5 def random_solution(problem_size)
6   return Array.new(problem_size){|i| (rand<0.5) ? "1" : "0"}
7 end
8
9 def random_neighbor(bitstring)
10  mutant = Array.new(bitstring)
11  pos = rand(bitstring.length)
12  mutant[pos] = (mutant[pos]=='1') ? '0' : '1'
13  return mutant
14 end
15
16 def search(max_iterations, problem_size)
17  candidate = {}
18  candidate[:vector] = random_solution(problem_size)
19  candidate[:cost] = cost(candidate[:vector])

```

```

20 | max_iterations.times do |iter|
21 |   neighbor = {}
22 |   neighbor[:vector] = random_neighbor(candidate[:vector])
23 |   neighbor[:cost] = cost(neighbor[:vector])
24 |   candidate = neighbor if neighbor[:cost] <= candidate[:cost]
25 |   puts " > iteration #{(iter+1)}, best=#{candidate[:cost]}"
26 |   break if candidate[:cost] == problem_size
27 | end
28 | return candidate
29 | end
30 |
31 | max_iterations = 1000
32 | problem_size = 64
33 |
34 | best = search(max_iterations, problem_size)
35 | puts "Done. Best Solution: c=#{best[:cost]}, v=#{best[:vector].join}"

```

Listing 2.3: Stochastic Hill Climbing algorithm in the Ruby Programming Language

2.4.6 References

Primary Sources

Perhaps the most popular implementation of the Stochastic Hill Climbing algorithm is by Forrest and Mitchell, who proposed the Random Mutation Hill Climbing (RMHC) algorithm (with communication from Richard Palmer) in a study that investigated the behavior of the genetic algorithm on a deceptive class of (discrete) bit-string optimization problem called ‘royal road’ functions [49]. The RMHC was compared to two other hill climbing algorithms in addition to the genetic algorithm, specifically: the Steepest-Ascent Hill Climber, and the Next-Ascent Hill Climber. This study was then followed up by Mitchell and Holland [100]

Jules and Wattenberg were also early to consider stochastic hill climbing as an approach to compare to the genetic algorithm [78]. Skalak applied the RMHC algorithm to a single long bit-string that represented a number of prototype vectors for use in classification [120].

Learn More

The Stochastic Hill Climbing algorithm is related to the genetic algorithm without crossover. Simplified version’s of the approach are investigated for bit-string based optimization problems with the population size of the genetic algorithm reduced to one. The general technique has been investigated under the names Iterated Hillclimbing [103], ES(1+1,m,hc) [104], Random Bit Climber [32], and (1+1)-Genetic Algorithm [4]. This main difference between RMHC and ES(1+1) is that the latter uses a fixed probability of a mutation for each discrete element of a solution (meaning the neighborhood size is probabilistic), whereas RMHC will only stochastically modify one element.

2.5 Iterated Local Search

Iterated Local Search, ILS.

2.5.1 Taxonomy

Iterated Local Search is a Metaheuristic and a Global Optimization technique. It is an extension of Mutli Start Search and may be considered a parent of many two-phase search approaches such as Greedy Randomized Adaptive Search Procedure (Section 2.8) and Variable Neighborhood Search (Section 2.7).

2.5.2 Strategy

The objective of Iterated Local Search is to improve upon stochastic Mutli-Restart Search by sampling in the broader neighborhood of candidate solutions and using a Local Search technique to refine solutions to their local optima. Iterated Local Search explores a sequence of solutions created as perturbations of the current best solution, the result of which is refined using an embedded heuristic.

2.5.3 Procedure

Algorithm 4 provides a pseudo-code listing of the Iterated Local Search algorithm for minimizing a cost function.

Algorithm 4: Pseudo Code for the Iterated Local Search algorithm.

Input:

Output: S_{best}

```

1  $S_{best} \leftarrow \text{ConstructInitialSolution}();$ 
2  $S_{best} \leftarrow \text{LocalSearch}();$ 
3  $\text{SearchHistory} \leftarrow S_{best};$ 
4 while  $\neg \text{StopCondition}()$  do
5    $S_{candidate} \leftarrow \text{Perturbation}(S_{best}, \text{SearchHistory});$ 
6    $S_{candidate} \leftarrow \text{LocalSearch}(S_{candidate});$ 
7    $\text{SearchHistory} \leftarrow S_{candidate};$ 
8   if  $\text{AcceptanceCriterion}(S_{best}, S_{candidate}, \text{SearchHistory})$  then
9      $S_{best} \leftarrow S_{candidate};$ 
10  end
11 end
12 return  $S_{best};$ 

```

2.5.4 Heuristics

- Iterated Local Search was designed for and has been predominately applied to discrete domains, such as combinatorial optimization problems.

- The perturbation of the current best solution should be in a neighborhood beyond the reach of the embedded heuristic and should not be easily undone.
- Perturbations that are too small make the algorithm too greedy, perturbations that are too large make the algorithm too stochastic.
- The embedded heuristic is most commonly a problem-specific local search technique.
- The starting point for the search may be a randomly constructed candidate solution, or constructed using a problem-specific heuristic (such as nearest neighbor).
- Perturbations can be made deterministically, although stochastic and probabilistic (adaptive based on history) are the most common.
- The procedure may store as much or as little history as needed to be used during perturbation and acceptance criteria. No history represents a random walk in a larger neighborhood of the best solution and is the most common implementation of the approach.
- The simplest and most common acceptance criteria is an improvement in the cost of constructed candidate solutions.

2.5.5 Code Listing

Listing 2.4 provides an example of the Iterated Local Search algorithm implemented in the Ruby Programming Language. The algorithm is applied to the Berlin52 instance of the Traveling Salesman Problem (TSP), taken from the TSPLIB. The problem seeks a permutation of the order to visit cities (called a tour) that minimized the total distance traveled. The optimal tour distance for Berlin52 instance is 7542 units.

The Iterated Local Search runs for a fixed number of iterations. The implementation is based on a common algorithm configuration for the TSP, where a ‘double-bridge move’ (4-opt) is used as the perturbation technique, and a stochastic 2-opt is used as the embedded Local Search heuristic. The double-bridge move involves partitioning a permutation into 4 pieces (a,b,c,d) and putting it back together in a specific and jumbled ordering (a,d,c,b).

```

1 def euc_2d(c1, c2)
2   Math::sqrt((c1[0] - c2[0])**2.0 + (c1[1] - c2[1])**2.0).round
3 end
4
5 def cost(permutation, cities)
6   distance = 0
7   permutation.each_with_index do |c1, i|
8     c2 = (i==permutation.length-1) ? permutation[0] : permutation[i+1]
9     distance += euc_2d(cities[c1], cities[c2])
10  end
11  return distance
12 end

```

```

13
14 def random_permutation(cities)
15     all = Array.new(cities.length) {|i| i}
16     return Array.new(all.length) {|i| all.delete_at(rand(all.length))}
17 end
18
19 def stochastic_two_opt(permutation)
20     perm = Array.new(permutation)
21     c1, c2 = rand(perm.length), rand(perm.length)
22     c2 = rand(perm.length) while c1 == c2
23     c1, c2 = c2, c1 if c2 < c1
24     perm[c1...c2] = perm[c1...c2].reverse
25     return perm
26 end
27
28 def local_search(best, cities, max_no_improvements)
29     count = 0
30     begin
31         candidate = {}
32         candidate[:vector] = stochastic_two_opt(best[:vector])
33         candidate[:cost] = cost(candidate[:vector], cities)
34         if candidate[:cost] < best[:cost]
35             count, best = 0, candidate
36         else
37             count += 1
38         end
39     end until count >= max_no_improvements
40     return best
41 end
42
43 def double_bridge_move(perm)
44     pos1 = 1 + rand(perm.length / 4)
45     pos2 = pos1 + 1 + rand(perm.length / 4)
46     pos3 = pos2 + 1 + rand(perm.length / 4)
47     return perm[0...pos1] + perm[pos3..perm.length] + perm[pos2...pos3] +
48         perm[pos1...pos2]
49 end
50
51 def perturbation(cities, best)
52     candidate = {}
53     candidate[:vector] = double_bridge_move(best[:vector])
54     candidate[:cost] = cost(candidate[:vector], cities)
55     return candidate
56 end
57
58 def search(cities, max_iterations, max_no_improvements)
59     best = {}
60     best[:vector] = random_permutation(cities)
61     best[:cost] = cost(best[:vector], cities)
62     best = local_search(best, cities, max_no_improvements)
63     max_iterations.times do |iter|
64         candidate = perturbation(cities, best)
65         candidate = local_search(candidate, cities, max_no_improvements)

```

```

65     best = candidate if candidate[:cost] < best[:cost]
66     puts " > iteration #{(iter+1)}, best=#{best[:cost]}"
67   end
68   return best
69 end
70
71 max_iterations = 100
72 max_no_improvements = 50
73 berlin52 = [[565,575],[25,185],[345,750],[945,685],[845,655],[880,660],[25,230],
74             [525,1000],[580,1175],[650,1130],[1605,620],[1220,580],[1465,200],[1530,5],
75             [845,680],[725,370],[145,665],[415,635],[510,875],[560,365],[300,465],
76             [520,585],[480,415],[835,625],[975,580],[1215,245],[1320,315],[1250,400],
77             [660,180],[410,250],[420,555],[575,665],[1150,1160],[700,580],[685,595],
78             [685,610],[770,610],[795,645],[720,635],[760,650],[475,960],[95,260],
79             [875,920],[700,500],[555,815],[830,485],[1170,65],[830,610],[605,625],
80             [595,360],[1340,725],[1740,245]]
81
82 best = search(berlin52, max_iterations, max_no_improvements)
83 puts "Done. Best Solution: c=#{best[:cost]}, v=#{best[:vector].inspect}"

```

Listing 2.4: Iterated Local Search algorithm in the Ruby Programming Language

2.5.6 References

Primary Sources

The definition and framework for Iterated Local Search was described by Stützle in his PhD dissertation [128]. Specifically he proposed constraints on what constitutes an Iterated Local Search algorithm as i) a single chain of candidate solutions, and ii) the method used to improve candidate solutions occurs within a reduced space by a black-box heuristic. Stützle does not take credit for the approach, instead highlighting specific instances of Iterated Local Search from the literature, such as ‘iterated descent’ [18], ‘large-step Markov chains’ [95], ‘iterated Lin-Kernighan’ [76], ‘chained local optimization’ [94], as well as [19] that introduces the principle, and [77] that summarized it (list taken from [112])

Learn More

Two early technical reports by Stützle that present applications of Iterated Local Search include a report on the Quadratic Assignment Problem [126], and another on the permutation flow shop problem [125]. Stützle and Hoos also published an early paper studying Iterated Local Search for the TSP [127]. Lourenco, Martin, and Stützle provide a concise presentation of the technique, related techniques and the framework, much as it is presented in Stützle’s dissertation [89]. The same author’s also present an authoritative summary of the approach and its applications as a book chapter [112].

2.6 Guided Local Search

Guided Local Search, GLS.

2.6.1 Taxonomy

The Guided Local Search algorithm is a Metaheuristic and a Global Optimization algorithm that makes use of an embedded Local Search algorithm. It is an extension to Local Search algorithms such as Hill Climbing (Section 2.4) and is similar in strategy to the Tabu Search algorithm (Section 2.10) and the Iterated Local Search algorithm (Section 2.5).

2.6.2 Strategy

The strategy for the Guided Local Search algorithm is to use penalties to encourage a Local Search technique to escape local optima and discover the global optima. A Local Search algorithm is run until it gets stuck in a local optima. The features from the local optima are evaluated and penalized, the results of which are used in an augmented cost function employed by the Local Search procedure. The Local Search is repeated a number of times using the last local optima discovered and the augmented cost function that guides exploration away from solutions with features present in discovered local optima.

2.6.3 Procedure

Algorithm 5 provides a pseudo-code listing of the Guided Local Search algorithm for minimization. The Local Search algorithm used by the Guided Local Search algorithm uses an augmented cost function in the form $h(s) = g(s) + \lambda \cdot \sum_{i=1}^M f_i$, where $h(s)$ is the augmented cost function, $g(s)$ is the problem cost function, λ is the ‘regularization parameter’ (a coefficient for scaling the penalties), s is a locally optimal solution of M features, and f_i is the i ’th feature in locally optimal solution. The augmented cost function is only used by the local search procedure, the Guided Local Search algorithm uses the problem specific cost function without augmentation.

Penalties are only updated for those features in a locally optimal solution that maximize utility, updated by adding 1 to the penalty for the feature (a counter). The utility for a feature is calculated as $U_{feature} = \frac{C_{feature}}{1+P_{feature}}$, where $U_{feature}$ is the utility for penalizing a feature (maximizing), $C_{feature}$ is the cost of the feature, and $P_{feature}$ is the current penalty for the feature.

2.6.4 Heuristics

- The Guided Local Search procedure is independent of the Local Search procedure embedded within it. A suitable domain-specific search procedure should be identified and employed.

Algorithm 5: Pseudo Code Listing for the Guided Local Search algorithm.

```

Input:  $Iter_{max}, \lambda$ 
Output:  $S_{best}$ 
1  $f_{penalties} \leftarrow 0$ ;
2  $S_{best} \leftarrow \text{RandomSolution}()$ ;
3 foreach  $Iter_i \in Iter_{max}$  do
4    $S_{curr} \leftarrow \text{LocalSearch}(S_{best}, \lambda, f_{penalties})$ ;
5    $f_{utilities} \leftarrow \text{CalculateFeatureUtilities}(S_{curr}, f_{penalties})$ ;
6    $f_{penalties} \leftarrow \text{UpdateFeaturePenalties}(S_{curr}, f_{penalties}, f_{utilities})$ ;
7   if  $\text{Cost}(S_{curr}) \leq \text{Cost}(S_{best})$  then
8      $S_{best} \leftarrow S_{curr}$ ;
9   end
10 end
11 return  $S_{best}$ ;

```

- The Guided Local Search procedure may need to be executed for thousands to hundreds-of-thousands of iterations, each iteration of which assumes a run of a Local Search algorithm to convergence.
- The algorithm was designed for discrete optimization problems where a solution is comprised of independently assessable ‘features’ such as Combinatorial Optimization, although it has been applied to continuous function optimization modeled as binary strings.
- The λ parameter is a scaling factor for feature penalization that must be in the same proportion to the candidate solution costs from the specific problem instance to which the algorithm is being applied. As such, the value for λ must be meaningful when used within the augmented cost function (such as when it is added to a candidate solution cost in minimization and subtracted from a cost in the case of a maximization problem).

2.6.5 Code Listing

Listing 2.5 provides an example of the Guided Local Search algorithm implemented in the Ruby Programming Language. The algorithm is applied to the Berlin52 instance of the Traveling Salesman Problem (TSP), taken from the TSPLIB. The problem seeks a permutation of the order to visit cities (called a tour) that minimized the total distance traveled. The optimal tour distance for Berlin52 instance is 7542 units.

The implementation of the algorithm for the TSP was based on the configuration specified by Voudouris in [133]. A TSP-specific local search algorithm is used called 2-opt that selects two points in a permutation and reconnects the tour, potentially untwisting the tour at the selected points. The stopping condition for 2-opt was configured to be a fixed number of non-improving moves.

The equation for setting λ for TSP instances is $\lambda = \alpha \cdot \frac{\text{cost}(\text{optima})}{N}$, where N is the number of cities, $\text{cost}(\text{optima})$ is the cost of a local optimum found by a local search, and $\alpha \in (0, 1]$ (around 0.3 for TSP and 2-opt). The cost of a local optima was fixed to the approximated value of 15000 for the Berlin52 instance. The utility function for features (edges) in the TSP is $U_{\text{edge}} = \frac{D_{\text{edge}}}{1 + P_{\text{edge}}}$, where U_{edge} is the utility for penalizing an edge (maximizing), D_{edge} is the cost of the edge (distance between cities) and P_{edge} is the current penalty for the edge.

```

1 def euc_2d(c1, c2)
2   Math::sqrt((c1[0] - c2[0])**2.0 + (c1[1] - c2[1])**2.0).round
3 end
4
5 def random_permutation(cities)
6   perm = Array.new(cities.length){|i|i}
7   for i in 0...perm.length
8     r = rand(perm.length-i) + i
9     perm[r], perm[i] = perm[i], perm[r]
10  end
11  return perm
12 end
13
14 def two_opt(permutation)
15   perm = Array.new(permutation)
16   c1, c2 = rand(perm.length), rand(perm.length)
17   c2 = rand(perm.length) while c1 == c2
18   c1, c2 = c2, c1 if c2 < c1
19   perm[c1...c2] = perm[c1...c2].reverse
20   return perm
21 end
22
23 def augmented_cost(permutation, penalties, cities, lambda)
24   distance, augmented = 0, 0
25   permutation.each_with_index do |c1, i|
26     c2 = (i==permutation.length-1) ? permutation[0] : permutation[i+1]
27     c1, c2 = c2, c1 if c2 < c1
28     d = euc_2d(cities[c1], cities[c2])
29     distance += d
30     augmented += d + (lambda * (permutation[c1][c2]))
31   end
32   return distance, augmented
33 end
34
35 def local_search(current, cities, penalties, max_no_improvements, lambda)
36   current[:cost], current[:acost] = augmented_cost(current[:vector], penalties, cities,
37     lambda)
38   count = 0
39   begin
40     perm = {}
41     perm[:vector] = two_opt(current[:vector])
42     perm[:cost], perm[:acost] = augmented_cost(perm[:vector], penalties, cities, lambda)
43     if perm[:acost] < current[:acost]
44       count, current = 0, perm
45     end
46   end
47 end

```

```

44     else
45         count += 1
46     end
47 end until count >= max_no_improvements
48 return current
49 end
50
51 def calculate_feature_utilities(penalties, cities, permutation)
52     utilities = Array.new(permutation.length,0)
53     permutation.each_with_index do |c1, i|
54         c2 = (i==permutation.length-1) ? permutation[0] : permutation[i+1]
55         c1, c2 = c2, c1 if c2 < c1
56         utilities[i] = euc_2d(cities[c1], cities[c2]) / (1.0 + penalties[c1][c2])
57     end
58     return utilities
59 end
60
61 def update_penalties!(penalties, cities, permutation, utilities)
62     max = utilities.max()
63     permutation.each_with_index do |c1, i|
64         c2 = (i==permutation.length-1) ? permutation[0] : permutation[i+1]
65         c1, c2 = c2, c1 if c2 < c1
66         penalties[c1][c2] += 1 if utilities[i] == max
67     end
68     return penalties
69 end
70
71 def search(max_iterations, cities, max_no_improvements, lambda)
72     best, current = nil, {}
73     current[:vector] = random_permutation(cities)
74     penalties = Array.new(cities.length){Array.new(cities.length,0)}
75     max_iterations.times do |iter|
76         current = local_search(current, cities, penalties, max_no_improvements, lambda)
77         utilities = calculate_feature_utilities(penalties, cities, current[:vector])
78         update_penalties!(penalties, cities, current[:vector], utilities)
79         if (best.nil? or current[:cost] < best[:cost])
80             best = current
81         end
82         puts " > iteration #{(iter+1)}, best=#{best[:cost]}"
83     end
84     return best
85 end
86
87 max_iterations = 100
88 max_no_improvements = 15
89 alpha = 0.3
90 local_search_optima = 15000.0
91 berlin52 = [[565,575],[25,185],[345,750],[945,685],[845,655],[880,660],[25,230],
92 [525,1000],[580,1175],[650,1130],[1605,620],[1220,580],[1465,200],[1530,5],
93 [845,680],[725,370],[145,665],[415,635],[510,875],[560,365],[300,465],
94 [520,585],[480,415],[835,625],[975,580],[1215,245],[1320,315],[1250,400],
95 [660,180],[410,250],[420,555],[575,665],[1150,1160],[700,580],[685,595],
96 [685,610],[770,610],[795,645],[720,635],[760,650],[475,960],[95,260],

```

```
97 [875,920],[700,500],[555,815],[830,485],[1170,65],[830,610],[605,625],  
98 [595,360],[1340,725],[1740,245]]  
99  
100 lambda = alpha * (local_search_optima/berlin52.length.to_f)  
101 best = search(max_iterations, berlin52, max_no_improvements, lambda)  
102 puts "Done. Best Solution: c=#{best[:cost]}, v=#{best[:vector].inspect}"
```

Listing 2.5: Guided Local Search algorithm in the Ruby Programming Language

2.6.6 References

Primary Sources

Guided Local Search emerged from an approach called GENET, which is a connectionist approach to constraint satisfaction [140, 131]. Guided Local Search was presented by Voudouris and Tsang in a series of technical reports (that were later published) that described the technique and provided example applications of it to constraint satisfaction [136], combinatorial optimization [139, 138], and function optimization [137]. The seminal work on the technique was Voudouris' PhD dissertation [133].

Learn More

Voudouris and Tsang provide a high-level introduction to the technique [135], and a contemporary summary of the approach in Glover and Kochenberger's 'Handbook of metaheuristics' [134] that includes a review of the technique, application areas, and demonstration applications on a diverse set of problem instances. Mills, et al. elaborated on the approach, devising an 'Extended Guided Local Search' (EGLS) technique that added 'aspiration criteria' and random moves to the procedure [99], work which culminated in Mills' PhD dissertation [98]. Lau and Tsang further extended the approach by integrating it with a Genetic Algorithm, called the 'Guided Genetic Algorithm' (GGA) [87], that also culminated in a PhD dissertation by Lau [86].

2.7 Variable Neighborhood Search

Variable Neighborhood Search, VNS.

2.7.1 Taxonomy

Variable Neighborhood Search is a Metaheuristic and a Global Optimization technique that manages a Local Search technique. It is related to the Iterative Local Search algorithm (Section 2.5).

2.7.2 Strategy

The strategy for the Variable Neighborhood Search involves iterative exploration of larger and larger neighborhoods for a given local optima until an improvement is located after which time the search across expanding neighborhoods is repeated. The strategy is motivated by three principles: i) a local minimum for one neighborhood structure may not be a local minimum for a different neighborhood structure, ii) a global minimum is a local minimum for all possible neighborhood structures, and iii) local minimum are relatively close to global minimum for many problem classes.

2.7.3 Procedure

Algorithm 6 provides a pseudo-code listing of the Variable Neighborhood Search algorithm for minimizing a cost function. The pseudo code shows that the systematic search of expanding neighborhoods for a local optimum is abandoned when a global improvement is achieved (shown with the **Break** jump).

2.7.4 Heuristics

- Approximation methods (such as stochastic hill climbing) are suggested for use as the Local Search procedure for large problem instances in order to reduce the running time.
- Variable Neighborhood Search has been applied to a very wide array of combinatorial optimization problems as well as clustering and continuous function optimization problems.
- The embedded Local Search technique should be specialized to the problem type and instance to which the technique is being applied.
- The Variable Neighborhood Descent (VND) can be embedded in the Variable Neighborhood Search as a the Local Search procedure and has been shown to be most effective.

Algorithm 6: Pseudo Code Listing for the Variable Neighborhood Search algorithm.

Input: Neighborhoods
Output: S_{best}

```

1  $S_{best} \leftarrow \text{RandomSolution}();$ 
2 while  $\neg \text{StopCondition}()$  do
3   foreach  $\text{Neighborhood}_i \in \text{Neighborhoods}$  do
4      $\text{Neighborhood}_{curr} \leftarrow \text{CalculateNeighborhood}(S_{best}, \text{Neighborhood}_i);$ 
5      $S_{candidate} \leftarrow \text{RandomSolutionInNeighborhood}(\text{Neighborhood}_{curr});$ 
6      $S_{candidate} \leftarrow \text{LocalSearch}(S_{candidate});$ 
7     if  $\text{Cost}(S_{candidate}) < \text{Cost}(S_{best})$  then
8        $S_{best} \leftarrow S_{candidate};$ 
9       Break;
10    end
11  end
12 end
13 return  $S_{best};$ 

```

2.7.5 Code Listing

Listing 2.6 provides an example of the Variable Neighborhood Search algorithm implemented in the Ruby Programming Language. The algorithm is applied to the Berlin52 instance of the Traveling Salesman Problem (TSP), taken from the TSPLIB. The problem seeks a permutation of the order to visit cities (called a tour) that minimized the total distance traveled. The optimal tour distance for Berlin52 instance is 7542 units.

The Variable Neighborhood Search uses a stochastic 2-opt procedure as the embedded local search. The procedure deletes two edges and reverses the sequence in-between the deleted edges, potentially removing ‘twists’ in the tour. The neighborhood structure used in the search is the number of times the 2-opt procedure is performed on a permutation, between 1 and 20 times. The stopping condition for the local search procedure is a maximum number of iterations without improvement. The same stop condition is employed by the higher-order Variable Neighborhood Search procedure, although with a lower boundary on the number of non-improving iterations.

```

1 def euc_2d(c1, c2)
2   Math::sqrt((c1[0] - c2[0])**2.0 + (c1[1] - c2[1])**2.0).round
3 end
4
5 def random_permutation(cities)
6   perm = Array.new(cities.length){|i|i}
7   for i in 0...perm.length
8     r = rand(perm.length-i) + i
9     perm[r], perm[i] = perm[i], perm[r]
10  end
11  return perm

```

```

12 end
13
14 def stochastic_two_opt!(perm)
15   c1, c2 = rand(perm.length), rand(perm.length)
16   c2 = rand(perm.length) while c1 == c2
17   c1, c2 = c2, c1 if c2 < c1
18   perm[c1...c2] = perm[c1...c2].reverse
19   return perm
20 end
21
22 def cost(permutation, cities)
23   distance = 0
24   permutation.each_with_index do |c1, i|
25     c2 = (i==permutation.length-1) ? permutation[0] : permutation[i+1]
26     distance += euc_2d(cities[c1], cities[c2])
27   end
28   return distance
29 end
30
31 def local_search(best, cities, max_no_improvements, neighborhood)
32   count = 0
33   begin
34     candidate = {}
35     candidate[:vector] = Array.new(best[:vector])
36     neighborhood.times{stochastic_two_opt!(candidate[:vector])}
37     candidate[:cost] = cost(candidate[:vector], cities)
38     if candidate[:cost] < best[:cost]
39       count, best = 0, candidate
40     else
41       count += 1
42     end
43   end until count >= max_no_improvements
44   return best
45 end
46
47 def search(cities, neighborhoods, max_no_improvements, max_no_improvements_ls)
48   best = {}
49   best[:vector] = random_permutation(cities)
50   best[:cost] = cost(best[:vector], cities)
51   iter, count = 0, 0
52   begin
53     neighborhoods.each do |neighborhood|
54       candidate = {}
55       candidate[:vector] = Array.new(best[:vector])
56       neighborhood.times{stochastic_two_opt!(candidate[:vector])}
57       candidate[:cost] = cost(candidate[:vector], cities)
58       candidate = local_search(candidate, cities, max_no_improvements_ls, neighborhood)
59       puts " > iteration #{(iter+1)}, neighborhood=#{neighborhood}, best=#{best[:cost]}"
60       iter += 1
61       if(candidate[:cost] < best[:cost])
62         best = candidate
63         count = 0
64         puts "New best, restarting neighborhood search."

```

```

65     break
66   else
67     count += 1
68   end
69 end
70 end until count >= max_no_improvements
71 return best
72 end
73
74 max_no_mprovements = 50
75 local_search_no_improvements = 70
76 neighborhoods = 1...20
77 berlin52 = [[565,575],[25,185],[345,750],[945,685],[845,655],[880,660],[25,230],
78 [525,1000],[580,1175],[650,1130],[1605,620],[1220,580],[1465,200],[1530,5],
79 [845,680],[725,370],[145,665],[415,635],[510,875],[560,365],[300,465],
80 [520,585],[480,415],[835,625],[975,580],[1215,245],[1320,315],[1250,400],
81 [660,180],[410,250],[420,555],[575,665],[1150,1160],[700,580],[685,595],
82 [685,610],[770,610],[795,645],[720,635],[760,650],[475,960],[95,260],
83 [875,920],[700,500],[555,815],[830,485],[1170,65],[830,610],[605,625],
84 [595,360],[1340,725],[1740,245]]
85
86 best = search(berlin52, neighborhoods, max_no_mprovements, local_search_no_improvements)
87 puts "Done. Best Solution: c=#{best[:cost]}, v=#{best[:vector].inspect}"

```

Listing 2.6: Variable Neighborhood Search algorithm in the Ruby Programming Language

2.7.6 References

2.7.7 Primary Sources

The seminal paper for describing Variable Neighborhood Search was by Mladenovic and Hansen in 1997 [102], although an early abstract by Mladenovic is sometimes cited [101]. The approach is explained in terms of three different variations on the general theme. Variable Neighborhood Descent (VND) refers to the use of a Local Search procedure and the deterministic (as opposed to stochastic or probabilistic) change of neighborhood size. Reduced Variable Neighborhood Search (RVNS) involves performing a stochastic random search within a neighborhood and no refinement via a local search technique. Basic Variable Neighborhood Search is the canonical approach described by Mladenovic and Hansen in the seminal paper.

2.7.8 Learn More

There are a large number of papers published on Variable Neighborhood Search, its applications and variations. Hansen and Mladenovic provide an overview of the approach that includes its recent history, extensions and a detailed review of the numerous areas of application [70]. For some additional useful overviews of the technique, its principles, and applications, see [68, 71, 69].

There are many extensions to Variable Neighborhood Search. Some popular examples include: Variable Neighborhood Decomposition Search (VNDS) that involves embedding a second heuristic or metaheuristic approach in VNS to replace the Local Search procedure [72], Skewed Variable Neighborhood Search (SVNS) that encourages exploration of neighborhoods far away from discovered local optima, and Parallel Variable Neighborhood Search (PVNS) that either parallelizes the local search of a neighborhood or parallelizes the searching of the neighborhoods themselves.

2.8 Greedy Randomized Adaptive Search

Greedy Randomized Adaptive Search Procedure, GRASP.

2.8.1 Taxonomy

The Greedy Randomized Adaptive Search Procedure is a Metaheuristic and Global Optimization algorithm, originally proposed for the Operations Research practitioners. The iterative application of an embedded Local Search technique relate the approach to Iterative Local Search (Section 2.5) and Multi-Start techniques.

2.8.2 Strategy

The objective of the Greedy Randomized Adaptive Search Procedure is to repeatedly sample stochastically greedy solutions, and then use a local search procedure to refine them to a local optima. The strategy of the procedure is centered on the stochastic and greedy step-wise construction mechanism that constrains the selection and order-of-inclusion of the components of a solution based on the value they are expected to provide.

2.8.3 Procedure

Algorithm 7 provides a pseudo-code listing of the Greedy Randomized Adaptive Search Procedure for minimizing a cost function.

Algorithm 7: Pseudo Code for the Greedy Randomized Adaptive Search Procedure.

Input: α
Output: S_{best}

```

1  $S_{best} \leftarrow \text{ConstructRandomSolution}();$ 
2 while  $\neg \text{StopCondition}()$  do
3    $S_{candidate} \leftarrow \text{GreedyRandomizedConstruction}(\alpha);$ 
4    $S_{candidate} \leftarrow \text{LocalSearch}(S_{candidate});$ 
5   if  $\text{Cost}(S_{candidate}) < \text{Cost}(S_{best})$  then
6      $S_{best} \leftarrow S_{candidate};$ 
7   end
8 end
9 return  $S_{best};$ 

```

Algorithm 8 provides the pseudo-code the Greedy Randomized Construction function. The function involves the step-wise construction of a candidate solution using a stochastically greedy construction process. The functions works by building a Restricted Candidate List (RCL) that constraints the components of a solution (features) that may be selected from each cycle. The RCL may be constrained by an explicit size, or by

using a threshold ($\alpha \in [0, 1]$) on the cost of adding each feature to the current candidate solution.

Algorithm 8: Pseudo Code Listing for the Greedy Randomized Construction function.

```

Input:  $\alpha$ 
Output:  $S_{candidate}$ 
1  $S_{candidate} \leftarrow 0$ ;
2 while  $S_{candidate} \neq \text{ProblemSize}$  do
3    $Feature_{costs} \leftarrow 0$ ;
4   for  $Feature_i \notin S_{candidate}$  do
5      $Feature_{costs} \leftarrow \text{CostOfAddingFeatureToSolution}(S_{candidate}, Feature_i)$ ;
6   end
7    $RCL \leftarrow 0$ ;
8    $Fcost_{min} \leftarrow \text{MinCost}(Feature_{costs})$ ;
9    $Fcost_{max} \leftarrow \text{MaxCost}(Feature_{costs})$ ;
10  for  $F_i cost \in Feature_{costs}$  do
11    if  $F_i cost \leq Fcost_{min} + \alpha \cdot (Fcost_{max} - Fcost_{min})$  then
12       $RCL \leftarrow Feature_i$ ;
13    end
14  end
15   $S_{candidate} \leftarrow \text{SelectRandomFeature}(RCL)$ ;
16 end
17 return  $S_{candidate}$ ;

```

2.8.4 Heuristics

- The α threshold defines the amount of greediness of the construction mechanism, where values close to 0 may be too greedy, and values close to 1 may be too generalized.
- As an alternative to using the α threshold, the RCL can be constrained to the top $n\%$ of candidate features that may be selected from each construction cycle.
- The technique was designed for discrete problem classes such as combinatorial optimization problems.

2.8.5 Code Listing

Listing 2.7 provides an example of the Greedy Randomized Adaptive Search Procedure implemented in the Ruby Programming Language. The algorithm is applied to the Berlin52 instance of the Traveling Salesman Problem (TSP), taken from the TSPLIB. The problem seeks a permutation of the order to visit cities (called a tour) that minimized the total distance traveled. The optimal tour distance for Berlin52 instance is 7542 units.

The stochastic and greedy step-wise construction of a tour involves evaluating candidate cities by the the cost they contribute as being the next city in the tour. The algorithm uses a stochastic 2-opt procedure for the Local Search with a fixed number of non-improving iterations as the stopping condition.

```

1 def euc_2d(c1, c2)
2   Math::sqrt((c1[0] - c2[0])**2.0 + (c1[1] - c2[1])**2.0).round
3 end
4
5 def cost(permutation, cities)
6   distance = 0
7   permutation.each_with_index do |c1, i|
8     c2 = (i==permutation.length-1) ? permutation[0] : permutation[i+1]
9     distance += euc_2d(cities[c1], cities[c2])
10  end
11  return distance
12 end
13
14 def stochastic_two_opt(permutation)
15  perm = Array.new(permutation)
16  c1, c2 = rand(perm.length), rand(perm.length)
17  c2 = rand(perm.length) while c1 == c2
18  c1, c2 = c2, c1 if c2 < c1
19  perm[c1...c2] = perm[c1...c2].reverse
20  return perm
21 end
22
23 def local_search(best, cities, max_no_improvements)
24  count = 0
25  begin
26    candidate = {}
27    candidate[:vector] = stochastic_two_opt(best[:vector])
28    candidate[:cost] = cost(candidate[:vector], cities)
29    if candidate[:cost] < best[:cost]
30      count, best = 0, candidate
31    else
32      count += 1
33    end
34  end until count >= max_no_improvements
35  return best
36 end
37
38 def construct_randomized_greedy_solution(cities, alpha)
39  candidate = {}
40  candidate[:vector] = [rand(cities.length)]
41  allCities = Array.new(cities.length) {|i| i}
42  while candidate[:vector].length < cities.length
43    candidates = allCities - candidate[:vector]
44    costs = Array.new(candidates.length) {|i| euc_2d(cities[candidate[:vector].last],
45      cities[i])}
46    rcl, max, min = [], costs.max, costs.min
47    costs.each_with_index { |c,i| rcl<<candidates[i] if c <= (min + alpha*(max-min)) }
48    candidate[:vector] << rcl[rand(rcl.length)]

```

```

48   end
49   candidate[:cost] = cost(candidate[:vector], cities)
50   return candidate
51 end
52
53 def search(cities, max_iterations, max_no_improvements, alpha)
54   best = nil
55   max_iterations.times do |iter|
56     candidate = construct_randomized_greedy_solution(cities, alpha);
57     candidate = local_search(candidate, cities, max_no_improvements)
58     best = candidate if best.nil? or candidate[:cost] < best[:cost]
59     puts " > iteration #{(iter+1)}, best=#{best[:cost]}"
60   end
61   return best
62 end
63
64 max_iterations = 50
65 max_no_improvements = 100
66 greediness_factor = 0.3
67 berlin52 = [[565,575],[25,185],[345,750],[945,685],[845,655],[880,660],[25,230],
68   [525,1000],[580,1175],[650,1130],[1605,620],[1220,580],[1465,200],[1530,5],
69   [845,680],[725,370],[145,665],[415,635],[510,875],[560,365],[300,465],
70   [520,585],[480,415],[835,625],[975,580],[1215,245],[1320,315],[1250,400],
71   [660,180],[410,250],[420,555],[575,665],[1150,1160],[700,580],[685,595],
72   [685,610],[770,610],[795,645],[720,635],[760,650],[475,960],[95,260],
73   [875,920],[700,500],[555,815],[830,485],[1170,65],[830,610],[605,625],
74   [595,360],[1340,725],[1740,245]]
75
76 best = search(berlin52, max_iterations, max_no_improvements, greediness_factor)
77 puts "Done. Best Solution: c=#{best[:cost]}, v=#{best[:vector].inspect}"

```

Listing 2.7: Greedy Randomized Adaptive Search Procedure algorithm in the Ruby Programming Language

2.8.6 References

Primary Sources

The seminal paper that introduces the general approach of stochastic and greedy step-wise construction of candidate solutions is by Feo and Resende [40]. The general approach was inspired by greedy heuristics by Hart and Shogan [73]. The seminal review paper that is cited with the preliminary paper is by Feo and Resende [38], and provides a coherent description of the GRASP technique, an example, and review of early applications. An early application was by Feo, Venkatraman and Bard for a machine scheduling problem [42]. Other early applications to scheduling problems include technical reports [39] (later published as [7]) and [41] (also later published as [43]).

Learn More

There are a vast number of review, application, and extension papers for GRASP. Pitsoulis and Resende provide an extensive contemporary overview of the field as a review chapter [110], as does Resende and Ribeiro that includes a clear presentation of the use of the α threshold parameter instead of a fixed size for the RCL [114]. Festa and Resende provide an annotated bibliography as a review chapter that provides some needed insight into large amount of study that has gone into the approach [44]. There are numerous extensions to GRASP, not limited to the popular Reactive GRASP for adapting α [111], the use of long term memory to allow the technique to learn from candidate solutions discovered in previous iterations, and parallel implementations of the procedure such as ‘Parallel GRASP’ [106].

2.9 Scatter Search

Scatter Search, SS.

2.9.1 Taxonomy

Scatter search is a Metaheuristic and a Global Optimization algorithm. It is also sometimes associated with the field of Evolutionary Computation given the use of a population and recombination in the structure of the technique. Scatter Search is a sibling of Tabu Search (Section 2.10), developed by the same author and based on similar origins.

2.9.2 Strategy

The objective of Scatter Search is to maintain a set of diverse and high-quality candidate solutions. The principle of the approach is that useful information about the global optima is stored in a diverse and elite set of solutions (the reference set) and that recombining samples from the set can exploit this information. The strategy involves an iterative process, where a population of diverse and high-quality candidate solutions that are partitioned into subsets and linearly recombined to create weighted centroids of sample-based neighborhoods. The results of recombination are refined using an embedded heuristic and assessed in the context of the reference set as to whether or not they are retained.

2.9.3 Procedure

Algorithm 9 provides a pseudo-code listing of the Scatter Search algorithm for minimizing a cost function. The procedure is based on the abstract form presented by Glover as a template for the general class of technique [58], with influences from an application of the technique to function optimization by Glover [58].

2.9.4 Heuristics

- Scatter search is suitable for both discrete domains such as combinatorial optimization as well as continuous domains such as non-linear programming (continuous function optimization).
- Small set sizes are preferred for the `ReferenceSet`, such as 10 or 20 members.
- Subset sizes can be 2,3,4 or more members that are all recombined to produce viable candidate solutions within the neighborhood of the members of the subset.
- Each subset should comprise at least one member added to the set in the previous algorithm iteration.
- The Local Search procedure should be a problem-specific improvement heuristic.

Algorithm 9: Pseudo Code for the Scatter Search algorithm.

Input: $DiverseSet_{size}$, $ReferenceSet_{size}$
Output: ReferenceSet

```

1 InitialSet  $\leftarrow$  ConstructInitialSolution( $DiverseSet_{size}$ );
2 RefinedSet  $\leftarrow$  0;
3 for  $S_i \in$  InitialSet do
4   | RefinedSet  $\leftarrow$  LocalSearch( $S_i$ );
5 end
6 ReferenceSet  $\leftarrow$  SelectInitialReferenceSet( $ReferenceSet_{size}$ );
7 while  $\neg$  StopCondition() do
8   Subsets  $\leftarrow$  SelectSubset(ReferenceSet);
9   CandidateSet  $\leftarrow$  0;
10  for  $Subset_i \in$  Subsets do
11    | RecombinedCandidates  $\leftarrow$  RecombineMembers( $Subset_i$ );
12    | for  $S_i \in$  RecombinedCandidates do
13      | | CandidateSet  $\leftarrow$  LocalSearch( $S_i$ );
14    | end
15  end
16  ReferenceSet  $\leftarrow$  Select(ReferenceSet, CandidateSet,  $ReferenceSet_{size}$ );
17 end
18 return ReferenceSet;
```

- The selection of members for the **ReferenceSet** at the end of each iteration favors solutions with higher quality and may also promote diversity.
- The **ReferenceSet** may be updated at the end of an iteration, or dynamically as candidates are created (a so-called steady-state population in some evolutionary computation literature).
- A lack of changes to the **ReferenceSet** may be used as a signal to stop the current search, and potentially restart the search with a newly initialized **ReferenceSet**.

2.9.5 Code Listing

Listing 2.8 provides an example of the Scatter Search algorithm implemented in the Ruby Programming Language. The example problem is an instance of a continuous function optimization that seeks $\min f(x)$ where $f = \sum_{i=1}^n x_i^2$, $-5.0 \leq x_i \leq 5.0$ and $n = 3$. The optimal solution for this basin function is $(v_1, \dots, v_n) = 0.0$.

The algorithm is an implementation of Scatter Search as described in an application of the technique to unconstrained non-linear optimization by Glover [61]. The seeds for initial solutions are generated as random vectors, as opposed to stratified samples. The example was further simplified by not including a restart strategy, and the exclusion of diversity maintenance in the **ReferenceSet**. A stochastic local search algorithm is used

as the embedded heuristic that uses a stochastic step size in the range of half a percent of the search space.

```

1  def cost(candidate_vector)
2    return candidate_vector.inject(0) {|sum, x| sum + (x ** 2.0)}
3  end
4
5  def random_solution(problem_size, search_space)
6    return Array.new(problem_size) do |i|
7      search_space[i][0] + ((search_space[i][1] - search_space[i][0]) * rand())
8    end
9  end
10
11 def take_step(current, search_space, step_size)
12   step = []
13   current.length.times do |i|
14     max, min = current[i]+step_size, current[i]-step_size
15     max = search_space[i][1] if max > search_space[i][1]
16     min = search_space[i][0] if min < search_space[i][0]
17     step << min + ((max - min) * rand)
18   end
19   return step
20 end
21
22 def local_search(best, search_space, max_no_improvements, step_size)
23   count = 0
24   begin
25     candidate = {}
26     candidate[:vector] = take_step(best[:vector], search_space, step_size)
27     candidate[:cost] = cost(candidate[:vector])
28     if candidate[:cost] < best[:cost]
29       count, best = 0, candidate
30     else
31       count += 1
32     end
33   end until count >= max_no_improvements
34   return best
35 end
36
37 def construct_initial_set(problem_size, search_space, div_set_size,
38   max_no_improvements, step_size)
39   diverse_set = []
40   begin
41     candidate = {}
42     candidate[:vector] = random_solution(problem_size, search_space)
43     candidate[:cost] = cost(candidate[:vector])
44     candidate = local_search(candidate, search_space, max_no_improvements, step_size)
45     diverse_set << candidate if !diverse_set.any? {|x| x[:vector]==candidate[:vector]}
46   end until diverse_set.length == div_set_size
47   return diverse_set
48 end
49
50 def euclidean(v1, v2)
51   sum = 0.0

```



```

51  v1.each_with_index {|v, i| sum += (v**2.0 - v2[i]**2.0) }
52  sum = sum + (0.0-sum) if sum < 0.0
53  return Math.sqrt(sum)
54 end
55
56 def distance(vector1, reference_set)
57   sum = 0.0
58   reference_set.each do |s|
59     sum += euclidean(vector1, s[:vector])
60   end
61   return sum
62 end
63
64 def diversify(diverse_set, numElite, ref_set_size)
65   diverse_set.sort!{|x,y| x[:cost] <=> y[:cost]}
66   reference_set = Array.new(numElite){|i| diverse_set[i]}
67   remainder = diverse_set - reference_set
68   remainder.sort!{|x,y| distance(y[:vector], reference_set) <=> distance(x[:vector],
69     reference_set)}
69   reference_set = reference_set + remainder[0..(ref_set_size-reference_set.length)]
70   return reference_set, reference_set[0]
71 end
72
73 def select_subsets(reference_set)
74   additions = reference_set.select{|c| c[:new]}
75   remainder = reference_set - additions
76   remainder = additions if remainder.empty?
77   subsets = []
78   additions.each{|a| remainder.each{|r| subsets << [a,r] if a!=r}}
79   return subsets
80 end
81
82 def recombine(subset, problem_size, search_space)
83   a, b = subset
84   d = rand(euclidean(a[:vector], b[:vector]))/2.0
85   children = []
86   subset.each do |p|
87     step = (rand<0.5) ? +d : -d
88     child = {}
89     child[:vector] = Array.new(problem_size){|i| p[:vector][i]+step}
90     child[:vector].each_with_index {|m,i| child[:vector][i]=search_space[i][0] if
91       m<search_space[i][0]}
91     child[:vector].each_with_index {|m,i| child[:vector][i]=search_space[i][1] if
92       m>search_space[i][1]}
92     child[:cost] = cost(child[:vector])
93     children << child
94   end
95   return children
96 end
97
98 def search(problem_size, search_space, max_iterations, ref_set_size, div_set_size,
99   max_no_improvements, step_size, max_elite)
100   diverse_set = construct_initial_set(problem_size, search_space, div_set_size,

```

```

max_no_improvements, step_size)
100 reference_set, best = diversify(diverse_set, max_elite, ref_set_size)
101 reference_set.each{|c| c[:new] = true}
102 max_iterations.times do |iter|
103   wasChange = false
104   subsets = select_subsets(reference_set)
105   reference_set.each{|c| c[:new] = false}
106   subsets.each do |subset|
107     candidates = recombine(subset, problem_size, search_space)
108     improved = Array.new(candidates.length) {|i| local_search(candidates[i],
109       search_space, max_no_improvements, step_size)}
110     improved.each do |c|
111       if !reference_set.any? {|x| x[:vector]==c[:vector]}
112         c[:new] = true
113         reference_set.sort!{|x,y| x[:cost] <=> y[:cost]}
114         if c[:cost]<reference_set.last[:cost]
115           reference_set.delete(reference_set.last)
116           reference_set << c
117           wasChange = true
118         end
119       end
120     end
121     reference_set.sort!{|x,y| x[:cost] <=> y[:cost]}
122     best = reference_set[0] if reference_set[0][:cost] < best[:cost]
123     puts " > iteration #{(iter+1)}, best=#{best[:cost]}"
124     break if !wasChange
125   end
126   return best
127 end
128
129 num_iterations = 100
130 problem_size = 3
131 search_space = Array.new(problem_size) {|i| [-5, +5]}
132 step_size = (search_space[0][1]-search_space[0][0])*0.005
133 max_no_improvements = 30
134 ref_set_size = 10
135 diverse_set_size = 20
136 no_elite = 5
137
138 best = search(problem_size, search_space, num_iterations, ref_set_size,
139   diverse_set_size, max_no_improvements, step_size, no_elite)
140 puts "Done. Best Solution: c=#{best[:cost]}, v=#{best[:vector].inspect}"

```

Listing 2.8: Scatter Search algorithm in the Ruby Programming Language

2.9.6 References

2.9.7 Primary Sources

A form of the Scatter Search algorithm was proposed by Glover for integer programming [52], based on Glover's earlier work on surrogate constraints. The approach remained

idle until it was revisited by Glover and combined with Tabu Search [57]. The modern canonical reference of the approach was proposed by Glover who provides a abstract template of the procedure that may be specialized for a given application domain [58].

2.9.8 Learn More

The primary reference for the approach is the book by Laguna and Martí that reviews the principles of the approach in detail and presents tutorials on applications of the approach on standard problems using the C programming language [85]. There are many review articles and chapters on Scatter Search that may be used to supplement an understanding of the approach, such as a detailed review chapter by Glover [59], a review of the fundamentals of the approach and its relationship to an abstraction called ‘path linking’ by Glover, Laguna, and Martí [60], and a modern overview of the technique by Martí, Laguna, and Glover [93].

2.10 Tabu Search

Tabu Search, TS, Taboo Search.

2.10.1 Taxonomy

Tabu Search is a Global Optimization algorithm and a Metaheuristic or Meta-strategy for controlling an embedded heuristic technique. Tabu Search is a parent for a large family of derivative approaches that introduce memory structures in Metaheuristics, such as Reactive Tabu Search (Section 2.11) and Parallel Tabu Search.

2.10.2 Strategy

The objective for the Tabu Search algorithm is to constrain an embedded heuristic from returning to recently visited areas of the search space, referred to as cycling. The strategy of the approach is to maintain a short term memory of the specific changes of recent moves within the search space and preventing future moves from undoing those changes. Additional intermediate-term memory structures may be introduced to bias moves toward promising areas of the search space, as well as longer-term memory structures that promote a general diversity in the search across the search space.

2.10.3 Procedure

Algorithm 10 provides a pseudo-code listing of the Tabu Search algorithm for minimizing a cost function. The listing shows the simple Tabu Search algorithm with short term memory, without intermediate and long term memory management.

2.10.4 Heuristics

- Tabu search was designed to manage an embedded hill climbing heuristic, although may be adapted to manage any neighborhood exploration heuristic.
- Tabu search was designed for, and has predominately been applied to discrete domains such as combinatorial optimization problems.
- Candidates for neighboring moves can be generated deterministically for the entire neighborhood or the neighborhood can be stochastically sampled to a fixed size, trading off efficiency for accuracy.
- Intermediate-term memory structures can be introduced (complementing the short-term memory) to focus the search on promising areas of the search space (intensification), called aspiration criteria.
- Long-term memory structures can be introduced (complementing the short-term memory) to encourage useful exploration of the broader search space, called diversification. Strategies may include generating solutions with rarely used components and biasing generation away from the most commonly used solution components.

Algorithm 10: Pseudo Code for the Tabu Search algorithm.

Input: $TabuList_{size}$
Output: S_{best}

```

1  $S_{best} \leftarrow \text{ConstructInitialSolution}();$ 
2  $TabuList \leftarrow 0;$ 
3 while  $\neg \text{StopCondition}()$  do
4    $CandidateList \leftarrow 0;$ 
5   for  $S_{candidate} \in S_{best_{neighborhood}}$  do
6     if  $\neg \text{ContainsAnyFeatures}(S_{candidate}, TabuList)$  then
7        $CandidateList \leftarrow S_{candidate};$ 
8     end
9   end
10   $S_{candidate} \leftarrow \text{LocateBestCandidate}(CandidateList);$ 
11  if  $\text{Cost}(S_{candidate}) \leq \text{Cost}(S_{best})$  then
12     $S_{best} \leftarrow S_{candidate};$ 
13     $TabuList \leftarrow \text{FeatureDifferences}(S_{candidate}, S_{best});$ 
14    while  $TabuList > TabuList_{size}$  do
15       $\text{DeleteFeature}(TabuList);$ 
16    end
17  end
18 end
19 return  $S_{best};$ 

```

2.10.5 Code Listing

Listing 2.9 provides an example of the Tabu Search algorithm implemented in the Ruby Programming Language. The algorithm is applied to the Berlin52 instance of the Traveling Salesman Problem (TSP), taken from the TSPLIB. The problem seeks a permutation of the order to visit cities (called a tour) that minimized the total distance traveled. The optimal tour distance for Berlin52 instance is 7542 units.

The algorithm is an implementation of the simple Tabu Search with a short term memory structure that executes for a fixed number of iterations. The starting point for the search is prepared using a random permutation that is refined using a stochastic 2-opt Local Search procedure. The stochastic 2-opt procedure is used as the embedded hill climbing heuristic with a fixed sized candidate list. The two edges that are deleted in each 2-opt move are stored on the tabu list. This general approach is similar to that used by Knox in his work on Tabu Search for symmetrical TSP [80] and Fiechter for the Parallel Tabu Search for the TSP [45].

```

1 def euc_2d(c1, c2)
2   Math::sqrt((c1[0] - c2[0])**2.0 + (c1[1] - c2[1])**2.0).round
3 end
4
5 def cost(permutation, cities)

```

```

6   distance = 0
7   permutation.each_with_index do |c1, i|
8     c2 = (i==permutation.length-1) ? permutation[0] : permutation[i+1]
9     distance += euc_2d(cities[c1], cities[c2])
10  end
11  return distance
12 end
13
14 def random_permutation(cities)
15   all = Array.new(cities.length) {|i| i}
16   return Array.new(all.length) {|i| all.delete_at(rand(all.length))}
17 end
18
19 def stochastic_two_opt(permutation)
20   perm = Array.new(permutation)
21   c1, c2 = rand(perm.length), rand(perm.length)
22   c2 = rand(perm.length) while c1 == c2
23   c1, c2 = c2, c1 if c2 < c1
24   perm[c1...c2] = perm[c1...c2].reverse
25   return perm, [[permutation[c1-1], permutation[c1]], [permutation[c2-1],
26     permutation[c2]]]
27 end
28
29 def generate_initial_solution(cities, max_no_improvements)
30   best = {}
31   best[:vector] = random_permutation(cities)
32   best[:cost] = cost(best[:vector], cities)
33   count = 0
34   begin
35     candidate = {}
36     candidate[:vector] = stochastic_two_opt(best[:vector])[0]
37     candidate[:cost] = cost(candidate[:vector], cities)
38     if candidate[:cost] <= best[:cost]
39       count, best = 0, candidate
40     else
41       count += 1
42     end
43   end until count >= max_no_improvements
44   return best
45 end
46
47 def is_tabu?(permutation, tabu_list)
48   permutation.each_with_index do |c1, i|
49     c2 = (i==permutation.length-1) ? permutation[0] : permutation[i+1]
50     tabu_list.each do |forbidden_edge|
51       return true if forbidden_edge == [c1, c2]
52     end
53   end
54   return false
55 end
56
57 def generate_candidate(best, tabu_list, cities)
58   permutation, edges = nil, nil

```

```

58   begin
59     permutation, edges = stochastic_two_opt(best[:vector])
60   end while is_tabu?(permutation, tabu_list)
61   candidate = {}
62   candidate[:vector] = permutation
63   candidate[:cost] = cost(candidate[:vector], cities)
64   return candidate, edges
65 end
66
67 def search(cities, tabu_list_size, candidate_list_size, max_iterations,
68           max_no_improvements)
69   best = generate_initial_solution(cities, max_no_improvements)
70   current = best
71   tabu_list = Array.new(tabu_list_size)
72   max_iterations.times do |iter|
73     candidates = Array.new(candidate_list_size) {|i| generate_candidate(current,
74                               tabu_list, cities)}
75     candidates.sort! {|x,y| x.first[:cost] <=> y.first[:cost]}
76     best_candidate = candidates.first[0]
77     best_candidate_edges = candidates.first[1]
78     if best_candidate[:cost] < current[:cost]
79       current = best_candidate
80       best = best_candidate if best_candidate[:cost] < best[:cost]
81       best_candidate_edges.each {|edge| tabu_list.push(edge)}
82       tabu_list.pop while tabu_list.length > tabu_list_size
83     end
84     puts " > iteration #{(iter+1)}, best=#{best[:cost]}"
85   end
86   return best
87 end
88
89 max_iterations = 100
90 max_no_improvements = 50
91 tabu_list_size = 15
92 max_candidates = 50
93 berlin52 = [[565,575],[25,185],[345,750],[945,685],[845,655],[880,660],[25,230],
94             [525,1000],[580,1175],[650,1130],[1605,620],[1220,580],[1465,200],[1530,5],
95             [845,680],[725,370],[145,665],[415,635],[510,875],[560,365],[300,465],
96             [520,585],[480,415],[835,625],[975,580],[1215,245],[1320,315],[1250,400],
97             [660,180],[410,250],[420,555],[575,665],[1150,1160],[700,580],[685,595],
98             [685,610],[770,610],[795,645],[720,635],[760,650],[475,960],[95,260],
99             [875,920],[700,500],[555,815],[830,485],[1170,65],[830,610],[605,625],
100            [595,360],[1340,725],[1740,245]]
101
102 best = search(berlin52, tabu_list_size, max_candidates, max_iterations,
103              max_no_improvements)
104 puts "Done. Best Solution: c=#{best[:cost]}, v=#{best[:vector].inspect}"

```

Listing 2.9: Tabu Search algorithm in the Ruby Programming Language

2.10.6 References

2.10.7 Primary Sources

Tabu Search was introduced by Glover applied to scheduling employees to duty rosters [62] and a more general overview in the context of the TSP [53], based on his previous work on surrogate constraints on integer programming problems [52]. Glover provided a seminal overview of the algorithm in a two-part journal article, the first part of which introduced the algorithm, and reviewed then-recent applications [54], and the second which focused on advanced topics and open areas of research [55].

2.10.8 Learn More

Glover provides a high-level introduction to Tabu Search in the form of a practical tutorial [56], as does Glover and Taillard in a user guide format [64]. The best source of information for Tabu Search is the book dedicated to the approach by Glover and Laguna that covers the principles of the technique in detail as well as an in-depth review of applications [65]. The approach appeared in Science, that considered a modification for its application to continuous function optimization problems [31]. Finally, Gendreau provides an excellent contemporary review of the algorithm, highlighting best practices and application heuristics collected from across the field of study [51].

2.11 Reactive Tabu Search

Reactive Tabu Search, RTS, R-TABU, Reactive Taboo Search.

2.11.1 Taxonomy

Reactive Tabu Search is a Metaheuristic and a Global Optimization algorithm. It is an extension of Tabu Search (Section 2.10) and the basis for a field of reactive techniques called Reactive Local Search and more broadly the field of Reactive Search Optimization.

2.11.2 Strategy

The objective of Tabu Search is to avoid cycles while applying a local search technique. The Reactive Tabu Search addresses this objective by explicitly monitoring the search and reacting to the occurrence of cycles and their repetition by adapting the tabu tenure (tabu list size). The strategy of the broader field of Reactive Search Optimization is to automate the process by which a practitioner configures a search procedure by monitoring its online behavior and to use machine learning techniques to adapt a techniques configuration.

2.11.3 Procedure

Algorithm 11 provides a pseudo-code listing of the Reactive Tabu Search algorithm for minimizing a cost function. The pseudo code is based on a the version of the Reactive Tabu Search described by Battiti and Tecchiolli in [15] with supplements like the `IsTabu` function from [13]. The procedure has been modified for brevity to exude the diversification procedure (escape move). Algorithm 12 describes the memory based reaction that manipulates the size of the `ProhibitionPeriod` in response to identified cycles in the ongoing search. Algorithm 13 describes the selection of the best move from a list of candidate moves in the neighborhood of a given solution. The function permits prohibited moves in the case where a prohibited move is better than the best know solution and the selected admissible move (called aspiration). Algorithm 14 determines whether a given neighborhood move is tabu based on the current `ProhibitionPeriod`, and is employed by sub-functions of the Algorithm 13 function.

2.11.4 Heuristics

- Reactive Tabu Search is an extension of Tabu Search and as such should exploit the best practices used for the parent algorithm.
- Reactive Tabu Search was designed for discrete domains such as combinatorial optimization, although has been applied to continuous function optimization.
- Reactive Tabu Search was proposed to use efficient memory data structures such as hash tables.

Algorithm 11: Pseudo Code for the Reactive Tabu Search algorithm.

Input: $Iteration_{max}$, Increase, Decrease, ProblemSize
Output: S_{best}

```

1  $S_{curr} \leftarrow \text{ConstructInitialSolution}();$ 
2  $S_{best} \leftarrow S_{curr};$ 
3 TabuList  $\leftarrow 0;$ 
4 ProhibitionPeriod  $\leftarrow 1;$ 
5 foreach  $Iteration_i \in Iteration_{max}$  do
6   MemoryBasedReaction(Increase, Decrease, ProblemSize);
7   CandidateList  $\leftarrow \text{GenerateCandidateNeighborhood}(S_{curr});$ 
8    $S_{curr} \leftarrow \text{BestMove}(\text{CandidateList});$ 
9   TabuList  $\leftarrow S_{curr}_{feature};$ 
10  if  $\text{Cost}(S_{curr}) \leq \text{Cost}(S_{best})$  then
11     $S_{best} \leftarrow S_{curr};$ 
12  end
13 end
14 return  $S_{best};$ 
```

Algorithm 12: Pseudo Code for the MemoryBasedReaction function in the Reactive Tabu Search algorithm.

Input: Increase, Decrease, ProblemSize
Output:

```

1 if HaveVisitedSolutionBefore( $S_{curr}$ , VisitedSolutions) then
2    $S_{curr}_t \leftarrow \text{RetrieveLastTimeVisited}(\text{VisitedSolutions}, S_{curr});$ 
3   RepetitionInterval  $\leftarrow Iteration_i - S_{curr}_t;$ 
4    $S_{curr}_t \leftarrow Iteration_i;$ 
5   if RepetitionInterval  $< 2 * \text{ProblemSize}$  then
6     RepetitionIntervalavg  $\leftarrow$ 
7        $0.1 * \text{RepetitionInterval} + 0.9 * \text{RepetitionInterval}_{avg};$ 
8     ProhibitionPeriod  $\leftarrow \text{ProhibitionPeriod} * \text{Increase};$ 
9     ProhibitionPeriodt  $\leftarrow Iteration_i;$ 
10  end
11 else
12   VisitedSolutions  $\leftarrow S_{curr};$ 
13    $S_{curr}_t \leftarrow Iteration_i;$ 
14 end
15 if  $Iteration_i - \text{ProhibitionPeriod}_t > \text{RepetitionInterval}_{avg}$  then
16   ProhibitionPeriod  $\leftarrow \text{Max}(1, \text{ProhibitionPeriod} * \text{Decrease});$ 
17   ProhibitionPeriodt  $\leftarrow Iteration_i;$ 
18 end
```

Algorithm 13: Pseudo Code for the `BestMove` function in the Reactive Tabu Search algorithm.

Input: `ProblemSize`
Output: S_{curr}

```

1  $CandidateList_{admissible} \leftarrow \text{GetAdmissibleMoves}(CandidateList);$ 
2  $CandidateList_{tabu} \leftarrow CandidateList - CandidateList_{admissible};$ 
3 if  $\text{Size}(CandidateList_{admissible}) < 2$  then
4   |  $ProhibitionPeriod \leftarrow ProblemSize - 2;$ 
5   |  $ProhibitionPeriod_t \leftarrow Iteration_i;$ 
6 end
7  $S_{curr} \leftarrow \text{GetBest}(CandidateList_{admissible});$ 
8  $S_{best_{tabu}} \leftarrow \text{GetBest}(CandidateList_{tabu});$ 
9 if  $\text{Cost}(S_{best_{tabu}}) < \text{Cost}(S_{best}) \wedge \text{Cost}(S_{best_{tabu}}) < \text{Cost}(S_{curr})$  then
10  |  $S_{curr} \leftarrow S_{best_{tabu}};$ 
11 end
12 return  $S_{curr};$ 

```

Algorithm 14: Pseudo Code for the `IsTabu` function in the Reactive Tabu Search algorithm.

Input:
Output: `Tabu`

```

1  $Tabu \leftarrow \text{FALSE};$ 
2  $S_{curr_{feature}}^t \leftarrow \text{RetrieveTimeFeatureLastUsed}(S_{curr_{feature}});$ 
3 if  $S_{curr_{feature}}^t \geq Iteration_{curr} - ProhibitionPeriod$  then
4   |  $Tabu \leftarrow \text{TRUE};$ 
5 end
6 return  $Tabu;$ 

```

- Reactive Tabu Search was proposed to use an long-term memory to diversify the search after a threshold of cycle repetitions has been reached.
- The `increase` parameter should be greater than one (such as 1.1 or 1.3) and the `decrease` parameter should be less than one (such as 0.9 or 0.8).

2.11.5 Code Listing

Listing 2.10 provides an example of the Reactive Tabu Search algorithm implemented in the Ruby Programming Language. The algorithm is applied to the Berlin52 instance of the Traveling Salesman Problem (TSP), taken from the TSPLIB. The problem seeks a permutation of the order to visit cities (called a tour) that minimized the total distance traveled. The optimal tour distance for Berlin52 instance is 7542 units.

The procedure is based on the code listing described by Battiti and Tecchiolli in [15]

with supplements like the `IsTabu` function from [13]. The implementation does not use efficient memory data structures such as hash tables. The algorithm is initialized with a stochastic 2-opt local search, and the neighborhood is generated as a fixed candidate list of stochastic 2-opt moves. The edges selected for changing in the 2-opt move are stored as features in the tabu list. The example does not implement the escape procedure for search diversification.

```

1  def euc_2d(c1, c2)
2    Math::sqrt((c1[0] - c2[0])**2.0 + (c1[1] - c2[1])**2.0).round
3  end
4
5  def cost(permutation, cities)
6    distance = 0
7    permutation.each_with_index do |c1, i|
8      c2 = (i==permutation.length-1) ? permutation[0] : permutation[i+1]
9      distance += euc_2d(cities[c1], cities[c2])
10   end
11   return distance
12 end
13
14 def random_permutation(cities)
15   all = Array.new(cities.length) {|i| i}
16   return Array.new(all.length) {|i| all.delete_at(rand(all.length))}
17 end
18
19 def stochastic_two_opt(permutation)
20   perm = Array.new(permutation)
21   c1, c2 = rand(perm.length), rand(perm.length)
22   c2 = rand(perm.length) while c1 == c2
23   c1, c2 = c2, c1 if c2 < c1
24   perm[c1...c2] = perm[c1...c2].reverse
25   return perm, [[permutation[c1-1], permutation[c1]], [permutation[c2-1],
26     permutation[c2]]]
27 end
28
29 def generate_initial_solution(cities, max_no_improvements)
30   best = {}
31   best[:vector] = random_permutation(cities)
32   best[:cost] = cost(best[:vector], cities)
33   count = 0
34   begin
35     candidate = {}
36     candidate[:vector] = stochastic_two_opt(best[:vector])[0]
37     candidate[:cost] = cost(candidate[:vector], cities)
38     if candidate[:cost] <= best[:cost]
39       count, best = 0, candidate
40     else
41       count += 1
42     end
43   end until count >= max_no_improvements
44   return best
45 end

```

```

46 def is_tabu?(edge, tabu_list, iteration, prohibition_period)
47   tabu_list.each do |entry|
48     if entry[:edge] == edge
49       if entry[:iteration] >= iteration-prohibition_period
50         return true
51       else
52         return false
53       end
54     end
55   end
56   return false
57 end
58
59 def make_tabu(tabu_list, edge, iteration)
60   tabu_list.each do |entry|
61     if entry[:edge] == edge
62       entry[:iteration] = iteration
63       return entry
64     end
65   end
66   entry = {}
67   entry[:edge] = edge
68   entry[:iteration] = iteration
69   tabu_list.push(entry)
70   return entry
71 end
72
73 def to_edge_list(permutation)
74   list = []
75   permutation.each_with_index do |c1, i|
76     c2 = (i==permutation.length-1) ? permutation[0] : permutation[i+1]
77     c1, c2 = c2, c1 if c1 > c2
78     list << [c1, c2]
79   end
80   return list
81 end
82
83 def equivalent_permutations(edgelist1, edgelist2)
84   edgelist1.each do |edge|
85     return false if !edgelist2.include?(edge)
86   end
87   return true
88 end
89
90 def generate_candidate(best, cities)
91   candidate = {}
92   candidate[:vector], edges = stochastic_two_opt(best[:vector])
93   candidate[:cost] = cost(candidate[:vector], cities)
94   return candidate, edges
95 end
96
97 def get_candidate_entry(visited_list, permutation)
98   edgeList = to_edge_list(permutation)

```

```

99   visited_list.each do |entry|
100     return entry if equivalent_permutations(edgeList, entry[:edgelist])
101   end
102   return nil
103 end
104
105 def store_permutation(visited_list, permutation, iteration)
106   entry = {}
107   entry[:edgelist] = to_edge_list(permutation)
108   entry[:iteration] = iteration
109   entry[:visits] = 1
110   visited_list.push(entry)
111   return entry
112 end
113
114 def sort_neighbourhood(candidates, tabu_list, prohibition_period, iteration)
115   tabu, admissable = [], []
116   candidates.each do |a|
117     if is_tabu?(a[1][0], tabu_list, iteration, prohibition_period) or
118        is_tabu?(a[1][1], tabu_list, iteration, prohibition_period)
119       tabu << a
120     else
121       admissable << a
122     end
123   end
124   return tabu, admissable
125 end
126
127 def search(cities, max_no_improvements, max_candidates, max_iterations, increase,
128            decrease)
129   current = generate_initial_solution(cities, max_no_improvements)
130   best = current
131   tabu_list, prohibition_period = [], 1
132   visited_list, avg_length, last_change = [], 1, 0
133   max_iterations.times do |iter|
134     candidate_entry = get_candidate_entry(visited_list, current[:vector])
135     if !candidate_entry.nil?
136       repetition_interval = iter - candidate_entry[:iteration]
137       candidate_entry[:iteration] = iter
138       candidate_entry[:visits] += 1
139       if repetition_interval < 2*(cities.length-1)
140         avg_length = 0.1*(iter-candidate_entry[:iteration]) + 0.9*avg_length
141         prohibition_period = (prohibition_period.to_f * increase)
142         last_change = iter
143       end
144     else
145       store_permutation(visited_list, current[:vector], iter)
146     end
147     if iter-last_change > avg_length
148       prohibition_period = [prohibition_period*decrease,1].max
149       last_change = iter
150     end
151     candidates = Array.new(max_candidates) {|i| generate_candidate(current, cities)}

```

```

151     candidates.sort! {|x,y| x.first[:cost] <=> y.first[:cost]}
152     tabu, admissible = sort_neighbourhood(candidates, tabu_list, prohibition_period,
153         iter)
154     if admissible.length < 2
155         prohibition_period = cities.length-2
156         last_change = iter
157     end
158     current, best_move_edges = (admissible.empty?) ? tabu.first : admissible.first
159     if !tabu.empty? and tabu.first[0][:cost]<best[:cost] and
160         tabu.first[0][:cost]<current[:cost]
161         current, best_move_edges = tabu.first
162     end
163     best_move_edges.each {|edge| make_tabu(tabu_list, edge, iter)}
164     best = candidates.first[0] if candidates.first[0][:cost] < best[:cost]
165     puts " > iteration #{(iter+1)}, tenure=#{prohibition_period.round},
166         best=#{best[:cost]}"
167     best=#{best[:cost]}"
168 end
169 return best
170 end
171
172 max_iterations = 300
173 max_no_improvements = 50
174 max_candidates = 50
175 increase = 1.3
176 decrease = 0.9
177 berlin52 = [[565,575],[25,185],[345,750],[945,685],[845,655],[880,660],[25,230],
178 [525,1000],[580,1175],[650,1130],[1605,620],[1220,580],[1465,200],[1530,5],
179 [845,680],[725,370],[145,665],[415,635],[510,875],[560,365],[300,465],
180 [520,585],[480,415],[835,625],[975,580],[1215,245],[1320,315],[1250,400],
181 [660,180],[410,250],[420,555],[575,665],[1150,1160],[700,580],[685,595],
182 [685,610],[770,610],[795,645],[720,635],[760,650],[475,960],[95,260],
183 [875,920],[700,500],[555,815],[830,485],[1170,65],[830,610],[605,625],
184 [595,360],[1340,725],[1740,245]]
185
186 best = search(berlin52, max_no_improvements, max_candidates, max_iterations, increase,
187     decrease)
188 puts "Done. Best Solution: c=#{best[:cost]}, v=#{best[:vector].inspect}"

```

Listing 2.10: Reactive Tabu Search algorithm in the Ruby Programming Language

2.11.6 References

2.11.7 Primary Sources

Reactive Tabu Search was proposed by Battiti and Tecchiolli as an extension to Tabu Search that included an adaptive tabu list size in addition to a diversification mechanism [13]. The technique also used efficient memory structures that were based on an earlier work by Battiti and Tecchiolli that considered a parallel tabu search [12]. Some early application papers by Battiti and Tecchiolli include a comparison to Simulated Annealing applied to the Quadratic Assignment Problem [14], benchmarked on instances of the knapsack problem and N-K models and compared with Repeated Local Minima Search,

Simulated Annealing, and Genetic Algorithms [15], and training neural networks on an array of problem instances [16].

2.11.8 Learn More

Reactive Tabu Search was abstracted to a form called Reactive Local Search that considers adaptive methods that learn suitable parameters for heuristics that manage an embedded local search technique [10, 11]. Under this abstraction, the Reactive Tabu Search algorithm is single example of the Reactive Local Search principle applied to the Tabu Search. This framework was further extended to the use of any adaptive machine learning techniques to adapt the parameters of an algorithm by reacting to algorithm outcomes online while solving a problem, called Reactive Search [17]. The best reference for this general framework is the book on Reactive Search Optimization by Battiti, Brunato, and Mascia [9]. Additionally, the review chapter by Battiti and Brunato provides a contemporary description [8].

Chapter 3

Physical Algorithms

3.1 Overview

todo

3.2 Simulated Annealing

The heading and alternate headings for the algorithm description.

3.2.1 Taxonomy

A small tree diagram showing related fields and algorithms.

3.2.2 Inspiration

A textual description of the inspiring system.

3.2.3 Metaphor

A textual description of the algorithm by analogy.

3.2.4 Strategy

A textual description of the information processing strategy.

3.2.5 Procedure

A pseudo code description of the algorithms procedure.

3.2.6 Heuristics

A bullet-point listing of best practice usage.

3.2.7 Code Listing

A code listing and a terse description of the listing.

3.2.8 References

An bullet-point annotated reference list of primary sources and useful resources.

3.3 Adaptive Simulated Annealing

The heading and alternate headings for the algorithm description.

3.3.1 Taxonomy

A small tree diagram showing related fields and algorithms.

3.3.2 Inspiration

A textual description of the inspiring system.

3.3.3 Metaphor

A textual description of the algorithm by analogy.

3.3.4 Strategy

A textual description of the information processing strategy.

3.3.5 Procedure

A pseudo code description of the algorithms procedure.

3.3.6 Heuristics

A bullet-point listing of best practice usage.

3.3.7 Code Listing

A code listing and a terse description of the listing.

3.3.8 References

An bullet-point annotated reference list of primary sources and useful resources.

3.4 Memetic Algorithm

The heading and alternate headings for the algorithm description.

3.4.1 Taxonomy

A small tree diagram showing related fields and algorithms.

3.4.2 Inspiration

A textual description of the inspiring system.

3.4.3 Metaphor

A textual description of the algorithm by analogy.

3.4.4 Strategy

A textual description of the information processing strategy.

3.4.5 Procedure

A pseudo code description of the algorithms procedure.

3.4.6 Heuristics

A bullet-point listing of best practice usage.

3.4.7 Code Listing

A code listing and a terse description of the listing.

3.4.8 References

An bullet-point annotated reference list of primary sources and useful resources.

3.5 Extremal Optimization

The heading and alternate headings for the algorithm description.

3.5.1 Taxonomy

A small tree diagram showing related fields and algorithms.

3.5.2 Inspiration

A textual description of the inspiring system.

3.5.3 Metaphor

A textual description of the algorithm by analogy.

3.5.4 Strategy

A textual description of the information processing strategy.

3.5.5 Procedure

A pseudo code description of the algorithms procedure.

3.5.6 Heuristics

A bullet-point listing of best practice usage.

3.5.7 Code Listing

A code listing and a terse description of the listing.

3.5.8 References

An bullet-point annotated reference list of primary sources and useful resources.

3.6 Cultural Algorithm

The heading and alternate headings for the algorithm description.

3.6.1 Taxonomy

A small tree diagram showing related fields and algorithms.

3.6.2 Inspiration

A textual description of the inspiring system.

3.6.3 Metaphor

A textual description of the algorithm by analogy.

3.6.4 Strategy

A textual description of the information processing strategy.

3.6.5 Procedure

A pseudo code description of the algorithms procedure.

3.6.6 Heuristics

A bullet-point listing of best practice usage.

3.6.7 Code Listing

A code listing and a terse description of the listing.

3.6.8 References

An bullet-point annotated reference list of primary sources and useful resources.

3.7 Summary

todo

Chapter 4

Evolutionary Algorithms

4.1 Overview

todo

4.2 Genetic Algorithm

The heading and alternate headings for the algorithm description.

4.2.1 Taxonomy

A small tree diagram showing related fields and algorithms.

4.2.2 Inspiration

A textual description of the inspiring system.

4.2.3 Metaphor

A textual description of the algorithm by analogy.

4.2.4 Strategy

A textual description of the information processing strategy.

4.2.5 Procedure

A pseudo code description of the algorithms procedure.

4.2.6 Heuristics

A bullet-point listing of best practice usage.

4.2.7 Code Listing

A code listing and a terse description of the listing.

4.2.8 References

An bullet-point annotated reference list of primary sources and useful resources.

4.3 Genetic Programming

The heading and alternate headings for the algorithm description.

4.3.1 Taxonomy

A small tree diagram showing related fields and algorithms.

4.3.2 Inspiration

A textual description of the inspiring system.

4.3.3 Metaphor

A textual description of the algorithm by analogy.

4.3.4 Strategy

A textual description of the information processing strategy.

4.3.5 Procedure

A pseudo code description of the algorithms procedure.

4.3.6 Heuristics

A bullet-point listing of best practice usage.

4.3.7 Code Listing

A code listing and a terse description of the listing.

4.3.8 References

An bullet-point annotated reference list of primary sources and useful resources.

4.4 Evolutionary Programming

The heading and alternate headings for the algorithm description.

4.4.1 Taxonomy

A small tree diagram showing related fields and algorithms.

4.4.2 Inspiration

A textual description of the inspiring system.

4.4.3 Metaphor

A textual description of the algorithm by analogy.

4.4.4 Strategy

A textual description of the information processing strategy.

4.4.5 Procedure

A pseudo code description of the algorithms procedure.

4.4.6 Heuristics

A bullet-point listing of best practice usage.

4.4.7 Code Listing

A code listing and a terse description of the listing.

4.4.8 References

An bullet-point annotated reference list of primary sources and useful resources.

4.5 Evolution Strategies

The heading and alternate headings for the algorithm description.

4.5.1 Taxonomy

A small tree diagram showing related fields and algorithms.

4.5.2 Inspiration

A textual description of the inspiring system.

4.5.3 Metaphor

A textual description of the algorithm by analogy.

4.5.4 Strategy

A textual description of the information processing strategy.

4.5.5 Procedure

A pseudo code description of the algorithms procedure.

4.5.6 Heuristics

A bullet-point listing of best practice usage.

4.5.7 Code Listing

A code listing and a terse description of the listing.

4.5.8 References

An bullet-point annotated reference list of primary sources and useful resources.

4.6 Learning Classifier System

The heading and alternate headings for the algorithm description.

4.6.1 Taxonomy

A small tree diagram showing related fields and algorithms.

4.6.2 Inspiration

A textual description of the inspiring system.

4.6.3 Metaphor

A textual description of the algorithm by analogy.

4.6.4 Strategy

A textual description of the information processing strategy.

4.6.5 Procedure

A pseudo code description of the algorithms procedure.

4.6.6 Heuristics

A bullet-point listing of best practice usage.

4.6.7 Code Listing

A code listing and a terse description of the listing.

4.6.8 References

An bullet-point annotated reference list of primary sources and useful resources.

4.7 Differential Evolution

The heading and alternate headings for the algorithm description.

4.7.1 Taxonomy

A small tree diagram showing related fields and algorithms.

4.7.2 Inspiration

A textual description of the inspiring system.

4.7.3 Metaphor

A textual description of the algorithm by analogy.

4.7.4 Strategy

A textual description of the information processing strategy.

4.7.5 Procedure

A pseudo code description of the algorithms procedure.

4.7.6 Heuristics

A bullet-point listing of best practice usage.

4.7.7 Code Listing

A code listing and a terse description of the listing.

4.7.8 References

An bullet-point annotated reference list of primary sources and useful resources.

4.8 Grammatical Evolution

The heading and alternate headings for the algorithm description.

4.8.1 Taxonomy

A small tree diagram showing related fields and algorithms.

4.8.2 Inspiration

A textual description of the inspiring system.

4.8.3 Metaphor

A textual description of the algorithm by analogy.

4.8.4 Strategy

A textual description of the information processing strategy.

4.8.5 Procedure

A pseudo code description of the algorithms procedure.

4.8.6 Heuristics

A bullet-point listing of best practice usage.

4.8.7 Code Listing

A code listing and a terse description of the listing.

4.8.8 References

An bullet-point annotated reference list of primary sources and useful resources.

4.9 Non-dominated Sorting Genetic Algorithm

The heading and alternate headings for the algorithm description.

4.9.1 Taxonomy

A small tree diagram showing related fields and algorithms.

4.9.2 Inspiration

A textual description of the inspiring system.

4.9.3 Metaphor

A textual description of the algorithm by analogy.

4.9.4 Strategy

A textual description of the information processing strategy.

4.9.5 Procedure

A pseudo code description of the algorithms procedure.

4.9.6 Heuristics

A bullet-point listing of best practice usage.

4.9.7 Code Listing

A code listing and a terse description of the listing.

4.9.8 References

An bullet-point annotated reference list of primary sources and useful resources.

4.10 Strength Pareto Evolutionary Algorithm

The heading and alternate headings for the algorithm description.

4.10.1 Taxonomy

A small tree diagram showing related fields and algorithms.

4.10.2 Inspiration

A textual description of the inspiring system.

4.10.3 Metaphor

A textual description of the algorithm by analogy.

4.10.4 Strategy

A textual description of the information processing strategy.

4.10.5 Procedure

A pseudo code description of the algorithms procedure.

4.10.6 Heuristics

A bullet-point listing of best practice usage.

4.10.7 Code Listing

A code listing and a terse description of the listing.

4.10.8 References

An bullet-point annotated reference list of primary sources and useful resources.

4.11 Island Population Genetic Algorithm

The heading and alternate headings for the algorithm description.

4.11.1 Taxonomy

A small tree diagram showing related fields and algorithms.

4.11.2 Inspiration

A textual description of the inspiring system.

4.11.3 Metaphor

A textual description of the algorithm by analogy.

4.11.4 Strategy

A textual description of the information processing strategy.

4.11.5 Procedure

A pseudo code description of the algorithms procedure.

4.11.6 Heuristics

A bullet-point listing of best practice usage.

4.11.7 Code Listing

A code listing and a terse description of the listing.

4.11.8 References

An bullet-point annotated reference list of primary sources and useful resources.

4.12 Summary

todo

Chapter 5

Probabilistic Algorithms

5.1 Overview

todo

5.2 Cross-Entropy Method

The heading and alternate headings for the algorithm description.

5.2.1 Taxonomy

A small tree diagram showing related fields and algorithms.

5.2.2 Inspiration

A textual description of the inspiring system.

5.2.3 Metaphor

A textual description of the algorithm by analogy.

5.2.4 Strategy

A textual description of the information processing strategy.

5.2.5 Procedure

A pseudo code description of the algorithms procedure.

5.2.6 Heuristics

A bullet-point listing of best practice usage.

5.2.7 Code Listing

A code listing and a terse description of the listing.

5.2.8 References

An bullet-point annotated reference list of primary sources and useful resources.

5.3 Population-Based Incremental Learning

The heading and alternate headings for the algorithm description.

5.3.1 Taxonomy

A small tree diagram showing related fields and algorithms.

5.3.2 Inspiration

A textual description of the inspiring system.

5.3.3 Metaphor

A textual description of the algorithm by analogy.

5.3.4 Strategy

A textual description of the information processing strategy.

5.3.5 Procedure

A pseudo code description of the algorithms procedure.

5.3.6 Heuristics

A bullet-point listing of best practice usage.

5.3.7 Code Listing

A code listing and a terse description of the listing.

5.3.8 References

An bullet-point annotated reference list of primary sources and useful resources.

5.4 Probabilistic Incremental Program Evolution

The heading and alternate headings for the algorithm description.

5.4.1 Taxonomy

A small tree diagram showing related fields and algorithms.

5.4.2 Inspiration

A textual description of the inspiring system.

5.4.3 Metaphor

A textual description of the algorithm by analogy.

5.4.4 Strategy

A textual description of the information processing strategy.

5.4.5 Procedure

A pseudo code description of the algorithms procedure.

5.4.6 Heuristics

A bullet-point listing of best practice usage.

5.4.7 Code Listing

A code listing and a terse description of the listing.

5.4.8 References

An bullet-point annotated reference list of primary sources and useful resources.

5.5 Compact Genetic Algorithm

The heading and alternate headings for the algorithm description.

5.5.1 Taxonomy

A small tree diagram showing related fields and algorithms.

5.5.2 Inspiration

A textual description of the inspiring system.

5.5.3 Metaphor

A textual description of the algorithm by analogy.

5.5.4 Strategy

A textual description of the information processing strategy.

5.5.5 Procedure

A pseudo code description of the algorithms procedure.

5.5.6 Heuristics

A bullet-point listing of best practice usage.

5.5.7 Code Listing

A code listing and a terse description of the listing.

5.5.8 References

An bullet-point annotated reference list of primary sources and useful resources.

5.6 Extended Compact Genetic Algorithm

The heading and alternate headings for the algorithm description.

5.6.1 Taxonomy

A small tree diagram showing related fields and algorithms.

5.6.2 Inspiration

A textual description of the inspiring system.

5.6.3 Metaphor

A textual description of the algorithm by analogy.

5.6.4 Strategy

A textual description of the information processing strategy.

5.6.5 Procedure

A pseudo code description of the algorithms procedure.

5.6.6 Heuristics

A bullet-point listing of best practice usage.

5.6.7 Code Listing

A code listing and a terse description of the listing.

5.6.8 References

An bullet-point annotated reference list of primary sources and useful resources.

5.7 Bayesian Optimization Algorithm

The heading and alternate headings for the algorithm description.

5.7.1 Taxonomy

A small tree diagram showing related fields and algorithms.

5.7.2 Inspiration

A textual description of the inspiring system.

5.7.3 Metaphor

A textual description of the algorithm by analogy.

5.7.4 Strategy

A textual description of the information processing strategy.

5.7.5 Procedure

A pseudo code description of the algorithms procedure.

5.7.6 Heuristics

A bullet-point listing of best practice usage.

5.7.7 Code Listing

A code listing and a terse description of the listing.

5.7.8 References

An bullet-point annotated reference list of primary sources and useful resources.

5.8 Hierarchical Bayesian Optimization Algorithm

The heading and alternate headings for the algorithm description.

5.8.1 Taxonomy

A small tree diagram showing related fields and algorithms.

5.8.2 Inspiration

A textual description of the inspiring system.

5.8.3 Metaphor

A textual description of the algorithm by analogy.

5.8.4 Strategy

A textual description of the information processing strategy.

5.8.5 Procedure

A pseudo code description of the algorithms procedure.

5.8.6 Heuristics

A bullet-point listing of best practice usage.

5.8.7 Code Listing

A code listing and a terse description of the listing.

5.8.8 References

An bullet-point annotated reference list of primary sources and useful resources.

5.9 Univariate Marginal Distribution Algorithm

The heading and alternate headings for the algorithm description.

5.9.1 Taxonomy

A small tree diagram showing related fields and algorithms.

5.9.2 Inspiration

A textual description of the inspiring system.

5.9.3 Metaphor

A textual description of the algorithm by analogy.

5.9.4 Strategy

A textual description of the information processing strategy.

5.9.5 Procedure

A pseudo code description of the algorithms procedure.

5.9.6 Heuristics

A bullet-point listing of best practice usage.

5.9.7 Code Listing

A code listing and a terse description of the listing.

5.9.8 References

An bullet-point annotated reference list of primary sources and useful resources.

5.10 Bivariate Marginal Distribution Algorithm

The heading and alternate headings for the algorithm description.

5.10.1 Taxonomy

A small tree diagram showing related fields and algorithms.

5.10.2 Inspiration

A textual description of the inspiring system.

5.10.3 Metaphor

A textual description of the algorithm by analogy.

5.10.4 Strategy

A textual description of the information processing strategy.

5.10.5 Procedure

A pseudo code description of the algorithms procedure.

5.10.6 Heuristics

A bullet-point listing of best practice usage.

5.10.7 Code Listing

A code listing and a terse description of the listing.

5.10.8 References

An bullet-point annotated reference list of primary sources and useful resources.

5.11 Gaussian Adaptation

The heading and alternate headings for the algorithm description.

5.11.1 Taxonomy

A small tree diagram showing related fields and algorithms.

5.11.2 Inspiration

A textual description of the inspiring system.

5.11.3 Metaphor

A textual description of the algorithm by analogy.

5.11.4 Strategy

A textual description of the information processing strategy.

5.11.5 Procedure

A pseudo code description of the algorithms procedure.

5.11.6 Heuristics

A bullet-point listing of best practice usage.

5.11.7 Code Listing

A code listing and a terse description of the listing.

5.11.8 References

An bullet-point annotated reference list of primary sources and useful resources.

5.12 Summary

todo

Chapter 6

Swarm Algorithms

6.1 Overview

todo

6.2 Particle Swarm Optimization

The heading and alternate headings for the algorithm description.

6.2.1 Taxonomy

A small tree diagram showing related fields and algorithms.

6.2.2 Inspiration

A textual description of the inspiring system.

6.2.3 Metaphor

A textual description of the algorithm by analogy.

6.2.4 Strategy

A textual description of the information processing strategy.

6.2.5 Procedure

A pseudo code description of the algorithms procedure.

6.2.6 Heuristics

A bullet-point listing of best practice usage.

6.2.7 Code Listing

A code listing and a terse description of the listing.

6.2.8 References

An bullet-point annotated reference list of primary sources and useful resources.

6.3 AntNet

The heading and alternate headings for the algorithm description.

6.3.1 Taxonomy

A small tree diagram showing related fields and algorithms.

6.3.2 Inspiration

A textual description of the inspiring system.

6.3.3 Metaphor

A textual description of the algorithm by analogy.

6.3.4 Strategy

A textual description of the information processing strategy.

6.3.5 Procedure

A pseudo code description of the algorithms procedure.

6.3.6 Heuristics

A bullet-point listing of best practice usage.

6.3.7 Code Listing

A code listing and a terse description of the listing.

6.3.8 References

An bullet-point annotated reference list of primary sources and useful resources.

6.4 Ant System

The heading and alternate headings for the algorithm description.

6.4.1 Taxonomy

A small tree diagram showing related fields and algorithms.

6.4.2 Inspiration

A textual description of the inspiring system.

6.4.3 Metaphor

A textual description of the algorithm by analogy.

6.4.4 Strategy

A textual description of the information processing strategy.

6.4.5 Procedure

A pseudo code description of the algorithms procedure.

6.4.6 Heuristics

A bullet-point listing of best practice usage.

6.4.7 Code Listing

A code listing and a terse description of the listing.

6.4.8 References

An bullet-point annotated reference list of primary sources and useful resources.

6.5 MAX-MIN Ant System

The heading and alternate headings for the algorithm description.

6.5.1 Taxonomy

A small tree diagram showing related fields and algorithms.

6.5.2 Inspiration

A textual description of the inspiring system.

6.5.3 Metaphor

A textual description of the algorithm by analogy.

6.5.4 Strategy

A textual description of the information processing strategy.

6.5.5 Procedure

A pseudo code description of the algorithms procedure.

6.5.6 Heuristics

A bullet-point listing of best practice usage.

6.5.7 Code Listing

A code listing and a terse description of the listing.

6.5.8 References

An bullet-point annotated reference list of primary sources and useful resources.

6.6 Rank-Based Ant System

The heading and alternate headings for the algorithm description.

6.6.1 Taxonomy

A small tree diagram showing related fields and algorithms.

6.6.2 Inspiration

A textual description of the inspiring system.

6.6.3 Metaphor

A textual description of the algorithm by analogy.

6.6.4 Strategy

A textual description of the information processing strategy.

6.6.5 Procedure

A pseudo code description of the algorithms procedure.

6.6.6 Heuristics

A bullet-point listing of best practice usage.

6.6.7 Code Listing

A code listing and a terse description of the listing.

6.6.8 References

An bullet-point annotated reference list of primary sources and useful resources.

6.7 Ant Colony System

The heading and alternate headings for the algorithm description.

6.7.1 Taxonomy

A small tree diagram showing related fields and algorithms.

6.7.2 Inspiration

A textual description of the inspiring system.

6.7.3 Metaphor

A textual description of the algorithm by analogy.

6.7.4 Strategy

A textual description of the information processing strategy.

6.7.5 Procedure

A pseudo code description of the algorithms procedure.

6.7.6 Heuristics

A bullet-point listing of best practice usage.

6.7.7 Code Listing

A code listing and a terse description of the listing.

6.7.8 References

An bullet-point annotated reference list of primary sources and useful resources.

6.8 Multiple Ant Colony System

The heading and alternate headings for the algorithm description.

6.8.1 Taxonomy

A small tree diagram showing related fields and algorithms.

6.8.2 Inspiration

A textual description of the inspiring system.

6.8.3 Metaphor

A textual description of the algorithm by analogy.

6.8.4 Strategy

A textual description of the information processing strategy.

6.8.5 Procedure

A pseudo code description of the algorithms procedure.

6.8.6 Heuristics

A bullet-point listing of best practice usage.

6.8.7 Code Listing

A code listing and a terse description of the listing.

6.8.8 References

An bullet-point annotated reference list of primary sources and useful resources.

6.9 Population-based Ant Colony Optimization

The heading and alternate headings for the algorithm description.

6.9.1 Taxonomy

A small tree diagram showing related fields and algorithms.

6.9.2 Inspiration

A textual description of the inspiring system.

6.9.3 Metaphor

A textual description of the algorithm by analogy.

6.9.4 Strategy

A textual description of the information processing strategy.

6.9.5 Procedure

A pseudo code description of the algorithms procedure.

6.9.6 Heuristics

A bullet-point listing of best practice usage.

6.9.7 Code Listing

A code listing and a terse description of the listing.

6.9.8 References

An bullet-point annotated reference list of primary sources and useful resources.

6.10 Bees Algorithm

The heading and alternate headings for the algorithm description.

6.10.1 Taxonomy

A small tree diagram showing related fields and algorithms.

6.10.2 Inspiration

A textual description of the inspiring system.

6.10.3 Metaphor

A textual description of the algorithm by analogy.

6.10.4 Strategy

A textual description of the information processing strategy.

6.10.5 Procedure

A pseudo code description of the algorithms procedure.

6.10.6 Heuristics

A bullet-point listing of best practice usage.

6.10.7 Code Listing

A code listing and a terse description of the listing.

6.10.8 References

An bullet-point annotated reference list of primary sources and useful resources.

6.11 Bacterial Foraging Optimization Algorithm

The heading and alternate headings for the algorithm description.

6.11.1 Taxonomy

A small tree diagram showing related fields and algorithms.

6.11.2 Inspiration

A textual description of the inspiring system.

6.11.3 Metaphor

A textual description of the algorithm by analogy.

6.11.4 Strategy

A textual description of the information processing strategy.

6.11.5 Procedure

A pseudo code description of the algorithms procedure.

6.11.6 Heuristics

A bullet-point listing of best practice usage.

6.11.7 Code Listing

A code listing and a terse description of the listing.

6.11.8 References

An bullet-point annotated reference list of primary sources and useful resources.

6.12 Summary

todo

Chapter 7

Immune Algorithms

7.1 Overview

todo

7.2 Clonal Selection Algorithm

The heading and alternate headings for the algorithm description.

7.2.1 Taxonomy

A small tree diagram showing related fields and algorithms.

7.2.2 Inspiration

A textual description of the inspiring system.

7.2.3 Metaphor

A textual description of the algorithm by analogy.

7.2.4 Strategy

A textual description of the information processing strategy.

7.2.5 Procedure

A pseudo code description of the algorithms procedure.

7.2.6 Heuristics

A bullet-point listing of best practice usage.

7.2.7 Code Listing

A code listing and a terse description of the listing.

7.2.8 References

An bullet-point annotated reference list of primary sources and useful resources.

7.3 Negative Selection Algorithm

The heading and alternate headings for the algorithm description.

7.3.1 Taxonomy

A small tree diagram showing related fields and algorithms.

7.3.2 Inspiration

A textual description of the inspiring system.

7.3.3 Metaphor

A textual description of the algorithm by analogy.

7.3.4 Strategy

A textual description of the information processing strategy.

7.3.5 Procedure

A pseudo code description of the algorithms procedure.

7.3.6 Heuristics

A bullet-point listing of best practice usage.

7.3.7 Code Listing

A code listing and a terse description of the listing.

7.3.8 References

An bullet-point annotated reference list of primary sources and useful resources.

7.4 Artificial Immune Recognition System

The heading and alternate headings for the algorithm description.

7.4.1 Taxonomy

A small tree diagram showing related fields and algorithms.

7.4.2 Inspiration

A textual description of the inspiring system.

7.4.3 Metaphor

A textual description of the algorithm by analogy.

7.4.4 Strategy

A textual description of the information processing strategy.

7.4.5 Procedure

A pseudo code description of the algorithms procedure.

7.4.6 Heuristics

A bullet-point listing of best practice usage.

7.4.7 Code Listing

A code listing and a terse description of the listing.

7.4.8 References

An bullet-point annotated reference list of primary sources and useful resources.

7.5 Immune Network Algorithm

The heading and alternate headings for the algorithm description.

7.5.1 Taxonomy

A small tree diagram showing related fields and algorithms.

7.5.2 Inspiration

A textual description of the inspiring system.

7.5.3 Metaphor

A textual description of the algorithm by analogy.

7.5.4 Strategy

A textual description of the information processing strategy.

7.5.5 Procedure

A pseudo code description of the algorithms procedure.

7.5.6 Heuristics

A bullet-point listing of best practice usage.

7.5.7 Code Listing

A code listing and a terse description of the listing.

7.5.8 References

An bullet-point annotated reference list of primary sources and useful resources.

7.6 Dendritic Cell Algorithm

The heading and alternate headings for the algorithm description.

7.6.1 Taxonomy

A small tree diagram showing related fields and algorithms.

7.6.2 Inspiration

A textual description of the inspiring system.

7.6.3 Metaphor

A textual description of the algorithm by analogy.

7.6.4 Strategy

A textual description of the information processing strategy.

7.6.5 Procedure

A pseudo code description of the algorithms procedure.

7.6.6 Heuristics

A bullet-point listing of best practice usage.

7.6.7 Code Listing

A code listing and a terse description of the listing.

7.6.8 References

An bullet-point annotated reference list of primary sources and useful resources.

7.7 Summary

todo

Part III

Extensions

Chapter 8

Advanced Topics

A chapter focused on applying, testing, visualizing, saving results, and comparing algorithms. The meta concerns once an algorithm is selected for a given practical problem solving scenario.

8.1 Programming Paradigms

Algorithms can be implemented on many different programming paradigms. Take the GA for example and realize it using a bunch of different paradigms.

8.1.1 Procedural Programming

The GA under a procedural paradigm

8.1.2 Object-Oriented Programming

The GA under a object oriented paradigm. Strategy pattern. modular operators, etc.

8.1.3 Agent Oriented Programming

A GA under an agent oriented programming paradigm. not really suited. algorithm as an agent with goals?

8.1.4 Functional Programming

The GA under a functional paradigm. closure etc

8.1.5 Meta-Programming

A GA under meta programming. A DSL i guess.

8.1.6 Flow Programming

A GA under a data flow or pipeline model.

8.1.7 Map Reduce

A GA under a map reduce paradigm.

8.2 Devising New Algorithms

A methodology for devising new unconventional optimization algorithms...

8.2.1 Conceptual Framework for Bio-Inspired Algorithms

A generic methodology for devising new biologically inspired algorithms

8.2.2 Information Processing Methodology

An info processing centric approach to devising new algorithms

8.2.3 Investigation

small models, rigor

8.2.4 Communication

you need to effectively describe them, like as in this book! goal is to be known and used, make it open and usable by anyone. like open source, documented, common languages, benchmark problems, a website, lots of papers

8.3 Testing Algorithms

This section will focus on the problem that ‘adaptive systems work even when they are not implemented correctly’ (they work in-spite of the developer). Topics will include unit testing algorithms, system testing software, specific concerns when testing inspired algorithms, examples of testing algorithms with the ruby unit testing framework, examples of testing algorithms with rspec.

8.3.1 Types of Testing

unit, TDD, system, user acceptance, black box, white box

8.3.2 Algorithm Testing Methodology

testing is hard these systems ‘work’ even with bugs, hard to test present a methodology for testing - discrete unit tests, behavior testing

8.3.3 Example

develop and show tests for the GA

8.4 Visualizing Algorithms

This section will focus on the use of visualization as a low-fidelity form of system testing. Topics will include free visualization packages such as R, GNUPlot and Processing. Examples visualizing a decision surface, a functions response surface, and candidate solutions.

8.4.1 Visualizing

we can do it as a form of testing. research aid - view on a complex process, can observe, take notes, formulate hypothesis think of all the measures you can, than measure them

Offline Plots

examples?

Online Plots

examples?

8.4.2 Visualization Tools

can use lots of things, can use lots of things

8.4.3 Example

Visualize genes through time for a ga run, with fitness graphs, and plots of domain

8.5 Saving Algorithm Results

This section will focus on algorithms and techniques as a fallible means to an end and the need to maintain save results. Topics will include check-pointing, storage in a database, storage on the filesystem, and algorithm restarting. Examples will be given for database, filesystem checkpointing and algorithm restarting.

8.5.1 Check-pointing

algorithms crash and it sucks, need to be able to pickup where you left off

8.5.2 Share Results

make them public with papers and source code

8.5.3 Example

show an example of check pointing

8.6 Comparing Algorithms

This section will focus on comparing algorithm's based on the solutions they provide. Topics will include the use statistical hypothesis testing and free software such as R, algorithm parameter selection, distribution testing, distribution comparisons. Examples will be given for algorithm parameter selection, result distribution classification, and pair-wise result distribution comparison.

8.6.1 No Free Lunch

all same over all problems with no prior info

8.6.2 Benchmarking

standard problem instances what problems? what algorithms? what configurations what are you measuring? what are you comparing?

8.6.3 Statistical Hypothesis Testing

you need stats or you will be killed by Zed Shaw need stats to compare results

8.6.4 Example

genetic algorithm vs something, use R to compare

8.7 Summary

We learned lots of advanced topics, there are more.

Index

Adaptive Random Search, [27](#)

Blind Search, [24](#)

GRASP, [49](#)

Greedy Randomized Adaptive Search, [49](#)

Guided Local Search, [39](#)

Hill Climbing, [32](#)

Iterated Local Search, [35](#)

Random Mutation Hill Climbing, [32](#)

Random Search, [24](#)

Reactive Tabu Search, [65](#)

Scatter Search, [54](#)

Stochastic Hill Climbing, [32](#)

Taboo Search, [60](#)

Tabu Search, [60](#)

Variable Neighborhood Search, [44](#)

bibliography

Bibliography

- [1] S. Aaronson. NP-complete problems and physical reality. *ACM SIGACT News (COLUMN: Complexity theory)*, 36(1):30–52, 2005.
- [2] M. M. Ali, C. Storey, and A. Trn. Application of stochastic global optimization algorithms to practical problems. *Journal of Optimization Theory and Applications*, 95(3):545–563, 1997.
- [3] C. Andrieu, N. de Freitas, A. Doucet, and M. I. Jordan. An introduction to mcmc for machine learning. *Machine Learning*, 50:5–43, 2003.
- [4] Thomas Back. Optimal mutation rates in genetic search. In *Proceedings of the Fifth International Conference on Genetic Algorithms*, pages 2–9, 1993.
- [5] Thomas Bäck, D.B Fogel, and Z Michalewicz, editors. *Evolutionary Computation 1: Basic Algorithms and Operators*. IoP, 2000.
- [6] Thomas Bäck, D.B Fogel, and Z Michalewicz, editors. *Evolutionary Computation 2: Advanced Algorithms and Operations*. IoP, 2000.
- [7] J.F. Bard, T.A. Feo, and S. Holland. A GRASP for scheduling printed wiring board assembly. *I.I.E. Trans.*, 28:155–165, 1996.
- [8] R. Battiti and M. Brunato. *Handbook of Metaheuristics*, chapter Reactive Search Optimization: Learning while Optimizing. Springer Verlag, 2nd edition, 2009.
- [9] R. Battiti, M. Brunato, and F. Mascia. *Reactive Search and Intelligent Optimization*. Springer, 2008.
- [10] R. Battiti and M. Protasi. Reactive local search for the maximum clique problem. Technical Report TR-95-052, International Computer Science Institute, Berkeley, CA, 1995.
- [11] R. Battiti and M. Protasi. Reactive local search for the maximum clique problem. *Algorithmica*, 29(4):610–637, 2001.
- [12] R. Battiti and G. Tecchiolli. Parallel biased search for combinatorial optimization: genetic algorithms and tabu. *Microprocessors and Microsystems*, 16(7):351–367, 1992.

- [13] R. Battiti and G. Tecchiolli. The reactive tabu search. *ORSA Journal on Computing*, 6(2):126–140, 1994.
- [14] R. Battiti and G. Tecchiolli. Simulated annealing and tabu search in the long run: a comparison on qap tasks. *Computer and Mathematics with Applications*, 28(6):1–8, 1994.
- [15] R. Battiti and G. Tecchiolli. Local search with memory: Benchmarking rts. *Operations Research Spektrum*, 17(2/3):67–86, 1995.
- [16] R. Battiti and G. Tecchiolli. Training neural nets with the reactive tabu search. *IEEE Transactions on Neural Networks*, 6(5):1185–1200, 1995.
- [17] Roberto Battiti. Machine learning methods for parameter tuning in heuristics. In *5th DIMACS Challenge Workshop: Experimental Methodology Day*, 1996.
- [18] E. B. Baum. Towards practical “neural” computation for combinatorial optimization problems. In *AIP conference proceedings: Neural Networks for Computing*, pages 53–64, 1986.
- [19] J. Baxter. Local optima avoidance in depot location. *Journal of the Operational Research Society*, 32:815–819, 1981.
- [20] Janine M. Benyus. *Biomimicry: innovation inspired by nature*. Quill, 1998.
- [21] D. Bergemann and J. Valimaki. Bandit problems. Cowles Foundation Discussion Papers 1551, Cowles Foundation, Yale University, January 2006.
- [22] C. M. Bishop. *Neural Networks for Pattern Recognition*. Oxford University Press, 1995.
- [23] Christian Blum and Andrea Roli. Metaheuristics in combinatorial optimization: Overview and conceptual comparison. *ACM Computing Surveys (CSUR)*, 35(3):268–308, 2003.
- [24] Eric Bonabeau, Marco Dorigo, and Guy Theraulaz. *Swarm intelligence: from natural to artificial systems*. Oxford University Press US, 1999.
- [25] S. Boyd and L. Vandenberghe. *Convex Optimization*. Cambridge University Press, 2004.
- [26] Samuel H. Brooks. A discussion of random methods for seeking maxima. *Operations Research*, 6(2):244–251, 1958.
- [27] E. K. Burke, E. Hart, G. Kendall, J. Newall, P. Ross, and S. Schulenburg. *Handbook of Metaheuristics*, chapter Hyper-heuristics: An emerging direction in modern search technology, pages 457–474. Kluwer, 2003.

- [28] E. K. Burke, G. Kendall, and E. Soubeiga. A tabu-search hyper-heuristic for timetabling and rostering. *Journal of Heuristics*, 9(6):451–470, 2003.
- [29] Leandro N. De Castro and Fernando J. Von Zuben. *Recent developments in biologically inspired computing*. Idea Group Inc, 2005.
- [30] David Corne, Marco Dorigo, and Fred Glover. *New ideas in optimization*. McGraw-Hill, 1999.
- [31] D. Cvijovic and J. Klinowski. Taboo search: An approach to the multiple minima problem. *Science*, 267:664–666, 1995.
- [32] L. Davis. Bit-climbing, representational bias, and test suite design. In *Proceedings of the fourth international conference on genetic algorithms*, pages 18–23, 1991.
- [33] L. N. de Castro and F. J. Von Zuben. *Recent developments in biologically inspired computing*, chapter From biologically inspired computing to natural computing. Idea Group, 2005.
- [34] Leandro N. de Castro and Jonathan Timmis. *Artificial Immune Systems: A New Computational Intelligence Approach*. Springer, 2002.
- [35] Marco Dorigo and Thomas Stützle. *Ant colony optimization*. MIT Press, 2004.
- [36] Stefan Droste, Thomas Jansen, and Ingo Wegener. Upper and lower bounds for randomized search heuristics in black-box optimization. *Theory of Computing Systems*, 39(4):525–544, 2006.
- [37] Andries P. Engelbrecht. *Computational intelligence: an introduction*. John Wiley and Sons, second edition, 2007.
- [38] T. A. Feo and M. G. C. Resende. Greedy randomized adaptive search procedures. *Journal of Global Optimization*, 6:109–133, 1995.
- [39] T.A. Feo, J. Bard, and S. Holland. A grasp for scheduling printed wiring board assembly. Technical Report TX 78712-1063, Operations Research Group, Department of Mechanical Engineering, The University of Texas at Austin, 1993.
- [40] T.A. Feo and M.G.C. Resende. A probabilistic heuristic for a computationally difficult set covering problem. *Operations Research Letters*, 8:67–71, 1989.
- [41] T.A. Feo, K. Sarathy, and J. McGahan. A grasp for single machine scheduling with sequence dependent setup costs and linear delay penalties. Technical Report TX 78712-1063, Operations Research Group, Department of Mechanical Engineering, The University of Texas at Austin, 1994.
- [42] T.A. Feo, K. Venkatraman, and J.F. Bard. A GRASP for a difficult single machine scheduling problem. *Computers & Operations Research*, 18:635–643, 1991.

- [43] Thomas A. Feo, Kishore Sarathy, and John McGahan. A grasp for single machine scheduling with sequence dependent setup costs and linear delay penalties. *Computers & Operations Research*, 23(9):881–895, 1996.
- [44] P. Festa and M. G. C. Resende. *Essays and Surveys on Metaheuristics*, chapter GRASP: An annotated bibliography, pages 325–367. Kluwer Academic Publishers, 2002.
- [45] C.-N. Fiechter. A parallel tabu search algorithm for large traveling salesman problems. *Discrete Applied Mathematics*, 3(6):243–267, 1994.
- [46] David Flanagan and Yukihiro Matsumoto. *The Ruby Programming Language*. O’Reilly Media, 2008.
- [47] N. Forbes. Biologically inspired computing. *Computing in Science and Engineering*, 2(6):83–87, 2000.
- [48] Nancy Forbes. *Imitation of Life: How Biology Is Inspiring Computing*. The MIT Press, 2005.
- [49] Stephanie Forrest and Melanie Mitchell. Relative building-block fitness and the building-block hypothesis. In *Foundations of Genetic Algorithms 2*, pages 109–126. Morgan Kaufmann, 1993.
- [50] K. Fukunaga. *Introduction to Statistical Pattern Recognition*. Academic Press, 1990.
- [51] Michel Gendreau. *Handbook of Metaheuristics*, chapter 2: An Introduction to Tabu Search, pages 37–54. Springer, 2003.
- [52] F. Glover. Heuristics for integer programming using surrogate constraints. *Decision Sciences*, 8(1):156–166, 1977.
- [53] F. Glover. Future paths for integer programming and links to artificial intelligence. *Computers and Operations Research*, 13(5):533–549, 1986.
- [54] F. Glover. Tabu search – Part I. *ORSA Journal on Computing*, 1(3):190–206, 1989.
- [55] F Glover. Tabu search – Part II. *ORSA Journal on Computing*, 2(1):4–32, 1990.
- [56] F. Glover. Tabu search: A tutorial. *Interfaces*, 4:74–94, 1990.
- [57] F. Glover. Tabu search for nonlinear and parametric optimization (with links to genetic algorithms). *Discrete Applied Mathematics*, 49:231–255, 1994.
- [58] F. Glover. *Artificial Evolution*, chapter A Template For Scatter Search And Path Relinking, page 13. Springer, 1998.

- [59] F. Glover. *New Ideas in Optimization*, chapter Scatter search and path relinking, pages 297–316. McGraw-Hill Ltd., 1999.
- [60] F. Glover, M. Laguna, and R. Martí. Fundamentals of scatter search and path relinking. *Control and Cybernetics*, 39(3):653–684, 2000.
- [61] F. Glover, M. Laguna, and R. Martí. *Advances in Evolutionary Computation: Theory and Applications*, chapter Scatter Search, pages 519–537. Springer-Verlag, 2003.
- [62] F. Glover and C. McMillan. The general employee scheduling problem: an integration of ms and ai. *Computers and Operations Research*, 13(5):536–573, 1986.
- [63] Fred Glover and Gary A. Kochenberger. *Handbook of metaheuristics*. Springer, 2003.
- [64] Fred Glover and Eric Taillard. A user’s guide to tabu search. *Annals of Operations Research*, 41(1):1–28, 1993.
- [65] Fred W. Glover and Manuel Laguna. *Tabu Search*. Springer, 1998.
- [66] David Edward Goldberg. *Genetic algorithms in search, optimization, and machine learning*. Addison-Wesley, 1989.
- [67] I Guyon and A Elisseeff. An introduction to variable and feature selection. *Journal of Machine Learning Research*, 3:1157–1182, 2003.
- [68] P. Hansen and N. Mladenovic. *Meta-heuristics, Advances and trends in local search paradigms for optimization*, chapter An introduction to Variable neighborhood search, pages 433–458. Kluwer Academic Publishers, 1998.
- [69] P. Hansen and N. Mladenovic. *Handbook of Applied Optimization*, chapter Variable neighbourhood search, pages 221–234. Oxford University Press, 2002.
- [70] P. Hansen and N. Mladenovic. *Handbook of metaheuristics*, chapter 6: Variable Neighborhood Search, pages 145–184. Springer, 2003.
- [71] Pierre Hansen and Nenad Mladenovic. Variable neighborhood search: Principles and applications. *European Journal of Operational Research*, 130(3):449–467, 2001.
- [72] Pierre Hansen, Nenad Mladenovic, and Dionisio Perez-Britos. Variable neighborhood decomposition search. *Journal of Heuristics*, 7(4):1381–1231, 2001.
- [73] J.P. Hart and A.W. Shogan. Semi-greedy heuristics: An empirical study. *Operations Research Letters*, 6:107–114, 1987.
- [74] John Henry Holland. *Adaptation in natural and artificial systems: an introductory analysis with applications to biology, control, and artificial intelligence*. University of Michigan Press, 1975.

- [75] R. Horst, P. M. Pardalos, and N. V. Thoai. *Introduction to Global Optimization*. Kluwer Academic Publishers, 2nd edition, 2000.
- [76] D. S. Johnson. Local optimization and the travelling salesman problem. In *Proceedings of the 17th Colloquium on Automata, Languages, and Programming*, pages 446–461, 1990.
- [77] D. S. Johnson and L. A. McGeoch. *Local Search in Combinatorial Optimization*, chapter The travelling salesman problem: A case study in local optimization, pages 215–310. John Wiley & Sons, 1997.
- [78] A. Juels and M. Wattenberg. Stochastic hill climbing as a baseline method for evaluating genetic algorithms. Technical report, University of California, Berkeley, 1994.
- [79] Dean C. Karnopp. Random search techniques for optimization problems. *Automatica*, 1(2–3):111–121, 1963.
- [80] John Knox. Tabu search performance on the symmetric traveling salesman problem. *Computers & Operations Research*, 21(8):867–876, 1994.
- [81] John R. Koza. *Genetic programming IV: routine human-competitive machine intelligence*. Springer, 2003.
- [82] J. Kregting and R. C. White. Adaptive random search. Technical Report TH-Report 71-E-24, Eindhoven University of Technology, Eindhoven, Netherlands, 1971.
- [83] M. Kudo and J. Sklansky. Comparison of algorithms that select features for pattern classifiers. *Pattern Recognition*, 33:25–41, 2000.
- [84] O. Yu. Kul’chitskii. Random-search algorithm for extrema in functional space under conditions of partial uncertainty. *Cybernetics and Systems Analysis*, 12(5):794–801, 1976.
- [85] Manuel Laguna and Rafael Martí. *Scatter search: methodology and implementations in C*. Kluwer Academic Publishers, 2003.
- [86] L. T. Lau. *Guided Genetic Algorithm*. PhD thesis, Department of Computer Science, University of Essex, 1999.
- [87] T.L. Lau and E.P.K. Tsang. The guided genetic algorithm and its application to the general assignment problems. In *IEEE 10th International Conference on Tools with Artificial Intelligence (ICTAI’98)*, 1998.
- [88] R. M. Lewis, V. T., and M. W. Trosset. Direct search methods: then and now. *Journal of Computational and Applied Mathematics*, 124:191–207, 2000.

- [89] H. R. Lourenco, O. Martin, and T. Stützle. A beginners introduction to iterated local search. In *Proceedings 4th Metaheuristics International Conference (MIC2001)*, 2001.
- [90] George F. Luger and William A. Stubblefield. *Artificial Intelligence: Structures and Strategies for Complex Problem Solving*. Benjamin/Cummings Pub. Co., second edition, 1993.
- [91] Sean Luke. *Essentials of Metaheuristics*. (self-published) <http://cs.gmu.edu/~sean/book/metaheuristics>, 2009.
- [92] P. Marrow. Nature-inspired computing technology and applications. *BT Technology Journal*, 18(4):13–23, 2000.
- [93] Rafael Martí, Manuel Lagunab, and Fred Glover. Principles of scatter search. *European Journal of Operational Research*, 169(1):359–372, 2006.
- [94] O. Martin and S. W. Otto. Combining simulated annealing with local search heuristics. *Annals of Operations Research*, 63:57–75, 1996.
- [95] O. Martin, S. W. Otto, and E. W. Felten. Large-step markov chains for the traveling salesman problems. *Complex Systems*, 5(3):299–326, 1991.
- [96] S. F. Masri, G. A. Bekey, and F. B. Safford. Global optimization algorithm using adaptive random search. *Applied Mathematics and Computation*, 7(4):353–376, 1980.
- [97] Zbigniew Michalewicz and David B. Fogel. *How to solve it: modern heuristics*. Springer, 2004.
- [98] P. Mills. *Extensions to Guided Local Search*. PhD thesis, Department of Computer Science, University of Essex, 2002.
- [99] Patrick Mills, Edward Tsang, and John Ford. Applying an extended guided local search on the quadratic assignment problem. *Annals of Operations Research*, 118:121–135, 2003.
- [100] Melanie Mitchell and John H. Holland. When will a genetic algorithm outperform hill climbing? In *Proceedings of the 5th International Conference on Genetic Algorithms*. Morgan Kaufmann Publishers Inc., 1993.
- [101] N. Mladenovic. A variable neighborhood algorithm - a new metaheuristic for combinatorial optimization. In *Abstracts of papers presented at Optimization Days*, 1995.
- [102] N. Mladenovic and P. Hansen. Variable neighborhood search. *Computers & Operations Research*, 24(11):1097–1100, 1997.

- [103] Heinz Muhlenbein. Evolution in time and space - the parallel genetic algorithm. In *Foundations of Genetic Algorithms*, 1991.
- [104] Heinz Muhlenbein. How genetic algorithms really work: I. mutation and hillclimbing. In *Parallel Problem Solving from Nature 2*, pages 15–26, 1992.
- [105] C. H. Papadimitriou and K. Steiglitz. *Combinatorial Optimization: Algorithms and Complexity*. Courier Dover Publications, 1998.
- [106] P.M. Pardalos, L.S. Pitsoulis, and M.G.C. Resende. A parallel grasp implementation for the quadratic assignment problem. In *Parallel Algorithms for Irregularly Structured Problems (Irregular94)*, pages 111–130. Kluwer Academic Publishers, 1995.
- [107] R. Paton. *Computing With Biological Metaphors*, chapter Introduction to computing with biological metaphors, pages 1–8. Chapman & Hall, 1994.
- [108] G. Paun. Bio-inspired computing paradigms (natural computing). *Unconventional Programming Paradigms*, 3566:155–160, 2005.
- [109] W. Pedrycz. *Computational Intelligence: An Introduction*. CRC Press, 1997.
- [110] L. Pitsoulis and M. G. C. Resende. *Handbook of Applied Optimization*, chapter Greedy randomized adaptive search procedures, pages 168–181. Oxford University Press, 2002.
- [111] M. Prais and C.C. Ribeiro. Reactive grasp: An application to a matrix decomposition problem in tdma traffic assignment. *INFORMS Journal on Computing*, 12:164–176, 2000.
- [112] Helena Ramalhinho-Loureno, Olivier C. Martin, and Thomas Stützle. *Handbook of Metaheuristics*, chapter Iterated Local Search, pages 320–353. Springer, 2003.
- [113] L. A. Rastrigin. The convergence of the random search method in the extremal control of a many parameter system. *Automation and Remote Control*, 24:1337–1342, 1963.
- [114] M. G. C. Resende and C. C. Ribeiro. *Handbook of Metaheuristics*, chapter Greedy randomized adaptive search procedures, pages 219–249. Kluwer Academic Publishers, 2003.
- [115] H. Robbins. Some aspects of the sequential design of experiments. *Bull. Amer. Math. Soc.*, 58:527–535, 1952.
- [116] Stuart Russell and Peter Norvig. *Artificial Intelligence: A Modern Approach*. Prentice Hall, third edition, 2009.
- [117] Gnther Schrack and Mark Choit. Optimized relative step size random searches. *Mathematical Programming*, 10(1):230–244, 1976.

- [118] M. Schumer and K. Steiglitz. Adaptive step size random search. *IEEE Transactions on Automatic Control*, 13(3):270–276, 1968.
- [119] Yuhui Shi, editor. *Swarm intelligence*. Morgan Kaufmann, 2001.
- [120] David B. Skalak. Prototype and feature selection by sampling and random mutation hill climbing algorithms. In *Proceedings of the eleventh international conference on machine learning*, pages 293–301. Morgan Kaufmann, 1994.
- [121] A. Sloman. *Evolving Knowledge in Natural Science and Artificial Intelligence*, chapter Must intelligent systems be scruffy? Pitman, 1990.
- [122] F. J. Solis and J. B. Wets. Minimization by random search techniques. *Mathematics of Operations Research*, 6:19–30, 1981.
- [123] James C. Spall. *Introduction to stochastic search and optimization: estimation, simulation, and control*. John Wiley and Sons, 2003.
- [124] James C. Spall. *Handbook of computational statistics: concepts and methods*, chapter 6. Stochastic Optimization, pages 169–198. Springer, 2004.
- [125] T. Stützle. Applying iterated local search to the permutation flow shop problem. Technical Report AIDA9804, FG Intellektik, TU Darmstadt, 1998.
- [126] T. Stützle. Iterated local search for the quadratic assignment problem. Technical Report AIDA-99-03, FG Intellektik, FB Informatik, TU Darmstadt, 1999.
- [127] Thomas Stützle and Holger H. Hoos. Analyzing the run-time behaviour of iterated local search for the tsp. In *Proceedings III Metaheuristics International Conference*, 1999.
- [128] Thomas G. Stützle. *Local Search Algorithms for Combinatorial Problems: Analysis, Improvements, and New Applications*. PhD thesis, Darmstadt University of Technology, Department of Computer Science, 1998.
- [129] El-Ghazali Talbi. *Metaheuristics: From Design to Implementation*. John Wiley and Sons, 2009.
- [130] Dave Thomas, Chad Fowler, and Andy Hunt. *Programming Ruby: The Pragmatic Programmers' Guide*. Pragmatic Bookshelf, second edition, 2004.
- [131] E. P. K Tsang and C. J. Wang. A generic neural network approach for constraint satisfaction problems. In Taylor G, editor, *Neural network applications*, pages 12–22, 1992.
- [132] A. Trn, M.M. Ali, and S. Viitanen. Stochastic global optimization: Problem classes and solution techniques. *Journal of Global Optimization*, 14:437–447, 1999.

- [133] C Voudouris. *Guided local search for combinatorial optimisation problems*. PhD thesis, Department of Computer Science, University of Essex, Colchester, UK, July 1997.
- [134] C. Voudouris and E. P. K. Tsang. *Handbook of metaheuristics*, chapter 7: Guided Local Search, pages 185–218. Springer, 2003.
- [135] C. Voudouris and E.P.K. Tsang. Guided local search joins the elite in discrete optimisation. In *Proceedings, DIMACS Workshop on Constraint Programming and Large Scale Discrete Optimisation*, 1998.
- [136] Chris Voudouris and Edward Tsang. The tunneling algorithm for partial csps and combinatorial optimization problems. Technical Report CSM-213, Department of Computer Science, University of Essex, Colchester, C04 3SQ, UK, 1994.
- [137] Chris Voudouris and Edward Tsang. Function optimization using guided local search. Technical Report CSM-249, Department of Computer Science University of Essex Colchester, CO4 3SQ, UK, 1995.
- [138] Chris Voudouris and Edward Tsang. Guided local search. Technical Report CSM-247, Department of Computer Science, University of Essex, Colchester, C04 3SQ, UK, 1995.
- [139] Edward Tsang & Chris Voudouris. Fast local search and guided local search and their application to british telecoms workforce scheduling problem. Technical Report CSM-246, Department of Computer Science University of Essex Colchester CO4 3SQ, 1995.
- [140] C. J. Wang and E. P. K. Tsang. Solving constraint satisfaction problems using neural networks. In *Proc. Second International Conference on Artificial Neural Networks*, pages 295–299, November 18–20, 1991.
- [141] Thomas Weise. *Global Optimization Algorithms - Theory and Application*. Thomas Weise, 2009-06-26 edition, 2007.
- [142] R. C. White. A survey of random methods for parameter optimization. *Simulation*, 17(1):197–205, 1971.
- [143] D. H. Wolpert and W. G. Macready. No free lunch theorems for search. Technical report, Santa Fe Institute, Santa Fe, NM, USA, 1995.
- [144] D. H. Wolpert and W. G. Macready. No free lunch theorems for optimization. *IEEE Transactions on Evolutionary Computation*, 1(67):67–82, 1997.
- [145] Z. B. Zabinsky. *Stochastic adaptive search for global optimization*. Kluwer Academic Publishers, 2003.
- [146] Lotfi Asker Zadeh, George J. Klir, and Bo Yuan. *Fuzzy sets, fuzzy logic, and fuzzy systems: selected papers*. World Scientific, 1996.

- [147] A. A. Zhigljavsky. *Theory of Global Random Search*. Kluwer Academic, 1991.