

Learning Classifier System*

Jason Brownlee
jasonb@CleverAlgorithms.com
The Clever Algorithms Project
<http://www.CleverAlgorithms.com>

April 11, 2010
Technical Report: CA-TR-20100411-1

Abstract

The Clever Algorithms project aims to describe a large number of Artificial Intelligence algorithms in a complete, consistent, and centralized manner, to improve their general accessibility. The project makes use of a standardized algorithm description template that uses well-defined topics that motivate the collection of specific and useful information about each algorithm described. This report describes the Learning Classifier System algorithm using the standardized template.

Keywords: Clever, Algorithms, Description, Optimization, Learning, Classifier, System

1 Introduction

The Clever Algorithms project aims to describe a large number of algorithms from the fields of Computational Intelligence, Biologically Inspired Computation, and Metaheuristics in a complete, consistent and centralized manner [3]. The project requires all algorithms to be described using a standardized template that includes a fixed number of sections, each of which is motivated by the presentation of specific information about the technique [4]. This report describes the Learning Classifier System algorithm using the standardized template.

2 Name

Learning Classifier System, LCS

3 Taxonomy

The Learning Classifier System algorithm is both an instance of an Evolutionary Algorithm from the field of Evolutionary Computation and an instance of a Reinforcement Learning algorithm from Machine Learning. The Learning Classifier System is a theoretical system with a number of implementations. Two streams of classifier are the Pittsburgh-style that seeks to optimize whole classifier, and the Michigan-style that optimize responsive rulesets. The Michigan-style Learning Classifier is the most common and is comprised of two versions: the ZCS (zeroth-level classifier system) and the XCS (accuracy-based classifier system).

*© Copyright 2010 Jason Brownlee. Some Rights Reserved. This work is licensed under a Creative Commons Attribution-Noncommercial-Share Alike 2.5 Australia License.

4 Strategy

The objective of the Learning Classifier System algorithm is to optimize payoff based on exposure to stimuli from a problem-specific environment. This is achieved by managing credit assignment for those rules that prove useful and searching for new rules and new variations on existing rules using an evolutionary process.

5 Procedure

The actors of the system include detectors, messages, effectors, feedback, and classifiers. Detectors are used by the system to perceive the state of the environment. Messages are the discrete information packets passed from the detectors into the system. The system performs information processing on messages, and messages may directly result in actions in the environment. Effectors control the actions of the system on and within the environment. In addition to the system actively perceiving via its detections, it may also receive directed feedback from the environment (payoff). Classifiers are condition-action rules that provides a filter for messages. If a message satisfies the conditional part of the classifier, the action of the classifier triggers. Rules act as message processors. Messages are defined at a fixed length using a binary alphabet. A classifier is defined as a binary string with a ternary alphabet of 1, 0, #, where the # represents do not care (matching both a 1 or 0).

The processing loop for the Learning Classifier system is as follows: i) Messages from the environment are placed on the message list. ii) The conditions of each classifier are checked to see if they are satisfied by at least one message in the message list. iii) All classifiers that are satisfied participate in a competition, those that win post their action to the message list. iv) All messages directed to the effectors are executed (causing actions in the environment). v) All messages on the message list from the previous cycle are deleted (messages persist for a single cycle). The algorithm may be described in terms of the main processing loop and two sub-algorithms: a reinforcement learning algorithm such as the bucket brigade algorithm or Q-learning, and a genetic algorithm for optimization of the system. Algorithm 1 provides a pseudo-code listing of the high-level processing loop of the Learning Classifier System, specifically the XCS as described by Butz and Wilson [6].

6 Heuristics

The majority of the heuristics in this section are specific to the XCS Learning Classifier System as described by Butz and Wilson [6].

- Learning Classifier Systems are suited for problems with the following characteristics: perpetually novel events with large amounts of noise, continual, and real-time requirements for action, implicitly or inexactly defined goals, and sparse payoff or reinforcement obtainable only through long sequences of tasks.
- The learning rate β for a classifiers expected payoff, error and fitness are typically in the range $\in [0.1, 0.2]$.
- The frequency of running the genetic algorithm θ_{GA} should be in the range $\in [25, 50]$.
- The discount factor used in multi-step programs γ are typically in the around 0.71.
- The minimum error for whereby classifiers are considered to have equal accuracy ϵ_0 are typically 10% of the maximum reward.
- The probability of crossover in the genetic algorithm χ are typically in the range $\in [0.5, 1.0]$.

Algorithm 1: Pseudo Code for the Learning Classifier System algorithm.

Input: env
Output: Population

```
1 env ← InitializeEnvironment(env);
2 Population ← InitializePopulation();
3 ActionSett-1 ← 0;
4 Inputt-1 ← 0;
5 Rewardt-1 ← 0;
6 while ¬StopCondition() do
7   Inputt ← env;
8   Matchset ← GenerateMatchSet(Population, Inputt);
9   Prediction ← GeneratePrediction(Matchset);
10  Action ← SelectionAction(Prediction);
11  ActionSett ← GenerateActionSet(Action, Matchset);
12  Rewardt ← ExecuteAction(Action, env);
13  if ActionSett-1 ≠ 0 then
14    Payofft ← CalculatePayoff(Rewardt-1, Prediction);
15    PerformLearning(ActionSett-1, Payofft, Population);
16    RunGeneticAlgorithm(ActionSett-1, Inputt-1, Population);
17  end
18  if LastStepOfTask(env, Action) then
19    Payofft ← Rewardt;
20    PerformLearning(ActionSett, Payofft, Population);
21    RunGeneticAlgorithm(ActionSett, Inputt, Population);
22    ActionSett-1 ← 0;
23  else
24    ActionSett-1 ← ActionSett;
25    Inputt-1 ← Inputt;
26    Rewardt-1 ← Rewardt;
27  end
28 end
```

- The probability of mutating a single position in a classifier in the genetic algorithm μ is typically in the range $\in [0.01, 0.05]$.
- The experience threshold during classifier deletion θ_{del} is typically about 20.
- The experience threshold for a classifier during subsumption θ_{sub} is typically around 20.
- The initial values for a classifiers expected payoff p_1 , error ϵ_1 , and fitness f_1 are typically small and close to zero.
- The probability of selecting a random action for the purposes of exploration p_{exp} is typically close to 0.5.
- The minimum number of different actions that must be specified in a match set θ_{mna} is usually the total number of possible actions in the environment for the input.
- Subsumption should be used on problem domains that are known contain well defined rules for mapping inputs to outputs.

7 Code Listing

Listing 1 provides an example of the Learning Classifier System algorithm implemented in the Ruby Programming Language. The problem is an instance of a Boolean multiplexer called the 6-multiplexer. It can be described as a classification problem, where each of the 2^6 patterns of bits is associated with a boolean class $\{1, 0\}$. For this problem instance, the first two bits may be decoded as an address into the remaining four bits that specify the class (for example in 100010, ‘10’ decode to the index of ‘2’ in the remaining 4 bits making the class ‘1’). In propositional logic this problem instance may be described as $F = (\neg x_0)(\neg x_1)x_2 + (\neg x_0)x_1x_3 + x_0(\neg x_1)x_4 + x_0x_1x_5$. The algorithm is an instance of XCS based on the description provided by Butz and Wilson [6] with the parameters based on the application of XCS to Boolean multiplexer problems by Wilson [16, 17]. The population is grown as needed, and subsumption which would be appropriate for the Boolean multiplexer problem was not used for brevity. The multiplexer problem is a single step problem, so the complexities of delayed payoff are not required. A number of parameters were hard coded to recommended values, specifically: $\alpha = 0.1, v = 5, \delta = 0.1$ and $P_{\#} = \frac{1}{3}$.

```

1 def new_classifier(condition, action, gen)
2   other = {}
3   other[:condition], other[:action], other[:lasttime] = condition, action, gen
4   other[:prediction], other[:error], other[:fitness] = 0.00001, 0.00001, 0.00001
5   other[:experience], other[:setsize], other[:num] = 0.0, 1.0, 1.0
6   return other
7 end
8
9 def copy_classifier(parent)
10  copy = {}
11  parent.keys.each {|k| copy[k] = (parent[k].kind_of? String) ? ""+parent[k] : parent[k]}
12  copy[:num] = 1
13  copy[:experience] = 0.0
14  return copy
15 end
16
17 def generate_problem_string(length)
18  return (0...(length)).inject("") {|s,i| s+((rand<0.5) ? "1" : "0")}
19 end
20
21 def neg(bit)
22  return (bit==1) ? 0 : 1
23 end
24
25 def target_function(s)
26  ints = Array.new(s.length){|i| s[i].chr.to_i}
27  x0,x1,x2,x3,x4,x5 = ints
28  return neg(x0)*neg(x1)*x2 + neg(x0)*x1*x3 + x0*neg(x1)*x4 + x0*x1*x5
29 end
30
31 def calculate_deletion_vote(classifier, pop, del_thresh)
32  vote = classifier[:setsize] * classifier[:num]
33  avg_fit = pop.inject(0.0){|s,c| s+c[:fitness]}/pop.inject(0.0){|s,c| s+c[:num]}
34  derated = classifier[:fitness] / classifier[:num]
35  if classifier[:experience] > del_thresh and derated < 0.1 * avg_fit
36    vote *= avg_fit / derated
37  end
38  return vote
39 end
40
41 def delete_from_pop(pop, pop_size, del_thresh)
42  total = pop.inject(0) {|s,c| s+c[:num]}
43  return if total < pop_size
44  pop.each {|c| c[:dvote] = calculate_deletion_vote(c, pop, del_thresh)}
45  vote_sum = pop.inject(0.0) {|s,c| s+c[:dvote]}

```

```

46 point = rand() * vote_sum
47 vote_sum, index = 0.0, 0
48 pop.each_with_index do |c,i|
49   vote_sum += c[:dvote]
50   if vote_sum > point
51     index = i
52     break
53   end
54 end
55 if pop[index][:num] > 1
56   pop[index][:num] -= 1
57 else
58   pop.delete_at(index)
59 end
60 end
61
62 def generate_random_classifier(input, actions, gen)
63   condition = ""
64   input.each_char {|s| condition << ((rand<1.0/3.0) ? '#' : s)}
65   action = actions[rand(actions.length)]
66   return new_classifier(condition, action, gen)
67 end
68
69 def does_match(input, condition)
70   i = 0
71   condition.each_char do |c|
72     return false if c!='#' and c!=input[i].chr
73     i += 1
74   end
75   return true
76 end
77
78 def get_actions(pop)
79   return [] if pop.empty?
80   set = {}
81   pop.each do |classifier|
82     key = classifier[:action]
83     set[key] = 0 if set[key].nil?
84     set[key] += 1
85   end
86   return set.keys
87 end
88
89 def generate_match_set(input, pop, all_actions, gen, pop_size, del_thresh)
90   match_set = pop.select{|c| does_match(input, c[:condition])}
91   actions = get_actions(match_set)
92   while actions.length < all_actions.length do
93     remaining = all_actions - actions
94     classifier = generate_random_classifier(input, remaining, gen)
95     pop << classifier
96     match_set << classifier
97     delete_from_pop(pop, pop_size, del_thresh)
98     actions << classifier[:action]
99   end
100   return match_set
101 end
102
103 def generate_prediction(input, match_set)
104   prediction = {}
105   match_set.each do |classifier|
106     key = classifier[:action]
107     prediction[key] = {:sum=>0.0,:count=>0.0,:weight=>0.0} if prediction[key].nil?
108     prediction[key][:sum] += classifier[:prediction]*classifier[:fitness]

```

```

109     prediction[key][:count] += classifier[:fitness]
110 end
111 prediction.keys.each do |key|
112     prediction[key][:weight] = prediction[key][:sum] / prediction[key][:count]
113 end
114 return prediction
115 end
116
117 def select_action(prediction_array, p_explore)
118     keys = prediction_array.keys
119     return true, keys[rand(keys.length)] if rand() < p_explore
120     keys.sort!{|x,y| prediction_array[y][:weight] <=> prediction_array[x][:weight]}
121     return false, keys.first
122 end
123
124 def update_set(action_set, payoff, l_rate)
125     action_set.each do |c|
126         c[:experience] += 1.0
127         pdiff = payoff - c[:prediction]
128         c[:prediction] += (c[:experience] < 1.0 / l_rate) ? pdiff / c[:experience] : l_rate * pdiff
129         diff = pdiff.abs - c[:error]
130         c[:error] += (c[:experience] < 1.0 / l_rate) ? diff / c[:experience] : l_rate * diff
131         sum = action_set.inject(0.0) {|s, other| s + other[:num] - c[:setsize]}
132         c[:setsize] += (c[:experience] < 1.0 / l_rate) ? sum / c[:experience] : l_rate * sum
133     end
134 end
135
136 def update_fitness(action_set, min_error, l_rate)
137     sum = 0.0
138     accuracy = Array.new(action_set.length)
139     action_set.each_with_index do |c, i|
140         accuracy[i] = (c[:error] < min_error) ? 1.0 : 0.1 * (c[:error] / min_error) ** -5.0
141         sum += accuracy[i] * c[:num]
142     end
143     action_set.each_with_index do |c, i|
144         c[:fitness] += l_rate * (accuracy[i] * c[:num] / sum - c[:fitness])
145     end
146 end
147
148 def can_run_genetic_algorithm(action_set, gen, ga_freq)
149     total = action_set.inject(0.0) {|s, c| s + c[:lasttime] * c[:num]}
150     sum = action_set.inject(0.0) {|s, c| s + c[:num]}
151     if gen - (total / sum) > ga_freq
152         return true
153     end
154     return false
155 end
156
157 def select_parent(pop)
158     sum = pop.inject(0.0) {|s, c| s + c[:fitness]}
159     point = rand() * sum
160     sum = 0
161     pop.each do |c|
162         sum += c[:fitness]
163         return c if sum > point
164     end
165 end
166
167 def mutation(classifier, p_mut, action_set, input)
168     classifier[:condition].length.times do |i|
169         if rand() < p_mut
170             if classifier[:condition][i].chr == '#'
171                 classifier[:condition][i] = input[i]

```

```

172     else
173         classifier[:condition][i] = '#'
174     end
175 end
176 end
177 if rand() < p_mut
178     new_action = nil
179     begin
180         new_action = action_set[rand(action_set.length)]
181     end until new_action != classifier[:action]
182     classifier[:action] = new_action
183 end
184 end
185
186 def uniform_crossover(string1, string2)
187     rs = ""
188     string1.length.times do |i|
189         rs << ((rand()<0.5) ? string1[i] : string2[i])
190     end
191     return rs
192 end
193
194 def insert_in_pop(classifier, pop)
195     pop.each do |c|
196         if classifier[:condition]==c[:condition] and classifier[:action]==c[:action]
197             c[:num] += 1
198             return
199         end
200     end
201     pop << classifier
202 end
203
204 def crossover(c1, c2, p1, p2)
205     c1[:condition] = uniform_crossover(p1[:condition], p2[:condition])
206     c2[:condition] = uniform_crossover(p1[:condition], p2[:condition])
207     c1[:prediction] = (p1[:prediction]+p2[:prediction])/2.0
208     c1[:error] = 0.25*(p1[:error]+p2[:error])/2.0
209     c1[:fitness] = 0.1*(p1[:fitness]+p2[:fitness])/2.0
210     c2[:prediction] = c1[:prediction]
211     c2[:error] = c1[:error]
212     c2[:fitness] = c1[:fitness]
213 end
214
215 def run_genetic_algorithm(all_actions, pop, action_set, input, gen, p_cross, p_mut, pop_size,
    del_thresh)
216     p1, p2 = select_parent(action_set), select_parent(action_set)
217     c1, c2 = copy_classifier(p1), copy_classifier(p2)
218     crossover(c1, c2, p1, p2) if rand() < p_cross
219     [c1,c2].each do |c|
220         mutation(c, p_mut, all_actions, input)
221         insert_in_pop(c, pop)
222     end
223     delete_from_pop(pop, pop_size, del_thresh)
224 end
225
226 def search(length, pop_size, max_gens, all_actions, p_explore, l_rate, min_error, ga_freq,
    p_cross, p_mut, del_thresh)
227     pop, abs = [], 0
228     max_gens.times do |gen|
229         input = generate_problem_string(length)
230         match_set = generate_match_set(input, pop, all_actions, gen, pop_size, del_thresh)
231         prediction_array = generate_prediction(input, match_set)
232         explore, action = select_action(prediction_array, p_explore)

```

```

233     action_set = match_set.select{|c| c[:action]==action}
234     expected = target_function(input)
235     payoff = ((expected-action.to_i)==0) ? 300.0 : 1.0
236     abs += (expected - action.to_i).abs.to_f
237     update_set(action_set, payoff, l_rate)
238     update_fitness(action_set, min_error, l_rate)
239     if can_run_genetic_algorithm(action_set, gen, ga_freq)
240         action_set.each {|c| c[:lasttime] = gen}
241         run_genetic_algorithm(all_actions, pop, action_set, input, gen, p_cross, p_mut, pop_size,
            del_thresh)
242     end
243     if (gen+1).modulo(50)==0
244         puts " >gen=#{gen+1} classifiers=#{pop.size}, error=#{abs.to_i}/50 (#{(abs/50*100)}%)"
245         abs = 0
246     end
247 end
248 return pop
249 end
250
251 max_gens, length, pop_size = 5000, 6, 150
252 all_actions = ['0', '1']
253 l_rate, min_error = 0.2, 0.01
254 p_explore, p_cross, p_mut = 0.10, 0.80, 0.04
255 ga_freq, del_thresh = 50, 20
256
257 pop = search(length, pop_size, max_gens, all_actions, p_explore, l_rate, min_error, ga_freq,
    p_cross, p_mut, del_thresh)
258 puts "done! Solution: classifiers=#{pop.size}"

```

Listing 1: Learning Classifier System algorithm in the Ruby Programming Language

8 References

8.1 Primary Sources

Early ideas on the theory of Learning Classifier Systems were proposed by Holland [7, 10], culminating in a standardized presentation a few years later [8]. A number of implementations of the theoretical system were investigated, although a taxonomy of the two main streams was proposed by De Jong [12]: 1) Pittsburgh-style proposed by Smith [14, 13] and 2) Holland-style or Michigan-style Learning classifiers that are further comprised of the Zeroth-level classifier (ZCS) [15] and the accuracy-based classifier (XCS) [16].

8.2 Learn More

Booker, Goldberg, and Holland provide a classical introduction to Learning Classifier Systems including an overview of the state of the field and the algorithm in detail [1]. Wilson and Goldberg also provide a classical introduction and review of the approach, although take a more critical stance [18]. Holmes, et al. provide a contemporary review of the field focusing both on the approach and application areas to which the approach has been demonstrated successfully [11]. Lanzi, Stolzmann, and Wilson provide a seminal book in the field as a collection of papers covering the basics, advanced topics, and demonstration applications. A particular highlight from this book is the first section that provides a concise description of Learning Classifier Systems by many leaders and major contributors to the field [9], providing rare insight. Another paper from this book by Lanzi and Riolo provides a detailed review of the development of the approach as it matured throughout the 1990s. Bull and Kovacs a second book introductory book to the field focusing on the theory of the approach and its practical application [5].

9 Conclusions

This report described the Learning Classifier System as a machine learning technique using both reinforcement learning and genetic algorithms. The content for this report based based on a previous work on the Learning Classifier System by this author [2].

10 Contribute

Found a typo in the content or a bug in the source code? Are you an expert in this technique and know some facts that could improve the algorithm description for all? Do you want to get that warm feeling from contributing to an open source project? Do you want to see your name as an acknowledgment in print?

Two pillars of this effort are i) that the best domain experts are people outside of the project, and ii) that this work is wrong by default. Please help to make this work less wrong by emailing the author ‘Jason Brownlee’ at jasonb@CleverAlgorithms.com or visit the project website at <http://www.CleverAlgorithms.com>.

References

- [1] L. B. Booker, D. E. Goldberg, and J. H. Holland. Classifier systems and genetic algorithms. *Artificial Intelligence*, 40:235–282, 1989.
- [2] Jason Brownlee. Learning classifier systems. Technical Report 070514A, Complex Intelligent Systems Laboratory, Centre for Information Technology Research, Faculty of Information and Communication Technologies, Swinburne University of Technology, May 2007.
- [3] Jason Brownlee. The clever algorithms project: Overview. Technical Report CA-TR-20100105-1, The Clever Algorithms Project <http://www.CleverAlgorithms.com>, January 2010.
- [4] Jason Brownlee. A template for standardized algorithm descriptions. Technical Report CA-TR-20100107-1, The Clever Algorithms Project <http://www.CleverAlgorithms.com>, January 2010.
- [5] Larry Bull and Tim Kovacs. *Foundations of learning classifier systems*. Springer, 2005.
- [6] M. V. Butz and S. W. Wilson. An algorithmic description of xcs. *Journal of Soft Computing*, 6(3–4):144–153, 2002.
- [7] J. H. Holland. *Progress in Theoretical Biology IV*, chapter Adaptation, pages 263–293. Academic Press, 1976.
- [8] J. H. Holland. Adaptive algorithms for discovering and using general patterns in growing knowledge-bases. *International Journal of Policy Analysis and Information Systems*, 4:217–240, 1980.
- [9] John H. Holland, Lashon B. Booker, Marco Colombetti, Marco Dorigo, David E. Goldberg, Stephanie Forrest, Rick L. Riolo, Robert E. Smith, Pier Luca Lanzi, Wolfgang Stolzmann, and Stewart W. Wilson. *Learning classifier systems: from foundations to applications*, chapter What is a learning classifier system?, pages 3–32. Springer, 2000.
- [10] John H. Holland and Judith S. Reitman. Cognitive systems based on adaptive algorithms. *ACM SIGART Bulletin*, 63:49, 1977.

- [11] J. H. Holmes, P. L. Lanzi, W. Stolzmann, and S. W. Wilson. Learning classifier systems: New models, successful applications. *Information Processing Letters*, 82:23–30, 2002.
- [12] Kenneth De Jong. Learning with genetic algorithms: An overview. *Machine Learning*, 3:121–138, 1988.
- [13] S. Smith. Flexible learning of problem solving heuristics through adaptive search. In *Proceedings 8th International Joint Conference on Artificial Intelligence*, pages 422–425, 1983.
- [14] Stephen Frederick Smith. *A learning system based on genetic adaptive algorithms*. PhD thesis, Department of Computer Science, University of Pittsburgh, 1980.
- [15] S. W. Wilson. Zcs: A zeroth level classifier systems. *Evolutionary Computation*, 2:1–18, 1994.
- [16] S. W. Wilson. Classifier fitness based on accuracy. *Evolutionary Computation*, 3:149–175, 1995.
- [17] S. W. Wilson. Generalization in the xcs classifier system. In *Genetic Programming 1998: Proceedings of the Third Annual Conference*, pages 665–674. Morgan Kaufmann, 1998.
- [18] Stewart W. Wilson and David E. Goldberg. A critical review of classifier systems. In *Proceedings of the third international conference on Genetic algorithms*, pages 244–255, 1989.